

Learned representations and kernel corrections for radiative-transfer emulation

Yitzchak Shmalo*

September 2026

Abstract

We compare neural networks, exact Matérn regression and their combinations on a six-dimensional radiative-transfer table over the 285-band EMIT grid. Across ten random splits, with validation-based selection and 18,884 training samples, an input-scaled kernel reaches 0.095% mean relative radiance error against 0.76% for cubic regression and 0.39% for a neural network; residual correction takes the network to 0.14%, and a kernel on wider-network features to 0.079%. The forward gain does not carry to the inverse problem: a convex stack of all heads raises the mean within-split 95th-percentile reflectance error from 7.57 to 22.71 percentage points. We prove that the reflectance error of a physically constrained retrieval is bounded, with constant one, by a transmission-weighted component error divided by the transmission, that the loss of rate where the transmission vanishes is unavoidable, and that it disappears when the conditional component error falls with the transmission. Under the protocol of a published multi-fidelity benchmark the residual-corrected network improves on the published values in all four metrics on both of its splits.

1 Introduction

Atmospheric correction for imaging spectroscopy requires repeated evaluations of a radiative-transfer model. Learned emulators reduce this computational cost by approximating the map from atmospheric and viewing conditions to spectral coefficients. Neural emulators have been incorporated into reflectance-retrieval systems [1, 2], while Gaussian-process methods provide accurate surrogates for the OCO-2 forward model [3]. Physics-guided multi-fidelity models offer another approach [4]. These applications raise a practical question: when is it preferable to fit a kernel directly on the atmospheric state, learn a neural representation, or combine both?

Kernel methods are competitive with neural architectures on a range of operator-learning benchmarks [5]. Networks such as DeepONet and Fourier neural operators learn representations for high-dimensional inputs and outputs [6, 7]; kernel regression instead specifies a metric and solves a regularized interpolation problem [8]. Their combination can take several forms. Deep kernel learning jointly trains a representation and a kernel [9], and kernel flows can regularize intermediate network layers [10]. Here we use the construction of [11]: a trained network supplies either a mean whose residual is fitted by a kernel, or fixed features on which a kernel regresses the target. We compare these predictors with direct kernel regression, polynomial regression, coordinatewise selection and convex stacking.

*Einstein Institute of Mathematics, The Hebrew University of Jerusalem, Jerusalem, Israel.
yitzchak.shmalo@gmail.com.

The main dataset contains 23,313 radiative-transfer evaluations over the 285-band EMIT wavelength grid [12]. Its six-dimensional input and low-rank spectral output make exact kernel regression feasible at the full training size. Ten random partitions provide matched comparisons of the model families, followed by experiments varying training size, output rank, network width and kernel tuning. Input-specific length scales substantially improve the direct kernel. Wider-network features further reduce radiance error, even though the corresponding network has a larger standalone error under the tested optimizer and training budget.

The forward and inverse objectives give different model rankings. A convex stack slightly reduces the wide feature kernel’s mean radiance error, but increases its 95th-percentile reflectance error at every split. We derive exact finite-perturbation identities that relate this distinction to the denominator of the reflectance inverse. We also give an exact error-alignment criterion for residual correction and separate its kernel-regression error from principal-component truncation. These calculations complement the empirical comparison by identifying the quantities required for improvement in a specified metric.

Section 2 places the constructions used here beside the literature they come from. Section 3 describes the data, evaluation protocol and inversion stability. Sections 4–6 present the models, results and ablations. Section 7 analyses residual correction. Section 8 proves the transfer of component error into retrieval error with constant one, the conditional error-profile theorem that decides the rate, and the same-data perturbation bound for learned corrections; Sections 9 to 12 measure the quantities those results name, replicate the three-member feature kernel, compare combination rules, and place the corpus beside a published multi-fidelity benchmark under its own protocol. Sections 13 and 14 compare the findings with OCO-2 and ten further emulation configurations. The study concerns interpolation of a fixed radiative-transfer table; the implications for retrieval and computational cost are discussed in Section 15.

2 Related work

Radiative transfer and its surrogates. The codes being replaced are physical: 6S for the solar spectrum [13], libRadtran for the band calculations behind most lookup tables [14, 15], and MODTRAN inside the atmospheric-correction tables used here [16, 2]. What is done with those tables is optimal estimation in the sense of [17], and the operational carbon-dioxide product of OCO-2 [18, 19] rests on the forward model and inverse method of [20, 21]. Surrogates enter because the forward model is the cost. Neural emulation of longwave radiation goes back to [22] and was carried into climate models by [23]; later work measures how accurate such emulators can be made [24, 25] and how learned parameterizations behave when the host model is stepped forward [26, 27]. In imaging spectroscopy and land-surface inference the same idea appears as emulator toolboxes for radiative-transfer models [28, 29], Gaussian-process emulation for parameter retrieval [30], and the wider use of Gaussian processes in earth observation [31]; [1] places an emulator inside a reflectance retrieval and [3] fits a cross-validated Gaussian process to the OCO-2 forward model. Replacing an expensive simulator by a statistical model is a classical problem [32], as is using a cheap model to inform an expensive one [33, 34, 35], surveyed for propagation and inference in [36]; the network of [4] is a multi-fidelity instance on the corpus rescored in Section 12. The measurement requirements behind the tables used here are those of the EMIT mission [12] and of the surface-imaging missions that follow it [37].

Kernel regression. The estimators are classical: interpolation in a reproducing-kernel Hilbert space [38], smoothing splines and the representer theorem [39, 40, 41], ridge regularization [42] read as spectral filtering of an ill-posed problem [43], and the Matérn family, whose smoothness index fixes the differentiability of the sample paths [44, 8]. Rates under source and capacity conditions are given in [45, 46, 47, 48]; the textbook accounts are [49, 50, 51]. Each component here is a vector of principal coefficients, so the multi-output formulation is the relevant one [52, 53]. Kernels also serve as solvers rather than regressors: Bayesian numerical homogenization [54], operator-adapted wavelets [55] and the Gaussian-process treatment of nonlinear partial differential equations [56] use the same estimator, and cross-validating the kernel itself gives the kernel-flow family [57, 58, 10].

Cost. An exact solve is cubic in the number of training rows, which has produced a long line of approximations: the Nyström method and its analysis [59, 60, 61], random features [62] and their operator-learning counterpart [63], preconditioned solvers at large scale [64, 65], inducing-point and variational Gaussian processes [66, 67, 68], and structured interpolation with hardware-aware linear algebra [69, 70]. At $n = 18,884$ the exact factorization is affordable, so approximation enters this study only as a perturbation whose size is measured (Section 9).

Choosing and combining kernels. A kernel can be built rather than picked: spectral mixtures [71], additive structure [72], compositional search over a grammar [73], and per-coordinate relevance fitted by marginal likelihood [8] or by cross-validation [57]. Several kernels can be combined instead of one being chosen, by semidefinite programming [74], by the algorithms surveyed in [75], or by alignment with the target Gram matrix [76, 77]. Combining predictors rather than kernels is older still [78, 79, 80]. Every selection made below is made on held-out rows [81, 82, 83], and Section 8.2 bounds what that selection costs.

Networks, and the kernels they induce. Feedforward networks are universal approximators [84]; those trained here use the standard adaptive optimizer [85] with decoupled weight decay [86] and early stopping [87], two of the three regularizers compared in Section 10, the third being dropout [88]. In the infinite-width limit a network is a Gaussian process [89, 90, 91] whose kernel is the arc-cosine kernel of [92], and gradient descent on a wide network follows the neural tangent kernel [93, 94]; both enter the kernel comparison of Section 10 as heads. The limit is not the whole story, since a finite network learns its representation [95, 96], which is what deep kernel learning [9] and the construction of [11] exploit. Kolmogorov–Arnold networks [97, 98] are the architecture behind the multi-fidelity model of [4]. Operator learning supplies the other comparison: DeepONet [6], Fourier neural operators [7] and the general formulation [99], with fair-comparison protocols [100], physics-informed training [101], the cost-accuracy trade-off [102], kernel competitors [5], and the benchmark collections [103, 104, 105] from which Section 14 draws.

3 Data and protocol

The data are $n = 23,313$ tabulated radiative-transfer evaluations. Each input $x \in \mathbb{R}^6$ collects the atmospheric and viewing state; the coordinate ranges are consistent with the standard lookup-table axes for imaging-spectroscopy atmospheric correction [2] — an aerosol optical depth in $[0.05, 1.2]$, a surface elevation in $[-0.5, 6]$ km, a water-vapor column in $[0.05, 7]$ cm, a relative

azimuth in $[0, 180]^\circ$, a solar zenith angle in $[0, 65]^\circ$, and a view zenith angle within 30° of nadir. Each output is a set of four spectral components on the same 285-point EMIT wavelength grid spanning 381–2493 nm: path radiance Y_1 , direct flux Y_2 , diffuse flux Y_3 , and spherical albedo Y_4 (Figure 1). The quantities retrieved downstream are the radiance at a test surface reflectance ρ ,

$$L(\rho) = Y_1 + \rho \frac{Y_2 + Y_3}{1 - \rho Y_4}, \quad (1)$$

and the reflectance recovered by inverting (1) with predicted components at the true radiance,

$$\hat{\rho} = \frac{L - \hat{Y}_1}{\hat{Y}_2 + \hat{Y}_3 + \hat{Y}_4(L - \hat{Y}_1)}, \quad (2)$$

both evaluated at $\rho = 0.7$.

The protocol is the same for every family. The table is split into a test block of 2,331 rows and a remaining block of 20,982; 2,098 of the latter are held out for validation, leaving 18,884 for fitting. The procedure is repeated with seeds 101–110. Standardization and principal components are fitted only on each split’s training rows. Hyperparameters, checkpoints, coordinatewise selections and combination weights use that split’s validation rows. Each fitted family is then evaluated on that split’s test rows. Unless stated otherwise, tables report the mean and sample standard deviation over the ten splits.

The model families and ablations were chosen during exploratory analysis of the same table, including inspection of test-set results. Repeated partitions therefore measure sensitivity to the split within this table, rather than independent external validation. Their test sets overlap, and their standard deviations are descriptive, not confidence intervals from ten independent datasets. Section 5.1 reports paired differences on common splits. Validation rows are shared by early stopping, tuning and model combination. Evaluation uses the single surface reflectance $\rho = 0.7$; other reflectances and atmospheric states outside the table are not tested.

Three empirical structural diagnostics motivate the model comparison.

The outputs are low rank. On standardized training outputs, 99.99% of the variance of every component lies in the first 11–25 principal components, and 99.9999% in at most 44 (Figure 2, left); the minimum adjacent-band correlation across the four components is 0.994952 on the training rows. All regressors below act on the first 64 principal coefficients, whose training unexplained-variance fraction is below 10^{-6} , except the networks, which are trained directly on the 285 standardized bands. A small aggregate variance tail is not a uniform relative-error guarantee in physical units; Section 7 states the projection error explicitly.

The measured input relevance is anisotropic. Shuffling the aerosol coordinate reduces a linear model’s validation R^2 by 0.89; the corresponding reduction is 0.54 for solar zenith and approximately 0.1 for elevation and water vapor. The two remaining angles have smaller effects (Figure 2, right). These validation diagnostics motivate input-specific kernel scales. They measure sensitivity of the fitted linear model, rather than identifying an intrinsic dimension of the regression problem.

Low-degree surrogates are accurate. Linear regression on the standardized state achieves $R^2 = 0.83$, compared with 0.74 for distance-weighted five-nearest-neighbors. Cubic polynomial ridge regression has mean relative component error of approximately two percent (Table 1). Together with the spectral compression, these results motivate smooth low-dimensional regression as a baseline.

Reflectance inversion requires additional care. In deep absorption windows, the transmitted flux $Y_2 + Y_3$ can be as small as 10^{-21} . At a fully opaque band, the radiance is independent of




Figure 1: Example output curves across the EMIT wavelength grid, five samples per panel. The curves are smooth, share the deep water-vapor absorption features near 1380 and 1900 nm, and differ mainly by smooth rescalings along the state axes.

surface reflectance and the inverse is non-identifiable. An auxiliary round-trip calculation on 2,332 states gives reflectance RMSE 0.0054 even with exact tabulated components. Its error is concentrated in 1,652 of 664,620 sample-band entries, whose inverse denominators are below 2×10^{-13} . All returned inverses are finite. This is a floating-point conditioning diagnostic, not a lower bound on emulator error. We report both the center and the tails of the reflectance-error distribution. Non-finite inverses are assigned zero by the scoring routine; their frequency is recorded separately.

Differentiating (2) separates the components. Write $D = \hat{Y}_2 + \hat{Y}_3 + \hat{Y}_4(L - \hat{Y}_1)$ for its denominator. Then

$$\frac{\partial \hat{\rho}}{\partial \hat{Y}_4} = -\hat{\rho}^2, \quad \frac{\partial \hat{\rho}}{\partial \hat{Y}_2} = \frac{\partial \hat{\rho}}{\partial \hat{Y}_3} = -\frac{\hat{\rho}}{D}, \quad \frac{\partial \hat{\rho}}{\partial \hat{Y}_1} = -\frac{\hat{Y}_2 + \hat{Y}_3}{D^2}. \quad (3)$$

At exact identifiable components, $\hat{\rho} = \rho = 0.7$, so the albedo derivative is -0.49 , independently of the positive transmitted flux. For general predictions it equals $-\hat{\rho}^2$ and can be large. The denominator factors in (3) make path-radiance and flux errors particularly important near opaque bands. Section 5.1 reports the resulting all-band error tails.

3.1 Finite-perturbation control of the inversion

The derivatives in (3) describe local conditioning. The following elementary identities give a finite-error test and distinguish forward accuracy from inversion stability. For one sample and




Figure 2: Left: explained-variance spectra of the standardized outputs. Every component concentrates 99.99% of its variance in 11–25 principal components out of 285 bands. Right: drop in linear validation R^2 when a single standardized input coordinate is shuffled. Aerosol optical depth and solar zenith angle dominate; the azimuth and view-zenith coordinates are almost inert at first order.

band, write $a = Y_1$, $t = Y_2 + Y_3$, $s = Y_4$ and $L = a + \rho t / (1 - \rho s)$. Set $e_a = \hat{a} - a$, $e_t = \hat{t} - t$, $e_s = \hat{s} - s$.

Proposition 1 (Conditional radiance and reflectance stability). *Suppose $0 \leq \rho \leq \rho_{\max}$, $t > 0$ and $q = 1 - \rho s > 0$. If also $\hat{q} = 1 - \rho \hat{s} > 0$, the forward error satisfies the exact identity*

$$\hat{L} - L = e_a + \frac{\rho e_t}{\hat{q}} + \frac{\rho^2 t e_s}{q \hat{q}}. \quad (4)$$

In particular, if $q, \hat{q} \geq \kappa > 0$,

$$|\hat{L} - L| \leq |e_a| + \rho_{\max} |e_t| / \kappa + \rho_{\max}^2 |t e_s| / \kappa^2.$$

For inversion at the true radiance, put $u = L - a = \rho t / q$, $D = t / q$, and $\hat{D} = \hat{t} + \hat{s}(L - \hat{a})$. Whenever $\hat{D} \neq 0$,

$$\hat{\rho} - \rho = -\frac{(1 - \rho \hat{s})e_a + \rho e_t + \rho u e_s}{\hat{D}}, \quad \hat{D} - D = e_t + u e_s - \hat{s} e_a. \quad (5)$$

If $B = |e_t| + |u e_s| + |\hat{s} e_a| \leq \eta D$ for some $0 \leq \eta < 1$, then $\hat{D} \geq (1 - \eta)D > 0$ and

$$|\hat{\rho} - \rho| \leq \frac{|1 - \rho \hat{s}| |e_a| + \rho |e_t| + \rho |u e_s|}{(1 - \eta)D}. \quad (6)$$

Proof. Subtract the two forward maps and use $1/\hat{q} - 1/q = \rho e_s / (q \hat{q})$ to obtain (4). Since $u = \rho D$, $\hat{D} = D + e_t + u e_s - \hat{s} e_a$ and $\hat{\rho} - \rho = (u - e_a - \rho \hat{D}) / \hat{D}$, which gives (5). The triangle inequality yields $|\hat{D} - D| \leq B$ and then (6). $\square$

When $\hat{q} \neq 0$, the same calculation gives the exact forward-inverse link

$$\hat{\rho} - \rho = -\frac{\hat{q}}{\hat{D}} (\hat{L} - L). \quad (7)$$

Consequently, on a common set where $|\hat{q}| \leq Q$ and $|\hat{D}| \geq \delta > 0$, the corresponding unweighted L^2 errors obey $\|\hat{\rho} - \rho\|_2 \leq (Q/\delta)\|\hat{L} - L\|_2$. Without the denominator condition there is no uniform implication from small spectral relative radiance error to small reflectance error. For example, take two bands with $a = s = 0$, $t = (1, \varepsilon)$ and $0 < \rho < 1$. A predictor with $\hat{t} = t$, $\hat{s} = 0$ and $\hat{a} = (0, c\varepsilon)$, where $0 < c < \rho$, has relative spectral radiance error $c\varepsilon/(\rho\sqrt{1+\varepsilon^2}) \rightarrow 0$ but retrieves $(\rho, \rho - c)$, so its reflectance RMSE remains $c/\sqrt{2}$. Every denominator is positive for $\varepsilon > 0$ and both retrieved reflectances remain physical. This example isolates the loss of inversion stability as the second-band flux vanishes.

The forward bound need not deteriorate when t tends to zero, whereas the inverse bound requires error control relative to $D = t/q$. Thus a small spectral radiance error does not certify stable reflectance recovery. At $t = 0$, $L = a$ is independent of ρ wherever the forward formula is defined, so reflectance is unidentifiable from this band. The identities are exact-arithmetic statements; floating-point cancellation in $L - a$ can further affect an implemented inverse.

For empirical comparison, a common application passband and conditioning threshold should be fixed before model evaluation. The utility `code/conditioned_reflectance.py` uses a positive flux scale determined on training data and a common truth-based mask, which is available in simulation benchmarks. It reports retained coverage, non-finite inversions and denominator failures separately. Masked errors require sample-level arrays and cannot be recovered from aggregate summaries. The EMIT comparisons in Section 5.1 therefore use the all-band statistics.

4 Models

All models are trained per component and evaluated in physical units after undoing the standardization; unless stated otherwise they regress the 64 leading principal coefficients of the standardized component, with the training variance criterion and projection qualification of Section 3. Every hyperparameter named below is selected on the validation set of the seed at hand. All experiments use CPU computation.

Reference family. Cubic polynomial ridge regression [42], acting on all monomials of degree at most three in the six standardized inputs, with the ridge weight selected from $\{10^{-10}, 10^{-8}, 10^{-6}, 10^{-4}, 10^{-2}\}$. It provides a low-cost, low-degree reference; its empirical error alone does not prove a smoothness class for the underlying function.

Exact kernel ridge regression. Three constructions, all Matérn kernels [44] on the standardized inputs with the length scale expressed as a multiple of the median pairwise distance of the fit set. The listed nugget is γ in $(K + n\gamma I)\alpha = Z$, not the unscaled diagonal term. The subset baseline is a Matérn-5/2 kernel fitted on a uniformly drawn subset of 4,000 training points — the landmark choice of a Nyström approximation [59, 60] — one eigendecomposition per length scale, multiplier in $\{1, 2\}$, nugget in $\{10^{-8}, 10^{-6}, 10^{-4}, 10^{-2}\}$. The second is the exact solve on every training row: smoothness $\nu \in \{3/2, 5/2\}$, multiplier in $\{0.5, 1, 2, 4\}$ and the same nugget grid, tuned on validation from a 6,000-row training subsample and the winner refit on all 18,884 rows by one Cholesky factorization [106]. The third gives each input its own length scale, an automatic relevance determination [89, 8]: starting from the isotropic winner, a coordinate search over per-input multipliers in $\{1/4, 1/2, 1, 2, 4\}$, two sweeps, scored on validation, after which the smoothness, scale and nugget are re-tuned at the found metric and the solve is

refit on every row. The third construction allows the kernel to adapt to the input anisotropy measured in Section 3.

Fully connected networks. Three ELU hidden layers of 512 units per component, trained on the 285 standardized bands with AdamW [85, 86] at learning rate 10^{-3} , cosine decay, batch size 1024, weight decay 10^{-6} , up to 150 epochs with early stopping on validation loss [87] (patience 25). Five networks are trained per component and seed from different initializations; the table reports the first and the mean of the five.

Kernel corrections of the neural mean. The construction of [11], with the exact kernel stages fitted on every training row. (i) *Residual correction*: an exact Matérn regression, on the standardized input space, of the network’s training residuals in coefficient space, added back to the network’s projection into that space, with smoothness, scale and nugget selected by the validation error of the corrected prediction; applied to the single network and to the five-network mean. (ii) *Kernel on features*: the same exact regression with the input metric replaced by the network’s last hidden layer, standardized, giving a kernel metric in the fitted feature representation. This is the post-hoc counterpart of deep kernel learning [9], in which the representation and the kernel are trained together, and of the kernel-flow regularization of a network’s inner layers [10]; here the representation is fixed first and the kernel is solved exactly on it. (iii) *Per-coordinate selection*: for each of the 64 output coefficients independently, pick on validation among the exact kernel, the network, the ensemble, the two corrected predictors and the feature kernel, and assemble the test prediction from the winners. (iv) *Convex stack*: per component, a convex combination of the same six heads with weights fitted on the validation points in the physical relative metric, in the manner of stacked generalization [79] and stacked regression [80]. The choice of weighting rule for pools of this kind, and its relation to constrained minimum-variance weighting, is studied separately and is not part of this comparison.

Sample-size, rank and width experiments. Section 6 repeats the isotropic exact kernel, the network, its residual correction and the feature kernel at nested training sizes, at other output ranks and at a wider network; the table of Section 5 is the full-size, rank-64, width-512 cell of that grid.

5 Results

Table 1 compares the model families at 18,884 training rows, output rank 64 and network width 512.

Polynomial and subset-kernel baselines. Cubic ridge regression gives 1.86% mean component error and 0.760% radiance error. The 4,000-point Matérn fit reduces these means to 0.88% and 0.332%. Its albedo error varies substantially across splits: approximately 1.9% on seven splits and 3.3–4.3% on the remaining three.

Full-data and input-scaled kernels. The isotropic full-data kernel reduces mean radiance error to 0.239% and path-radiance error to 0.08%, compared with 0.17% for the subset fit. Its albedo error is 2.3–2.9% across splits: less variable, but larger than the subset fit on the latter’s seven most accurate splits. Input-specific length scales give 0.095% radiance error, 0.63% mean

model	component rel. L^2 [%]				radiance	reflectance
	Y_1	Y_2	Y_3	Y_4	rel. L^2 [%]	med. $ \hat{\rho} - \rho $ [pp]
Cubic ridge	1.16±0.02	2.18±0.04	1.11±0.03	2.98±0.06	0.760±0.011	0.471±0.010
Matérn KRR, 4000-point fit	0.17±0.01	0.61±0.04	0.29±0.02	2.44±0.87	0.332±0.029	0.178±0.030
Matérn KRR, exact on all rows	0.08±0.01	0.42±0.03	0.20±0.00	2.62±0.18	0.239±0.016	0.141±0.018
Matérn KRR, per-input scales	0.08±0.01	0.16±0.02	0.08±0.01	2.18±0.16	0.095±0.009	0.036±0.004
FC-DNN (3×512)	0.48±0.02	0.97±0.04	0.41±0.02	3.02±0.18	0.387±0.012	0.245±0.008
Five-network mean	0.37±0.01	0.67±0.03	0.28±0.01	2.51±0.09	0.251±0.006	0.163±0.003
DNN + residual KRR	0.10±0.00	0.30±0.02	0.16±0.01	2.43±0.08	0.141±0.006	0.089±0.008
Ensemble + residual KRR	0.08±0.00	0.27±0.03	0.13±0.00	2.42±0.17	0.123±0.007	0.080±0.008
Kernel on network features	0.18±0.01	0.17±0.01	0.10±0.01	1.99±0.08	0.098±0.007	0.054±0.004
Per-coordinate selection	0.19±0.01	0.44±0.10	0.13±0.02	2.25±0.18	0.193±0.022	0.123±0.023
Convex stack	0.07±0.00	0.16±0.01	0.09±0.00	1.94±0.06	0.087±0.006	0.048±0.003

Table 1: Test errors under the protocol of Section 3, mean $\pm$ sample standard deviation over ten splits. Component and radiance columns report mean relative L^2 error; radiance is evaluated at $\rho = 0.7$. The last column is median absolute reflectance error in percentage points. The inverse is evaluated across all bands, including deep absorption windows, with the non-finite convention of Section 3. Table 3 reports the associated 95th-percentile errors and RMSE.

component error and 0.036 percentage points of median reflectance error. The fitted scales are consistent with weak azimuth dependence of the flux components. The relative-azimuth length scale is sixteen to thirty-two times larger than the other scales, with median factors sixteen for direct flux and albedo and thirty-two for diffuse flux. The longer azimuth scale occurs on every split for the flux components and nine of ten for albedo. View-zenith scales are two to sixteen times longer, whereas the path-radiance metric is approximately isotropic.

Networks and kernel corrections. The network has 1.22% mean component error and 0.387% radiance error. Averaging five initializations reduces these to 0.96% and 0.251%. Residual kernel regression reduces the single-network radiance error to 0.141% and the ensemble error to 0.123%. The feature kernel has 0.61% mean component error and 0.098% radiance error. Its radiance comparison with the input-scaled kernel is close: the feature kernel is lower on six of ten splits, although its mean is slightly larger. Its median reflectance error is 0.054 points, compared with 0.036 for the input-scaled kernel, and is larger on every split. The separately tuned residual correction substantially improves the network, but has larger radiance error than the input-scaled and feature kernels.

Selection and stacking. Both combiners use the isotropic full-data kernel, the network, its five-member ensemble, the two residual-corrected predictors and the feature kernel. The input-scaled kernel is not in this pool. Coordinatewise validation selection gives 0.193% radiance error, larger than several individual members. The convex stack gives 0.087% radiance error, 0.57% mean component error and 0.048 points of median reflectance error. Its radiance gain is approximately 8% relative to the input-scaled kernel and 11% relative to its best individual member, the feature kernel. On mean component error, the stack improves on the input-scaled kernel at nine splits, with paired difference 0.06 ± 0.05 points, and on the feature kernel at all ten, with difference 0.047 ± 0.006 points.

The residual second moments provide a diagnostic of complementarity. On path radiance, the isotropic kernel and five-network ensemble have residual correlation 0.27 ± 0.09 , below their

candidate / reference	mean gain $\pm$ SD [pp]	gain range [pp]	wins/10
Input-scaled / isotropic KRR	0.1436 ± 0.0122	[0.1196, 0.1585]	10
Residual correction / network	0.2464 ± 0.0090	[0.2365, 0.2615]	10
Narrow stack / input-scaled KRR	0.0083 ± 0.0066	[-0.0054, 0.0185]	9
Wide feature kernel / input-scaled KRR	0.0165 ± 0.0066	[0.0054, 0.0295]	10
Wide / narrow feature kernel	0.0191 ± 0.0079	[0.0096, 0.0352]	10
Wide stack / wide feature kernel	0.0006 ± 0.0026	[-0.0053, 0.0037]	8

Table 2: Paired radiance contrasts on seeds 101–110. Gain is the reference error minus the candidate error, in percentage points (pp). SD is the sample standard deviation of the ten paired differences; the range and win count describe those same ten differences. They are not confidence intervals or tests of a population-level superiority claim.

error ratio of approximately 0.34 at every split. For a quadratic error metric, this is the classical condition under which a convex combination can improve on the better member [78, 11]. The stack instead minimizes a mean relative norm, and its reflectance objective is nonlinear; the diagnostic does not guarantee improvement in those metrics. The two residual-corrected predictors have much larger correlations, approximately 0.88 on path radiance and 0.99 on albedo. Coordinatewise selection also differs from stacking in its objective and number of independently chosen coefficients.

Spherical albedo. Albedo has the largest relative error for every family. Its small short-wave-infrared values make relative errors sensitive to absolute prediction errors in that region. The width-512 stack achieves 1.94% albedo error. The learning curves below show little improvement with sample size on this component; the experiments do not isolate whether the plateau arises from approximation, optimization or the relative-error weighting.

5.1 Paired comparisons and the radiance–reflectance trade-off

Paired differences compare candidate and reference errors on the same split. Table 2 reports those differences for the main and wide pipelines. Positive gains favor the candidate. The wide feature kernel improves on the input-scaled kernel at all ten splits, by 0.0165 ± 0.0066 radiance percentage points. By contrast, the wide stack’s improvement over its own feature kernel is only 0.0006 ± 0.0026 points, with eight wins and two losses. The mean improvement is small relative to the variability of the paired differences. These are descriptive paired summaries: repeated splits of one finite table, overlapping test rows and adaptive study development do not supply ten independent external validation experiments.

Radiance and reflectance also order the heads differently. Table 3 reports the all-band reflectance tails alongside the medians. Each entry averages a within-split statistic; the mean of ten 95th percentiles is not a pooled 95th percentile. For the width-512 pipeline, the convex stack lowers radiance error relative to the feature kernel, but raises the average 95th-percentile reflectance error from 7.75 to 29.44 percentage points. The tail is worse at all ten splits, by factors between 2.63 and 4.94. The width-2000 comparison repeats the reversal: the feature kernel has 7.57 points at the 95th percentile and the stack 22.71, with the stack worse at all ten splits by factors between 2.36 and 3.79. Thus the near-tie at 0.079% versus 0.078% radiance error conceals a substantial difference in the inverse metric.

model	radiance [%]	median [pp]	95th percentile [pp]	RMSE [pp]
Cubic ridge	0.760±0.011	0.471±0.010	70.68±0.02	4.18e+03
Isotropic KRR	0.239±0.016	0.141±0.018	69.53±0.42	4.18e+03
Input-scaled KRR	0.095±0.009	0.036±0.004	31.16±6.21	1.34e+04
Narrow feature kernel	0.098±0.007	0.054±0.004	7.75±1.13	1.78e+03
Narrow convex stack	0.087±0.006	0.048±0.003	29.44±5.89	1.25e+03
Wide feature kernel	0.079±0.009	0.045±0.006	7.57±0.95	1.40e+03
Wide convex stack	0.078±0.011	0.043±0.005	22.71±4.42	5.81e+03

Table 3: All-band metrics from the complete ten-seed main and wide campaign records. Radiance is mean relative L^2 error. Reflectance median, 95th-percentile absolute error and RMSE are in percentage points at $\rho = 0.7$. The first three columns show mean $\pm$ sample SD across splits; RMSE shows its mean only, with its spread retained in the machine-readable summary. The inverse is neither physically clipped nor conditioned on an application passband. All inverses in these records are finite, so their zero-fill branch does not explain the large errors. Large finite values near denominator zeros dominate RMSE. These are diagnostic all-band errors, not operational retrieval accuracy.

The input-scaled kernel illustrates the same distinction: its median error is only 0.0365 points but its mean 95th percentile is 31.16 points. Conversely, the feature kernels have slightly larger medians and appreciably smaller tails. Model rankings depend on the metric and spectral domain. The identities of Section 3.1 relate inverse error to the predicted denominator as well as forward error. This explains why optimizing componentwise or radiance error alone need not control reflectance tails; it does not identify the component-error correlations responsible for a particular observed tail.

The tables are generated from the twenty individual records. The generator checks split configurations, agreement of recorded array digests, and equality of common deterministic baselines. These checks establish consistency of the summaries; reproducing model training requires the raw arrays.

6 Learning curves, output rank and network width

The following experiments vary training size, output rank and network width, using the same validation-based selection. Learning curves include cubic ridge, isotropic and input-scaled kernels, the network, residual correction and the feature kernel. Ensemble comparisons use the full training size. The number of available repetitions is given with each experiment.

Learning curves. Each rung uses the first n rows of the seed’s training block, $n = 500, 1,000, 2,000, 4,000, 8,000$ and the full 18,884, with the validation and test blocks unchanged, so the rungs of one seed are nested and the full block is the top rung (Table 4, Figure 3). Cubic error changes little: from 0.85% to 0.76% over a factor of 38 in data, with most of its decrease achieved by a thousand rows. On the first three components every other family decays at close to $n^{-1/2}$ (Table 5): the least-squares slopes of log error against log n over the six rungs lie between -0.38 and -0.61 for the kernel, the network, the corrected network and the feature kernel on Y_1, Y_2 and Y_3 , while on the albedo every slope lies between -0.09 and $+0.01$: the albedo shows no

training rows	seeds	Cubic ridge	KRR, isotropic	KRR, per-input	FC-DNN	DNN + residual KRR	kernel on features
500	10	0.846±0.031	0.639±0.022	0.415±0.024	1.961±0.057	0.815±0.035	0.456±0.020
1,000	10	0.802±0.024	0.519±0.034	0.260±0.066	1.862±0.015	0.590±0.027	0.316±0.040
2,000	10	0.780±0.016	0.424±0.051	0.217±0.052	1.232±0.034	0.376±0.041	0.240±0.049
4,000	10	0.769±0.010	0.328±0.028	0.154±0.037	0.830±0.019	0.256±0.021	0.164±0.033
8,000	10	0.763±0.010	0.279±0.020	0.119±0.027	0.584±0.020	0.190±0.014	0.121±0.017
18,884	10	0.760±0.011	0.239±0.016	0.095±0.009	0.387±0.012	0.141±0.006	0.098±0.007

Table 4: Learning curves: radiance relative L^2 error [%] against the number of training rows, mean $\pm$ standard deviation over the ten splits reported in the archived aggregate. At $n = 8,000$, one individual record per kernel campaign is absent from the public archive; the table reports the ten-run aggregate. Each rung uses the first n rows of the seed’s training block; the validation and test blocks are those of Table 1, so the top rung is that table. The per-input kernel repeats its coordinate search at every rung. Component columns are in Figure 3 and in the per-seed records.




Figure 3: Learning curves on log–log axes for the radiance and for three of the four components, mean and standard deviation over ten splits. Cubic error changes little over the sample-size range. The kernel, network, residual-corrected network and feature kernel decrease approximately as $n^{-1/2}$ on the first three components, with fitted slopes in Table 5. Albedo error shows no sustained decrease; unusually large errors on individual splits increase its spread at 1,000 and 2,000 rows.

sustained decrease over this budget range. This plateau does not distinguish approximation bias, optimization, relative-error weighting or limitations of the training data.

These are slopes of six mean errors over a finite sample-size range, not established asymptotic rates. Standard Sobolev or kernel regression rate formulas require sampling, source and regularization assumptions [45, 47]; neither those assumptions nor the smoothness norm of the fitted network residual is verified by this experiment. The isotropic kernel’s decrease is slower at the largest sample size. Its slope over the first five rungs is -0.59 , -0.44 and -0.47 on the smooth components, and its last step, from 8,000 to 18,884 rows, is -0.38 , -0.19 and -0.17 , while the network’s last step is as steep as its average (-0.59 , -0.60 and -0.55 ; Table 5). Part of this is the tuning protocol: the kernel’s smoothness, scale and nugget are chosen on a 6,000-row subsample at every rung, so above 6,000 rows the hyperparameters are those of a smaller problem, and Section 6.1 measures what tuning with every row in the solve recovers: a substantial part of the measured change in slope. The consequence for the comparison is that the ratio of the network’s radiance error to the kernel’s falls from 3.1 at 500 rows to 1.6 at the full block, or to 1.8 with the kernel tuned on every row. The kernel’s margin over the network on this table is a margin at this training size; whether it would survive ten times the data depends on where

family	slope over the six rungs					slope of the last step				
	Y_1	Y_2	Y_3	Y_4	radiance	Y_1	Y_2	Y_3	Y_4	radiance
Cubic ridge	-0.03	-0.03	-0.03	-0.02	-0.03	-0.00	-0.01	-0.00	-0.08	-0.01
KRR, isotropic	-0.55	-0.38	-0.40	-0.01	-0.28	-0.38	-0.19	-0.17	0.10	-0.18
KRR, per-input	-0.45	-0.43	-0.43	0.01	-0.40	-0.32	-0.26	-0.29	0.12	-0.26
FC-DNN	-0.50	-0.49	-0.57	-0.09	-0.48	-0.59	-0.60	-0.55	0.07	-0.48
DNN + residual KRR	-0.52	-0.53	-0.61	-0.04	-0.50	-0.39	-0.33	-0.39	0.12	-0.35
Kernel on features	-0.45	-0.44	-0.46	-0.03	-0.43	-0.28	-0.22	-0.31	0.08	-0.25

Table 5: Learning-curve exponents: the least-squares slope of log error against log n over the six rungs of Table 4, and the slope of the last step alone, from 8,000 to 18,884 rows, per family and component.

PCA rank	variance outside the basis	KRR, isotropic	FC-DNN	DNN + residual KRR	kernel on features
16	7.5×10^{-4}	0.270 ± 0.015	0.390 ± 0.015	0.175 ± 0.007	0.148 ± 0.009
32	1.5×10^{-5}	0.243 ± 0.018	0.390 ± 0.015	0.139 ± 0.006	0.102 ± 0.009
64	8.2×10^{-8}	0.242 ± 0.018	0.390 ± 0.015	0.143 ± 0.006	0.100 ± 0.009
128	5.4×10^{-9}	0.242 ± 0.018	0.390 ± 0.015	0.150 ± 0.005	0.100 ± 0.009
285	0	0.242 ± 0.018	0.390 ± 0.015	0.172 ± 0.004	0.100 ± 0.009

Table 6: Output-rank ablation: radiance relative L^2 error [%] when the kernels and corrections regress the leading r principal coefficients, mean $\pm$ standard deviation over five splits. The second column is the fraction of the standardized variance the basis leaves out for the slowest-decaying component, the spherical albedo; at rank 285 it is zero. The network regresses the 285 bands directly and is the same predictor in every row.

the kernel’s curve settles, which the bend leaves open.

Relative to the isotropic kernel, the per-input metric changes the measured finite-range slope as well as the error level. Its coordinate search, repeated at every rung, gives a slope of -0.399 on the radiance against the isotropic kernel’s -0.279 , and the ratio between them falls from 0.65 at 500 rows to 0.40 at the full block, with one step against that direction (0.50 at 1,000 rows, 0.51 at 2,000) and a fall at every other. The relative error reduction increases over this sample-size range. The selected metric, smoothness and nugget all change with sample size, so these measurements do not isolate a change of asymptotic dimension or establish a faster universal rate. Against the kernel on network features the margin closes as the data grows. The per-input metric has the lower mean at all six rungs, but the split-by-split count runs 10, 9, 10, 7, 5 and 6 wins of ten from 500 rows to the full block, and the gap itself narrows from 0.415 against 0.456 at 500 rows to 0.095 against 0.098 at the top; the feature kernel’s own slope, -0.435 , is the steeper of the two. The input-scaled kernel has a larger advantage at smaller sample sizes, while the feature kernel approaches it as the training set grows. On the albedo the per-input slope is $+0.01$: with no sustained improvement over this range.

The corrected network is below both of its parents on the radiance at every one of the ten splits from 2,000 rows on, and at no split at 500 or 1,000 rows, where its error sits between the network’s and the kernel’s. On Y_2 and Y_3 it is below both at every split at the two largest sizes; on Y_1 it never is, the isotropic kernel alone being the better head at every size (0.084% against 0.099% at the full block). Its exponent is close to the network’s and its error level is closer to the kernel’s in these runs. Section 7 gives the exact residual decomposition, while distinguishing it from a proof of these empirical slopes or of the component-specific exception.

network	head	component rel. L^2 [%]				radiance	reflectance
		Y_1	Y_2	Y_3	Y_4	rel. L^2 [%]	med. $ \hat{\rho} - \rho $ [pp]
3×512	network	0.48 ± 0.02	0.97 ± 0.04	0.41 ± 0.02	3.02 ± 0.18	0.387 ± 0.012	0.245 ± 0.008
	network + residual KRR	0.10 ± 0.00	0.30 ± 0.02	0.16 ± 0.01	2.43 ± 0.08	0.141 ± 0.006	0.089 ± 0.008
	kernel on its features	0.18 ± 0.01	0.17 ± 0.01	0.10 ± 0.01	1.99 ± 0.08	0.098 ± 0.007	0.054 ± 0.004
	per-coordinate selection	0.19 ± 0.01	0.44 ± 0.10	0.13 ± 0.02	2.25 ± 0.18	0.193 ± 0.022	0.123 ± 0.023
	convex stack	0.07 ± 0.00	0.16 ± 0.01	0.09 ± 0.00	1.94 ± 0.06	0.087 ± 0.006	0.048 ± 0.003
3×2000	network	0.93 ± 0.20	2.11 ± 0.49	0.76 ± 0.19	3.27 ± 0.19	0.748 ± 0.116	0.508 ± 0.110
	network + residual KRR	0.12 ± 0.01	0.35 ± 0.03	0.19 ± 0.03	2.35 ± 0.13	0.160 ± 0.016	0.106 ± 0.010
	kernel on its features	0.13 ± 0.02	0.13 ± 0.02	0.08 ± 0.01	2.10 ± 0.34	0.079 ± 0.009	0.045 ± 0.006
	per-coordinate selection	0.14 ± 0.02	0.30 ± 0.26	0.09 ± 0.01	2.55 ± 0.26	0.187 ± 0.057	0.099 ± 0.033
	convex stack	0.08 ± 0.01	0.13 ± 0.02	0.08 ± 0.00	2.00 ± 0.24	0.078 ± 0.011	0.043 ± 0.005

Table 7: The pipeline of Section 4 run with a 3×2000 network in place of the 3×512 one, at the same ten splits: the network (one member, up to 500 epochs under the same optimizer, learning rate and early stopping), its residual correction, the exact kernel on its last-layer features, and the per-coordinate selection and convex stack over the wide pipeline’s heads (the isotropic kernel, the network, its correction and its feature kernel), beside the corresponding rows of Table 1; mean $\pm$ standard deviation over the ten splits.

Output rank. The kernels and corrections regress the leading r principal coefficients of the standardized component; the network regresses the 285 bands and is the same predictor at every rank. Table 6 refits at $r = 16, 32, 128$ and 285 at five of the seeds beside the rank-64 rows of those seeds. The kernel’s measured error changes little above rank 32: its radiance error reads 0.243, 0.242, 0.242 and 0.242 at 32, 64, 128 and 285, and 0.270 at 16, where the truncation leaves 7.5×10^{-4} of the albedo’s variance outside the basis. The kernel on features behaves the same way. The residual correction does not: it is best at 32 (0.139), a little worse at 64 (0.143), and worse again at 128 and 285 (0.150 and 0.172). The two regressions fit different targets: the direct kernel fits the component, whereas the correction fits the projected network residual. Increasing rank therefore exposes the correction to additional network-error directions, even when the target variance in those directions is small. The lowest measured correction errors occur at ranks 32–64. The projection identity in Section 7 separates these two error terms.

A wider network. A network of three hidden layers of 2,000 units, trained for up to 500 epochs under the same optimizer, learning rate and early stopping as the 512-unit network, was run at all ten seeds with the rest of the pipeline built on it (Table 7). It is a worse emulator at every seed, 0.75% against 0.39% on the radiance, and its early stopping fired between 42 and 210 epochs of the 500 allowed: at this learning rate and batch size this particular wider-network training configuration gives worse standalone predictions. Width and epoch budget change together; these results do not rule out a better-optimized wide network. The exact kernel on its last-layer features reaches 0.079% on the radiance against 0.098% on the 512-unit features, below the per-input kernel’s 0.095% at all ten seeds and below the convex stack of Table 1 at nine of them, and on the two flux components, at 0.13% and 0.08%, it is at or below every row of that table. The convex stack over the wide pipeline’s heads, which puts 0.88 to 0.96 of its weight on that feature kernel, has the lowest mean radiance error in the ten-split comparison at 0.078% on the radiance and 0.043 reflectance points, below the stack of Table 1 at nine of the ten seeds; the per-coordinate selection over the same heads does not transfer here either. The residual correction of the wider network is a little worse than that of the narrower one,

kernel	component rel. L^2 [%]				radiance	reflectance
	Y_1	Y_2	Y_3	Y_4	rel. L^2 [%]	med. $ \hat{\rho} - \rho $ [%]
isotropic, tuned on 6,000 rows	0.08±0.01	0.42±0.03	0.20±0.00	2.62±0.18	0.239±0.016	0.141±0.018
isotropic, tuned on every row	0.08±0.00	0.36±0.02	0.17±0.01	2.58±0.14	0.210±0.008	0.107±0.006
per-input search, tuned on 6,000 rows	0.08±0.01	0.16±0.02	0.08±0.01	2.18±0.16	0.095±0.009	0.036±0.004
per-input search, re-tuned on every row	0.08±0.00	0.16±0.02	0.08±0.01	2.09±0.12	0.096±0.009	0.036±0.004
per-input from the cubic’s gradients, $p = 1$	0.34±0.01	0.19±0.02	0.14±0.01	2.17±0.14	0.159±0.007	0.078±0.004
per-input from the cubic’s gradients, $p = 1/2$	0.11±0.00	0.19±0.02	0.09±0.01	2.27±0.15	0.125±0.008	0.055±0.004
per-input from the cubic’s gradients, $p = 1/4$	0.08±0.00	0.25±0.02	0.14±0.01	2.38±0.20	0.157±0.015	0.077±0.013

Table 8: Where the kernel’s hyperparameters are chosen, and how its metric is obtained, at the full training block and the same ten splits. The first pair tunes smoothness, scale and nugget on a 6,000-row subsample and on every training row; the last three rows replace the coordinate search by a metric read off the cubic surrogate’s gradients, at three exponents. Selecting the exponent on validation, per seed, chooses $p = 1/2$ at all ten, so that row would duplicate the $p = 1/2$ row and is not printed; the per-seed choices are in the stored numbers. Mean $\pm$ standard deviation over the ten splits.

although that ordering is not guaranteed merely because its base network is worse (Section 7). A 3×1000 rung at the same ten seeds sits between the two: the network at 0.72%, the kernel on its features at 0.088% against 0.098% at width 512 and 0.079% at width 2,000. From 512 to 2,000 the network’s own error rises with width while the error of a kernel on its features falls: the fitted features improve kernel prediction despite the larger standalone network error. At 4,000 units, run at two seeds, the network is worse again at 0.95% and the kernel on its features turns back up to 0.097%, so the gain in the representation does not continue past 2,000.

A three-member wide-network experiment, run at three seeds, supplies the exact kernel with the concatenated last-layer features of all three members at once, 6,000 coordinates rather than 2,000. It reads 0.074% on the radiance against 0.083% for the kernel on a single member’s features on the same splits, lower at all three seeds, and 0.043 against 0.051 reflectance points, again at all three. That puts one exact solve level on the radiance with the convex stack of six heads fitted on those same splits (0.074% against 0.074% on the mean, though the stack is the lower at two of the three seeds and on the reflectance median at all three). The six-predictor stack and the single kernel on concatenated features have similar mean radiance errors on these three splits. The comparison does not isolate all optimization and kernel-selection effects. Three seeds provide exploratory evidence rather than a ten-split replication, and the concatenated-feature construction was evaluated only at width 2,000.

6.1 Full-data tuning and gradient-based metrics

We compare hyperparameter selection on a 6,000-row subsample with selection on all training rows, and compare coordinate search with a metric derived from a cubic surrogate’s gradients (Table 8). The reported aggregates contain ten splits; the per-split archive for full-data tuning contains five of them.

Tuning with every row in the solve. The kernel families of Section 4 choose smoothness, scale and nugget by the validation error of a solve on a 6,000-row subsample, and refit the winner on all 18,884 rows. Doing the tuning solve on every row instead lowers the isotropic kernel’s radiance error from 0.239% to 0.210% and its reflectance median from 0.141 to 0.107 points, better at nine of the ten splits, the ratio averaging 0.882 and never above 1.001. The

selected scales differ: over the forty component-seed cells the full block picks a shorter length scale than the subsample in 26, the same in 12, and a longer one in 2. In this experiment, subsample selection usually chooses a longer length scale than full-data selection, consistent with the smaller tuning fit favoring greater smoothing.

Below 6,000 rows the tuning subsample *is* the whole training block and the rungs are already full-block tuned; at 8,000 rows it is three quarters of it; only the top rung is a third. Applying the measured ratio to the top rung of Table 4 takes the kernel’s last step on the radiance from -0.183 to -0.334 , against a slope of -0.305 over the five lower rungs: the top-step slope becomes closer to that of the lower rungs. One retuned endpoint does not establish a power law over or beyond the observed range. On the two flux components the same correction takes the last step from -0.187 and -0.171 to -0.352 and -0.358 against averages of -0.437 and -0.468 , reducing, but not eliminating, the difference between the last-step and lower-rung slopes. The consequence for the comparison of Section 5 is that the kernel’s margin over the network at the full block is wider than Table 4 shows: 1.8 times rather than 1.6, under this retuned endpoint. No conclusion about unobserved sample sizes follows.

Full-data tuning has little effect on the input-scaled kernel: 0.096% against 0.095%, better at four of the ten splits, ratio 1.005. Its coordinate search already selects a scale of once the median distance at the metric it finds, which is what the full-block tune selects too. The input-scaled kernel is therefore less sensitive to the tuning sample size in these experiments.

Gradient-based metric. The coordinate search costs a sweep over the six inputs at five multipliers, twice, per component and per seed. The alternative sets the metric before any kernel is fitted: the multiplier of coordinate j is the root-mean-square gradient of the fitted cubic surrogate along j over the training rows, raised to a power p , normalized to geometric mean one and clipped to the search’s own range $[1/16, 4]$, so weakly varying coordinates receive longer length scales. The exponent p controls the strength of the rescaling. It costs one cubic fit and six finite differences.

At $p = 1$, radiance error is 0.159%, below the isotropic kernel at every one of the ten splits by a factor 0.669 on average and above the searched metric at every one of them by a factor 1.680. The dependence on p is nonmonotone: $p = 1/2$ gives 0.125% and $p = 1/4$ gives 0.157%, so the best exponent is interior, and choosing it the way every other hyperparameter in this study is chosen — the lowest mean validation error over the four components, per seed, with the test block untouched — selects $p = 1/2$ at all ten seeds. At that exponent the metric closes 79% of the distance from the isotropic kernel to the searched one on the radiance and 82% on the reflectance median, for one cubic fit, six finite differences and a choice among three numbers.

The componentwise errors favor different exponents: Y_2 and Y_3 are best at $1/2$, the albedo at 1 (2.18% against 2.27%), and the path radiance at $1/4$. On that component the isotropic kernel already reads 0.084% and the searched metric 0.085%: coordinate search gives little improvement. The proportional metric nonetheless reads 0.34% there, approximately four times the isotropic error. A possible explanation is that the surrogate’s gradient along a coordinate measures how fast the map moves with it and not how much predictive information it carries, so gradient magnitude alone need not provide an appropriate kernel metric. Softening the exponent to $1/4$ takes Y_1 back to 0.085%, exactly the searched metric’s value, at the cost of Y_2 and Y_3 . The validated gradient metric recovers much of the observed radiance gain with fewer metric searches. Coordinate search additionally allows independent componentwise adjustments, whereas the gradient construction selects one exponent for all four components.

7 Residual correction and error alignment

Fix the training design, neural mean and selected kernel hyperparameters. The following identities hold without sampling assumptions. We first work with one output coordinate; the vector-valued statements follow coordinatewise.

7.1 A frozen regression operator

Fix a design $X_n = (x_1, \dots, x_n)$, a positive-semidefinite kernel with Gram matrix K , and $\lambda > 0$. Define

$$(Sg)(x) = k(x, X_n)(K + \lambda I)^{-1}g(X_n), \quad R = I - S. \quad (8)$$

Here λ is the diagonal regularizer in the linear system. The campaign code uses $\lambda = n\gamma$, where γ is its reported nugget; the distinction matters when the training size changes. Although a selected kernel can depend on the targets, S is linear in the response only *after that selection is frozen*.

Proposition 2 (Frozen residual correction). *Let f be the target and h any fitted mean, and set $\hat{f}_c = h + S(f - h)$. Then:*

$$f - \hat{f}_c = R(f - h), \quad (9)$$

$$(Rg)(X_n) = \lambda(K + \lambda I)^{-1}g(X_n). \quad (10)$$

If $Ku_j = \mu_j u_j$, the design residual in direction u_j is multiplied by $\lambda/(\mu_j + \lambda) \in (0, 1]$. The factor is one when $\mu_j = 0$ and strictly below one when $\mu_j > 0$; this is shrinkage, not a hard spectral projection.

Suppose, in addition, that $\mathcal{S}$ is a normed vector space containing f and h and that the following bound holds uniformly on that space for the realized operator:

$$\|Rg\|_{L^2(P)} \leq C_n \|g\|_{\mathcal{S}} \quad (g \in \mathcal{S}).$$

Then

$$\|f - \hat{f}_c\|_{L^2(P)} \leq C_n \|f - h\|_{\mathcal{S}}. \quad (11)$$

Consequently, $\|f - h\|_{\mathcal{S}} < \|f\|_{\mathcal{S}}$ makes this displayed upper bound smaller than $C_n \|f\|_{\mathcal{S}}$ when $C_n > 0$. It is neither a necessary nor a sufficient condition for the corrected predictor's actual error to be smaller than that of Sf .

Proof. Subtract $h + S(f - h)$ from f to obtain (9). On the design, $I - K(K + \lambda I)^{-1} = \lambda(K + \lambda I)^{-1}$, which proves (10) and its spectral form. Apply the assumed uniform inequality to $g = f - h$ for (11). The examples below distinguish this bound from actual error comparisons. $\square$

Why the bounds do not order the errors. On a two-point design take $K = \text{diag}(1, 99)$ and $\lambda = 1$, so $R = \text{diag}(1/2, 1/100)$ and $C_n = 1/2$ is its Euclidean operator norm. With $f = (1, 0)$ and $f - h = (0, 2)$, the residual norm is twice the target norm, yet $\|R(f - h)\| = 0.02 < 0.50 = \|Rf\|$. Conversely, with $f = (0, 1)$ and $f - h = (1/2, 0)$, the residual norm is smaller, but $\|R(f - h)\| = 0.25 > 0.01 = \|Rf\|$. These examples also show why source-norm accuracy alone does not order corrected prediction errors.

7.2 An exact criterion for improvement

The relevant comparison can be made exactly in any Hilbert error metric. This includes a finite weighted squared error and $L^2(P)$, but not directly the mean of per-sample relative norms or the nonlinear radiance and reflectance metrics in Table 1.

Proposition 3 (Error alignment). *For the same frozen operator S , put $a = Rf$ and $b = Rh$. Then*

$$\|f - \hat{f}_c\|^2 - \|f - Sf\|^2 = \|b\|^2 - 2\langle a, b \rangle. \quad (12)$$

Thus full correction improves on Sf exactly when $2\langle a, b \rangle > \|b\|^2$. More generally, the interpolation $p_t = Sf + tRh$, $0 \leq t \leq 1$, has minimum squared error at

$$t_* = \min\{1, \max\{0, \langle a, b \rangle / \|b\|^2\}\} \quad \text{if } b \neq 0; \quad (13)$$

if $b = 0$, every t is equivalent. If the direct fit uses S_0 and the correction uses S_1 , (12) instead holds with $a = (I - S_0)f$ and $b = (S_1 - S_0)f + (I - S_1)h$.

Proof. For a common operator the corrected error is $a - b$, so expansion of its squared norm gives (12). The error of p_t is $a - tb$; minimize the resulting quadratic on $[0, 1]$ to obtain (13). For distinct operators, $f - h - S_1(f - h) = (I - S_0)f - [(S_1 - S_0)f + (I - S_1)h]$, after which the same expansion applies. $\square$

The criterion depends on the unknown target error. Replacing its inner products by validation errors gives a validation-optimal interpolation, not a guarantee on new data. The reported experiments use separately tuned direct and residual kernels, so their comparison must use the final, two-operator form with their actual selected operators. Neither the criterion nor the training-design contraction establishes that correction will beat both parents at every sample size.

7.3 What the principal-component implementation computes

The campaign trains a network on all standardized bands but projects its predictions into the retained principal-component basis before correction. Thus its corrected predictor is based on the projected network mean. To state the implementation precisely, let f and h now denote the target and network output after the common fitted standardization and PCA centering. Let Π be the orthogonal projection onto the retained spectral coordinates. In these centered coordinates, the implemented prediction is

$$\hat{f}_{c,\Pi} = \Pi h + S\Pi(f - h).$$

The output mean and diagonal band scaling are restored afterwards. Since the same scalar kernel operator acts on each retained output coordinate,

$$f - \hat{f}_{c,\Pi} = (I - \Pi)f + (I - S)\Pi(f - h). \quad (14)$$

The terms on the right are orthogonal in the standardized spectral Euclidean norm at each query, hence their squared norms add there. They need not be orthogonal in physical units after unequal band rescaling. This decomposition separates truncation error from residual regression and explains why output rank is a genuine part of the model specification. It does not by itself prove the observed optimum near rank 32. Comparisons with the unprojected network include the effect of projection as well as residual fitting.

Conditional rate bound. Kernel learning-rate results impose source, spectral, sampling and regularization conditions [45, 46, 47]. If a uniform residual bound holds with $C_n = O(n^{-\alpha})$ and $\|f - h_n\|_{\mathcal{S}} = O(n^{-\beta})$, then (11) gives $\|f - \hat{f}_c\|_{L^2(P)} = O(n^{-(\alpha+\beta)})$. The source-norm bound and the uniformity required for a data-dependent residual are additional hypotheses. They do not follow from empirical network loss. The experiments measure finite-range error slopes, not these source norms or an asymptotic rate.

The code checks the residual identities, regularizer convention, semidefinite boundary, alignment formula, counterexamples and projected implementation on explicit and randomized finite designs.

8 Retrieval-error transfer and the sensitivity of learned corrections

8.1 Sensitivity of the correction

For a common frozen operator $R = I - S_\lambda$ and the norm of Section 7, the exact criterion $2\langle Rf, Rh \rangle > \|Rh\|^2$ is necessary and sufficient for strict improvement of the corrected predictor over the standalone kernel, since $\|Rf - Rh\|^2 < \|Rf\|^2$ is that inequality rearranged; the comparison of smoothness norms is neither necessary nor sufficient. Write the design values of the targets, of the fitted mean and of the corrected predictor as y , h_X and $\hat{y}$, with $\hat{y} = h_X + S_\lambda(y - h_X)$ and $S_\lambda = K(K + \lambda I)^{-1}$; in the nugget convention of the experiments, $\lambda = n\gamma$. The mean is trained on the same labels, so $h_X = h_X(y)$.

Proposition 4. *1. If h_X is differentiable in y , the Jacobian of the trained pipeline is $J_{\hat{y}} = S_\lambda + (I - S_\lambda)J_h$. In particular the correction stage alone does not bound the sensitivity of the pipeline: with $K = I$, $\lambda = 1$ and $h_X(y) = 2y$ one has $\hat{y} = \frac{3}{2}y$ although $\|S_\lambda\|_{\text{op}} = \frac{1}{2}$; and if the mean interpolates the training labels, $h_X(y) = y$, then $\hat{y} = y$ and the correction changes nothing on the design.*

2. With the mean and the kernel frozen, $y \mapsto \hat{y}$ is affine with linear part S_λ , its effective degrees of freedom are $\text{tr } S_\lambda = \sum_j \mu_j / (\mu_j + \lambda)$, and $\|S_\lambda\|_{\text{op}} = \max_j \mu_j / (\mu_j + \lambda) < 1$.

3. Suppose the mean, the kernel and the ridge are fixed independently of the correction noise, conditional on the design, and the correction is fitted on an independent sample with $y = f_X + \varepsilon$, $\mathbb{E}\varepsilon = 0$, $\mathbb{E}[\varepsilon\varepsilon^\top] = \sigma^2 I$. With $r = f_X - h_X$ and r_j its coordinates in the eigenbasis of K ,

$$\frac{1}{n} \mathbb{E}_\varepsilon \|\hat{y} - f_X\|^2 = \frac{1}{n} \sum_j \left(\frac{\lambda}{\mu_j + \lambda} \right)^2 r_j^2 + \frac{\sigma^2}{n} \sum_j \left(\frac{\mu_j}{\mu_j + \lambda} \right)^2.$$

4. At a new input x , with $a_\lambda(x) = (K + \lambda I)^{-1}k(X_n, x)$, the noise variance of the correction is $\sigma^2 \|a_\lambda(x)\|_2^2$.

Proof. (i) is the chain rule applied to $\hat{y} = S_\lambda y + (I - S_\lambda)h_X(y)$, and the two examples are direct substitutions. (ii) With the mean and kernel frozen the map is affine with linear part S_λ ; the trace is the sum of the eigenvalues $\mu_j / (\mu_j + \lambda)$, and the operator norm of a symmetric positive semidefinite matrix is its largest eigenvalue. (iii) $\hat{y} - f_X = -(I - S_\lambda)r + S_\lambda\varepsilon$; the cross term has zero mean, and both matrices are diagonal in the eigenbasis of K . (iv) The correction at x is $a_\lambda(x)^\top(y - h_X)$, linear in ε with coefficient vector $a_\lambda(x)$. $\square$

Item (iii) is what the ridge-path experiments report: the residual's spectral energy $\sum_j r_j^2$ by eigendirection, and the two sums as functions of λ , together with the out-of-sample factor of item (iv). A ridge chosen from fits that depend on the correction noise is outside the hypothesis of (iii).

8.2 A finite-class selection guarantee

Every predictor below is chosen on held-out rows rather than by an information criterion [81, 82, 83], so what has to be bounded is the price of that choice.

Proposition 5. *Let $p_1, \dots, p_M$ be predictors frozen before an independent selection sample of m i.i.d. units, let the loss take values in $[0, B]$, and let $\hat{p}$ minimise the empirical risk on that sample. Then with probability at least $1 - \delta$,*

$$R(\hat{p}) \leq \min_{1 \leq j \leq M} R(p_j) + 2B \sqrt{\frac{\log(2M/\delta)}{2m}}.$$

Proof. Hoeffding's inequality [107] for each candidate and a union bound over the M candidates. $\square$

The squared error of the constrained retrieval of Section 8.3 is bounded by R^2 , so the guarantee covers it. It covers the damped correction $p_{\alpha, \lambda}(x) = h(x) + \alpha k(x, X_n)(K + \lambda I)^{-1}(y - h_X)$ at finitely many $\alpha \in [0, 1]$, with $\alpha = 0$ and the standalone kernel among the candidates. It does not cover a stack whose weights are optimised; the stacks of Section 8.5 are therefore fitted on development rows, frozen as a small collection, and selected on an independent state-level sample, with bands and prescribed reflectances averaged within each state, so that m is the number of independent states.

8.3 Transfer of component error into retrieval error

Section 3.1 controls the inverse where the denominator stays away from zero. The theorem below describes a whole evaluation distribution, with entries of vanishing transmission included, with the constant one, and shows that the resulting loss of rate is not an artefact of the argument. Let P be the evaluation distribution over entries (a state, a band, and the surface reflectance when it is varied), in fixed physical units; expectations are conditional on the fitted emulator. Assume

$$0 \leq \rho \leq R \leq 1, \quad t > 0, \quad 0 \leq s \leq S, \quad RS < 1. \quad (15)$$

The emulator predicts $(\hat{a}, \hat{t}, \hat{s})$. The *constrained retrieval* replaces $\hat{t}$ by $\max\{\hat{t}, 0\}$ and $\hat{s}$ by its projection onto $[0, S]$, neither of which increases $|e_t|$ or $|e_s|$, and sets

$$\bar{\rho} = \Pi_{[0, R]} \left(\frac{L - \hat{a}}{\hat{t} + \hat{s}(L - \hat{a})} \right) \quad (16)$$

when the denominator is positive, and otherwise returns a prescribed value in $[0, R]$ and records a failure. It is an estimator in its own right and is reported beside the unconstrained inverse and its failure counts, never in place of them. The *retrieval-weighted component error*, computed after the projections, is

$$e_R = |e_a| + R|e_t| + R^2 t |e_s|, \quad \varepsilon_R^2 = \mathbb{E}_P[e_R^2], \quad (17)$$

which is at most $e = |e_a| + |e_t| + t|e_s|$ since $R \leq 1$; ε_R is not the mean relative radiance error of the tables. Put $F_t(\tau) = \mathbb{P}(t \leq \tau)$.

Theorem 6 (error transfer with constant one). *Under (15), for the constrained retrieval (16):*

1. $|\bar{\rho} - \rho| \leq \min\{R, e_R/t\} \leq \min\{R, e/t\}$ at every entry, and no universal constant below one is possible.
2. For every $\tau > 0$,

$$\mathbb{E}_P |\bar{\rho} - \rho|^2 \leq \mathbb{E}_P \min\left\{R^2, \frac{e_R^2}{t^2}\right\} \leq R^2 F_t(\tau) + \mathbb{E}_P \left[\frac{e_R^2}{t^2} \mathbf{1}\{t > \tau\} \right] \leq R^2 F_t(\tau) + \frac{\varepsilon_R^2}{\tau^2}.$$

3. If $F_t(u) \leq Au^\beta$ for $0 < u \leq u_0$ with some $\beta > 0$, then for $\varepsilon_R \leq u_0^{(\beta+2)/2}$, $\mathbb{E}_P |\bar{\rho} - \rho|^2 \leq (R^2 A + 1) \varepsilon_R^{2\beta/(\beta+2)}$; a component-error bound $\varepsilon_{R,n} = O(n^{-r})$ yields $\|\bar{\rho}_n - \rho\|_{L^2(P)} = O(n^{-r\beta/(\beta+2)})$.
4. If $t \geq t_0 > 0$ P -almost surely, then $\|\bar{\rho} - \rho\|_{L^2(P)} \leq \varepsilon_R/t_0$.

Proof. Write $q = 1 - \rho s$, $u = L - a = \rho t/q$, $D = t/q$, $\hat{D} = \hat{t} + \hat{s}(L - \hat{a})$ and $z = \bar{\rho}$. Suppose first $\hat{D} > 0$. The map $F(v) = \hat{D}v - (L - \hat{a})$ is increasing and affine, z is the projection of its root onto $[0, R]$ and $\rho \in [0, R]$, so $(z - \rho)F(z) \leq 0$: at an interior root $F(z) = 0$; at $z = R$ the root exceeds R , so $F(R) < 0$ and $z - \rho \geq 0$; at $z = 0$ the root is negative, so $F(0) > 0$ and $z - \rho \leq 0$. Since $L - \hat{a} = u - e_a$ and, by (5), $\hat{D} = D + e_t + ue_s - \hat{s}e_a$ with $u = \rho D$,

$$F(z) = D(z - \rho) + (1 - z\hat{s})e_a + ze_t + zue_s,$$

hence $D(z - \rho)^2 \leq |z - \rho|((1 - z\hat{s})|e_a| + z|e_t| + zu|e_s|)$ and, with $D = t/q$,

$$|z - \rho| \leq q(1 - z\hat{s})\frac{|e_a|}{t} + qz\frac{|e_t|}{t} + z\rho|e_s|.$$

Now $0 < 1 - RS \leq q(1 - z\hat{s}) \leq 1$, $qz \leq R$ and $z\rho \leq R^2$, so $|z - \rho| \leq e_R/t$; and $|z - \rho| \leq R$ because both lie in $[0, R]$. If $\hat{D} \leq 0$, then $\hat{t} \leq \hat{s}(e_a - u) \leq \hat{s}|e_a|$ and $t = \hat{t} - e_t \leq |e_t| + \hat{s}|e_a|$, so $Rt \leq R|e_t| + R\hat{s}|e_a| \leq R|e_t| + |e_a| \leq e_R$; the fallback value and ρ lie in $[0, R]$, so $|z - \rho| \leq R = \min\{R, e_R/t\}$. For optimality take $s = \hat{s} = 0$, $\hat{t} = t$, $\rho = r_0$ and $\hat{a} = a + ct$ with $0 < c < r_0 < R$: then $|\bar{\rho} - \rho| = c = e_R/t$ exactly, with a positive denominator and no projection active. This proves (i). For (ii) integrate (i), split at $t = \tau$ and bound e_R^2/t^2 by e_R^2/τ^2 on $\{t > \tau\}$. For (iii) take $\tau = \varepsilon_R^{2/(\beta+2)} \leq u_0$; then $R^2 A \tau^\beta + \varepsilon_R^2/\tau^2 = (R^2 A + 1) \varepsilon_R^{2\beta/(\beta+2)}$. (iv) is (i) with $t \geq t_0$. $\square$

The middle expression in (ii) keeps the association between the errors and the transmission and is the one the tables evaluate; the last needs no power law. Both are computable on a finite test table; (iii) is a population statement whose tail hypothesis a finite table cannot establish down to $t = 0$. The exponent in (iii) cannot be improved from its hypotheses.

Proposition 7 (the exponent is attained). *For every $\beta > 0$, every $0 < c < r_0 < R$ and every $\tau \in (0, 1)$ there are an evaluation distribution and a predictor satisfying (15) with $S = 0$ and $F_t(u) = u^\beta$ on $(0, 1)$, namely $\rho = r_0$, $a = s = 0$, $\hat{t} = t$, $\hat{s} = 0$ and $\hat{a} = ct \mathbf{1}\{t \leq \tau\}$, for which $\bar{\rho} = r_0 - c \mathbf{1}\{t \leq \tau\}$,*

$$\varepsilon_R^2 = \frac{c^2 \beta}{\beta + 2} \tau^{\beta+2}, \quad \mathbb{E}_P |\bar{\rho} - \rho|^2 = c^2 \tau^\beta, \quad \text{so} \quad \mathbb{E}_P |\bar{\rho} - \rho|^2 = c^{4/(\beta+2)} \left(\frac{\beta + 2}{\beta} \right)^{\beta/(\beta+2)} \varepsilon_R^{2\beta/(\beta+2)}.$$

Proof. $L = r_0 t$ and $\hat{\rho} = (r_0 t - \hat{a})/t$; $e_R = |\hat{a}| = ct \mathbf{1}\{t \leq \tau\}$; with density $\beta u^{\beta-1}$ on $(0, 1)$, $\mathbb{E}_P[t^2 \mathbf{1}\{t \leq \tau\}] = \beta \tau^{\beta+2}/(\beta + 2)$; the retrieval error equals c exactly on $\{t \leq \tau\}$, an event of probability τ^β . Eliminate τ . $\square$

This is sharpness of the transfer inequality for the physical inverse, not a minimax statement over retrieval algorithms. What decides the rate is where the errors sit relative to the transmission.

Theorem 8 (conditional error profile). *Suppose $\mathbb{E}_P[e_R^2 | t] \leq \sigma^2 t^{2\nu}$ P -almost surely for some $0 \leq \nu \leq 1$. Then*

$$\mathbb{E}_P|\bar{\rho} - \rho|^2 \leq \mathbb{E}_P \min\{R^2, \sigma^2 t^{-2(1-\nu)}\}.$$

Under the tail hypothesis of Theorem 6(iii) and $\nu < 1$, as $\sigma \rightarrow 0$,

$$\mathbb{E}_P|\bar{\rho} - \rho|^2 = \begin{cases} O(\sigma^{\beta/(1-\nu)}), & \beta < 2(1-\nu), \\ O(\sigma^2 \log(1/\sigma)), & \beta = 2(1-\nu), \\ O(\sigma^2), & \beta > 2(1-\nu), \end{cases}$$

and at $\nu = 1$, $\mathbb{E}_P|\bar{\rho} - \rho|^2 \leq \sigma^2$ with no tail hypothesis. In the model of Proposition 7 the predictor $\hat{a} = a + \min\{ct, \sigma t^\nu\}$ has $e_R \leq \sigma t^\nu$ and $|\bar{\rho} - \rho| = \min\{c, \sigma t^{\nu-1}\}$, and attains each order.

Proof. Conditional on t the squared error is at most both R^2 and e_R^2/t^2 by Theorem 6(i), so $\mathbb{E}[|\bar{\rho} - \rho|^2 | t] \leq \min\{R^2, \mathbb{E}[e_R^2 | t]/t^2\} \leq \min\{R^2, \sigma^2 t^{2\nu-2}\}$; integrate. With $t_* = (\sigma/R)^{1/(1-\nu)}$, $\mathbb{E}_P \min\{R^2, \sigma^2 t^{2\nu-2}\} = R^2 F_t(t_*) + \sigma^2 \mathbb{E}_P[t^{2\nu-2} \mathbf{1}\{t > t_*\}]$; the first term is at most $R^2 A t_*^\beta = O(\sigma^{\beta/(1-\nu)})$ once $t_* \leq u_0$, and $\mathbb{E}_P[t^{2\nu-2} \mathbf{1}\{t > t_*\}] = \int_0^{t_*^{2\nu-2}} \mathbb{P}(t_* < t < v^{-1/(2-2\nu)}) dv \leq u_0^{2\nu-2} + A \int_{u_0^{2\nu-2}}^{t_*^{2\nu-2}} v^{-\beta/(2-2\nu)} dv$, which is $O(1)$ for $\beta > 2(1-\nu)$, $O(\log(1/\sigma))$ for $\beta = 2(1-\nu)$ and $O(t_*^{\beta-2(1-\nu)})$ for $\beta < 2(1-\nu)$. At $\nu = 1$ the bound is σ^2 before any tail enters. In the model, $c^{-2} \mathbb{E}_P|\bar{\rho} - \rho|^2 = \mathbb{E}_P \min\{1, (\sigma/c)^2 t^{2\nu-2}\}$, the same integral with equality. $\square$

At $\nu = 0$ the theorem recovers the uniform-error regimes under the weaker hypothesis of a bounded conditional second moment; at $\nu = 1$ an emulator whose conditional root-mean-square error decreases proportionally to the transmission loses no rate. The central diagnostic is therefore the profile $t \mapsto \mathbb{E}[e_R^2 | t]$ and not only the tail $t \mapsto F_t(t)$: the same overall error concentrated in the opaque entries gives the worst-case loss, and proportional to the transmission gives none.

8.4 Perturbing a learned-mean kernel predictor

Fix the training sample and write a corrected predictor as $p(x) = h(x) + k_x^\top H^{-1}(Y - M)$ with $H = K + \lambda I$ positive definite, $Y \in \mathbb{R}^{n \times q}$ the targets and $M \in \mathbb{R}^{n \times q}$ the training-row means used as correction labels. Let a second fitted system be $\tilde{p}(x) = \tilde{h}(x) + \tilde{k}_x^\top \tilde{H}^{-1}(Y - \tilde{M})$ with $\tilde{H}$ symmetric positive definite; both systems may depend on Y through the mean, the features, the selected ridge and any kernel approximation. Put $B = H^{-1/2}(\tilde{H} - H)H^{-1/2}$, whose least eigenvalue exceeds -1 , $Q = (I + B)^{-1}$, $\gamma_Q = \|Q\|_{\text{op}} = 1/(1 + \lambda_{\min}(B))$, $\eta_Q = \|Q - I\|_{\text{op}} = \max_i |b_i|/(1 + b_i)$ over the eigenvalues b_i of B , and

$$g = H^{-1/2} k_x, \quad \tilde{g} = H^{-1/2} \tilde{k}_x, \quad v = \tilde{g} - g, \quad W = H^{-1/2}(Y - M), \quad U = H^{-1/2}(\tilde{M} - M).$$

A conditioning argument bounds such a difference by the condition number of H times a norm of the change [108, 106, 109]. The statement below instead keeps the residual action W and the label shift U separate, which is what makes it readable on data.

Proposition 9 (same-data perturbation). 1. $\tilde{p}(x) - p(x) = (\tilde{h}(x) - h(x)) + v^\top W - (Q\tilde{g})^\top (BW + U)$.

$$2. \|\tilde{p}(x) - p(x)\|_2 \leq \|\tilde{h}(x) - h(x)\|_2 + \|W^\top v\|_2 + \|Q\tilde{g}\|_2 \|BW + U\|_{\text{op}}.$$

$$3. \|\tilde{p}(x) - p(x)\|_2 \leq \|\tilde{h}(x) - h(x)\|_2 + \|W\|_F (\gamma_Q \|v\|_2 + \eta_Q \|g\|_2) + \gamma_Q \|U\|_F \|\tilde{g}\|_2.$$

$$4. \text{ If } \tilde{H} = H + sI \text{ with the kernel and the means unchanged, then } \|\tilde{p}(x) - p(x)\|_2 \leq \|g\|_2 \|W\|_F s / (\lambda_{\min}(H) + s).$$

Proof. $\tilde{H} = H^{1/2}(I + B)H^{1/2}$ gives $\tilde{H}^{-1} = H^{-1/2}QH^{-1/2}$ and $Q = I - QB$. Hence $\tilde{k}_x^\top \tilde{H}^{-1}(Y - \tilde{M}) - k_x^\top H^{-1}(Y - M) = \tilde{g}^\top Q(W - U) - g^\top W = \tilde{g}^\top W - \tilde{g}^\top QBW - \tilde{g}^\top QU - g^\top W = v^\top W - (Q\tilde{g})^\top (BW + U)$, which is (i) since Q is symmetric; (ii) bounds its two terms. For (iii) write the same difference as $v^\top QW + g^\top (Q - I)W - \tilde{g}^\top QU$ and use $\|Q\|_{\text{op}} = \gamma_Q$ and $\|Q - I\|_{\text{op}} = \|QB\|_{\text{op}} = \eta_Q$. For (iv), $v = 0$, $U = 0$ and $B = sH^{-1}$ has eigenvalues $s/\lambda_i(H) \geq 0$, so $\eta_Q = \max_i (s/\lambda_i) / (1 + s/\lambda_i) = s / (\lambda_{\min}(H) + s)$. $\square$

Only positive definiteness of $\tilde{H}$ is needed: a large positive eigenvalue of B , from a larger ridge or an approximation that dominates the exact Gram matrix nowhere, does not make the bound vacuous. Item (ii) keeps the interaction between the perturbation and the residual: with $v = 0$, $U = 0$ and $BW = 0$ the predictor does not move however large $\|B\|$ is. The bounds are stated in the residual W , not in Y , they use no independence between the representation, the mean, the ridge and Y , and they transfer between fitted predictors without being generalisation statements.

For component j the physical target at band b is $f_{j,b} = \mu_{j,b} + \sum_\ell (A_j)_{b\ell} z_{j,\ell} + \pi_{j,b}$, with A_j the standardiser scale times the retained PCA basis and π_j the truncation residual, and the pipeline predicts the retained coordinates. For a frozen realised system $p_\ell(x) = h_\ell(x) + a_x^\top (y_\ell - h_{\ell,X})$ with $y_\ell = f_{\ell,X} + \varepsilon_\ell$, the conditional power-function bound of the kernel paper gives

$$|p_\ell(x) - f_\ell(x)| \leq B_\ell \tilde{P}_\lambda(x) + |a_x^\top \varepsilon_\ell|, \quad (18)$$

with B_ℓ the residual RKHS norm, a hypothesis and not an observed certificate, and the second term zero for deterministic simulator targets, which is the form used for EMIT. A perturbed system adds the bound $\Delta_\ell(x)$ of Proposition 9. If $|\hat{z}_{j,\ell} - z_{j,\ell}| \leq u_{j,\ell}$, then

$$|\hat{f}_{j,b} - f_{j,b}| \leq |\pi_{j,b}| + \sum_\ell |(A_j)_{b\ell}| u_{j,\ell}, \quad (19)$$

and a change of basis or standardiser between two systems adds its own reconstruction term. Assembling $e_R = |e_a| + R|e_t| + R^2 t|e_s|$ from (19) for the three physical components and applying Theorem 6 closes the chain. If the inverse is evaluated at $L_{\text{obs}} = L + \eta$, the proof of Theorem 6 goes through with e_a replaced by $e_a - \eta$, so $|\bar{\rho} - \rho| \leq \min\{R, (|e_a - \eta| + R|e_t| + R^2 t|e_s|)/t\}$. The theorem is an exact-arithmetic statement: forming $L = a + \rho t/q$ in floating point loses the term $\rho t/q$ once t is below $2^{-52}a$, and the operational inverse then returns zero even with exact components; the tables separate the operational computation from an exact-arithmetic check on those entries.

8.5 The stacking objective

With every base head projected first ($\hat{t} \mapsto \max\{\hat{t}, 0\}$, $\hat{s} \mapsto \Pi_{[0, S]}(\hat{s})$) and combined by simplex weights w per component, the combination stays in the constraint set, its predictions are affine in w , and

$$J_\tau(w) = \mathbb{E}_P \left[\frac{e_a(w)^2 + R^2 e_t(w)^2 + R^4 t^2 e_s(w)^2}{(t \vee \tau)^2} \right] \quad (20)$$

is a convex quadratic in w . By Theorem 6(i) and $e_R^2 \leq 3(e_a^2 + R^2 e_t^2 + R^4 t^2 e_s^2)$,

$$\mathbb{E}_P |\bar{\rho}_w - \rho|^2 \leq R^2 F_t(\tau) + 3J_\tau(w). \quad (21)$$

Projecting an unrestricted combination inside the loss would destroy the convexity. The threshold τ is fixed before any test read under a declared coverage requirement, because a threshold can make the task easier by discounting the difficult bands; candidate stacks at a few values of τ are frozen and selected under Proposition 5.

8.6 The interpolation regime

Two results on the Gaussian kernel at a bandwidth that does not scale with the dimension describe a regime the corpora of this paper are not in; they are stated for the record and proved elsewhere.

Theorem 10. *For every fixed $h > 0$, q_0 , B and σ , and every $\epsilon > 0$, $\sum_\ell \Pr\{\mathcal{R}_\ell > \epsilon\} < \infty$ along $d_\ell = 2^{2^\ell}$, $n_\ell = 2^{2^{2^\ell}}$. In particular $\mathcal{R}_\ell \rightarrow 0$ in probability, and almost surely under any coupling of the levels.*

Theorem 11. *For every fixed $h > 0$, $\liminf_{d \rightarrow \infty} \mathbb{E} \mathcal{E}_{2^d, d} / \sigma^2 \geq 5/36$, the expectation being over the design and the noise.*

The first concerns uniform inputs on a sphere of growing dimension, independent Gaussian noise and a signal in a fixed finite sum of harmonic spaces; the second a noise-only problem at $n = 2^d$, and the same argument gives larger constants, so 5/36 is not sharp. A network with two or six thousand features does not place a Matérn regression on those features in either setting. The bridge to the corpora is Proposition 4 together with Theorems 6 and 8.

9 Retrieval error and the transmission profile

Theorem 6 makes the reflectance error at an entry a function of two things, the retrieval-weighted component error e_R and the transmission t , and Theorem 8 makes the rate a function of where e_R sits relative to t . Table 9 reports, for the wide-network heads of Section 6 on the test rows of each split, the quantities those theorems name: ε_R and the exponent 2ν of the conditional profile $t \mapsto \mathbb{E}[e_R^2 \mid t]$ fitted over the observed transmission range, the raw inverse (2) with its non-finite and non-positive-denominator failures, the constrained retrieval (16) with its failures and clips, the pointwise quantity $\mathbb{E} \min\{R^2, e_R^2/t^2\}$, the number of entries at which the pointwise inequality of Theorem 6(i) fails, and the ratio of the middle bound of Theorem 6(ii), minimised over a threshold ladder, to the measured mean-square error of the constrained retrieval. The albedo bound S is the largest physical value on the training rows; the table carries 454 training entries with spherical albedo above one, all in deep absorption bands at vanishing flux, which

family	ε_R	2ν (obs.)	raw inverse p95 / p99	constrained p95 / p99	failures / clipped at R [%]	$\mathbb{E} \min\{R^2, \varepsilon_R^2/t^2\}$	violations	bound / MSE
input-scaled kernel	0.0886±0.0331	0.16±0.11	0.696±0.000 / 0.97±0.02	0.266±0.018 / 0.70±0.00	1.09±0.03 / 3.56±0.07	0.064±0.000	0	3.9±0.2
network	0.1664±0.0156	0.20±0.14	0.677±0.027 / 1.08±0.25	0.254±0.101 / 0.70±0.00	1.20±0.90 / 3.47±1.26	0.057±0.003	0	3.6±0.8
network ensemble	0.1085±0.0254	0.18±0.14	0.625±0.013 / 0.86±0.08	0.239±0.078 / 0.70±0.01	0.61±0.33 / 2.56±0.92	0.051±0.002	0	3.7±0.9
network + residual kernel	0.0654±0.0338	0.23±0.12	0.486±0.047 / 0.89±0.01	0.200±0.000 / 0.70±0.00	0.80±0.03 / 2.50±0.10	0.044±0.001	0	3.9±1.1
ensemble + residual kernel	0.0690±0.0381	0.22±0.14	0.464±0.033 / 0.88±0.01	0.200±0.000 / 0.70±0.00	0.85±0.01 / 2.43±0.07	0.043±0.001	0	3.9±1.3
kernel on features	0.0472±0.0404	0.28±0.14	0.076±0.012 / 0.89±0.05	0.077±0.013 / 0.70±0.00	0.71±0.02 / 1.53±0.09	0.031±0.001	0	4.8±2.1
kernel on concatenated features	0.0634±0.0680	0.28±0.16	0.060±0.004 / 0.92±0.03	0.061±0.004 / 0.70±0.00	0.72±0.02 / 1.49±0.02	0.030±0.000	0	5.3±2.4
coordinatewise selection	0.0721±0.0648	0.25±0.15	0.098±0.031 / 0.85±0.03	0.099±0.032 / 0.70±0.00	0.73±0.02 / 1.61±0.13	0.033±0.002	0	4.7±2.0
convex stack	0.0600±0.0661	0.25±0.14	0.143±0.015 / 1.09±0.05	0.151±0.015 / 0.70±0.00	0.84±0.05 / 1.78±0.08	0.036±0.001	0	4.4±2.0

Table 9: Transmission-conditioned retrieval on the wide-network heads at $\rho = 0.7$: the retrieval-weighted error, the observed-range profile exponent, the raw and constrained inverses with their failure and clip rates, the pointwise bound of Theorem 6, its violation count over the test entries, and the ratio of the distributional bound to the measured error. Mean and standard deviation over the splits listed in the file.

are outside the hypothesis and are counted rather than used, and no held-out entry exceeds S . The surface reflectance is $\rho = 0.7$ and $R = 0.9$.

The table’s columns line up with the theorem. The retrieval-weighted error ε_R runs from 0.047 for the kernel on features to 0.166 for the single network; the observed-range profile exponents lie between 0.16 and 0.28; the pointwise inequality of Theorem 6(i) is violated at no test entry of any head; and the distributional bound of item (ii), minimised over the threshold ladder, exceeds the measured mean-square error of the constrained retrieval by a factor between 3.6 and 5.3, the slack being the tail mass term. The raw inverse and the constrained retrieval separate the heads in the same order but at different scales: the raw 95th percentile is 0.06–0.10 for the feature kernels and the selection against 0.46–0.70 for the networks and the input kernel, and the constrained 95th percentile 0.06–0.15 against 0.20–0.27, with 0.6 to 1.2% of the entries failing the positivity test and 1.5 to 3.6% clipped at R . Three facts carry the comparison between the feature kernel and its stack. One entry in twenty has transmission below 3×10^{-4} of the training flux scale and one in a hundred below 10^{-10} ; the observed-range tail slope of F_t is about 0.23. The absolute component errors of every head are not concentrated at low transmission: the share of ε_R^2 below 0.03 of the scale is under a thousandth for every head. The retrieval tails are therefore produced by dividing an ordinary absolute error by a vanishing t , which is the mechanism of Theorem 6(i), and the all-band 95th and 99th percentiles measure the opaque bands rather than the heads. The profile exponents in the table, between 0.17 and 0.34 on the observed range, are far from the value $2\nu = 2$ at which Theorem 8 gives a rate without loss, so on the all-band distribution the rate loss of Theorem 6(iii) is the operative regime, while on a passband $t \geq t_0$ the error transfers at the rate of the components by Theorem 6(iv). The comparison between heads is accordingly made decile by decile in transmission and on the constrained retrieval, whose error is bounded by R and is covered by Proposition 5; the raw inverse and its failure counts are reported beside it.

10 Representation, correction and combination

Table 10 replicates the three-member concatenated-feature kernel of Section 6 at four further splits, with the three members trained at width 2,000 for 500 epochs and every prediction stored in double precision. The concatenated kernel has a mean radiance error of 0.069% against 0.076% for the kernel on a single member’s features; the paired gain is 0.0076 ± 0.0053 radiance points, positive at three of the four splits and -0.0002 at the fourth. Its median reflectance error is 0.033 points against 0.040, and its within-split 95th percentile 6.0 points against 7.5: at these splits

seed	kernel on features	kernel on concatenated features	convex stack	network + residual kernel	gain of concatenation (radiance points)
104	0.0924 / 0.044 / 9.25	0.0810 / 0.025 / 6.25	0.0781 / 0.024 / 12.77	0.1843 / 0.115 / 54.71	+0.0114
105	0.0611 / 0.031 / 7.31	0.0508 / 0.028 / 6.02	0.0475 / 0.024 / 14.50	0.1349 / 0.087 / 42.61	+0.0103
106	0.0789 / 0.043 / 7.67	0.0702 / 0.035 / 6.28	0.0687 / 0.034 / 16.29	0.1608 / 0.108 / 51.84	+0.0087
107	0.0731 / 0.042 / 5.82	0.0732 / 0.044 / 5.41	0.0833 / 0.043 / 13.84	0.1476 / 0.103 / 46.58	-0.0002
mean	0.0764±0.0130 / 0.040±0.006 / 7.51±1.41	0.0688±0.0128 / 0.033±0.009 / 5.99±0.40	0.0694±0.0158 / 0.031±0.009 / 14.35±1.48	0.1569±0.0211 / 0.103±0.012 / 48.93±5.40	0.0076±0.0053

Table 10: The kernel on the concatenated last-layer features of three wide members against the kernel on a single member’s features, the convex stack and the residual-corrected network, per split, with the paired gain of the concatenation in radiance points.

predictor	radiance rel. L^2 [%]	ε_R	constrained p95 / p99	raw inverse p95
network	0.8609±0.1051	0.1664±0.0156	0.254±0.101 / 0.70±0.00	0.677±0.027
network + residual kernel	0.1560±0.0184	0.0654±0.0338	0.200±0.000 / 0.70±0.00	0.486±0.047
kernel on features	0.0755±0.0114	0.0472±0.0404	0.077±0.013 / 0.70±0.00	0.076±0.012
input-scaled kernel	0.2454±0.0106	0.0886±0.0331	0.266±0.018 / 0.70±0.00	0.696±0.000
network ensemble	0.5050±0.0592	0.1085±0.0254	0.239±0.078 / 0.70±0.01	0.625±0.013
ensemble + residual kernel	0.1379±0.0157	0.0690±0.0381	0.200±0.000 / 0.70±0.00	0.464±0.033
kernel on concatenated features	0.0688±0.0128	0.0634±0.0680	0.061±0.004 / 0.70±0.00	0.060±0.004
stack, mean relative norm	0.0754±0.0116	0.0573±0.0575	0.141±0.033 / 0.70±0.00	0.133±0.029
stack, row-relative squared error	0.1130±0.0454	0.0518±0.0414	0.180±0.044 / 0.70±0.00	0.281±0.122
stack, J_τ at $\tau = 0.001$	0.2203±0.1042	0.0633±0.0298	0.135±0.060 / 0.69±0.01	0.213±0.227
stack, J_τ at $\tau = 0.003$	0.2159±0.1100	0.0736±0.0522	0.103±0.055 / 0.69±0.01	0.182±0.238
stack, J_τ at $\tau = 0.01$	0.2303±0.1023	0.0761±0.0508	0.107±0.053 / 0.59±0.22	0.185±0.236
stack, J_τ at $\tau = 0.03$	0.2304±0.1024	0.0761±0.0508	0.107±0.053 / 0.59±0.22	0.185±0.235
stack, J_τ at $\tau = 0.1$	0.2298±0.1013	0.0739±0.0463	0.110±0.052 / 0.59±0.22	0.188±0.235
stack, J_τ at $\tau = 0.3$	0.1491±0.0508	0.0615±0.0468	0.117±0.048 / 0.69±0.03	0.183±0.196
stack, J_τ at $\tau = 1.0$	0.1175±0.0422	0.0553±0.0469	0.156±0.032 / 0.70±0.00	0.152±0.032
stack, J_τ at $\tau = \infty$	0.1262±0.0406	0.0559±0.0420	0.168±0.039 / 0.70±0.00	0.194±0.095

Table 11: Combinations of the heads fitted on the validation rows and read once on the test rows, at $\rho = 0.7$: the single heads, the stack under the mean relative norm per component (the objective of the campaign stack), the stack under the row-relative squared error, and the stack under J_τ of (20) on projected heads at several thresholds. Mean and standard deviation over the lanes listed in the file.

the concatenation improves the tail as well as the mean. The convex stack of the two kernels and the residual-corrected network has the same mean radiance error as the concatenated kernel (0.069%) and a 95th percentile of 14 points, the tail cost of stacking seen in Section 6; the residual-corrected network alone stands at 0.157% and 49 points. The split-level pairing is the unit of comparison; the test sets overlap and the standard deviations describe the spread over splits, not confidence intervals.

The stacking objective decides the albedo head. Under any squared objective, whether in physical units or row-relative, the simplex weight on the spherical albedo goes to the network, whose albedo error is smaller in mean square, and the radiance error of the combination is worse than the feature kernel alone: the stack under the row-relative squared error reaches 0.113% against 0.076% for the kernel on features and 0.069% for the kernel on concatenated features. Under the mean relative norm the weight goes to the kernel, whose albedo errors are mostly small with a heavier tail, and the combination ties the single-member feature kernel (0.075%) without reaching the concatenated kernel. The retrieval tails move the other way: the constrained 95th percentile of the mean-relative stack is 0.14 against 0.06 for the concatenated kernel and 0.08 for the single-member kernel, so the stack that matches the kernel’s forward error does not match its inverse tail. The theorem’s objective J_τ inherits the squared form and therefore the network’s

predictor	Y_1	Y_2	Y_3	Y_4
Matérn-1/2	0.179±0.008	0.672±0.003	0.418±0.012	2.326±0.103
Matérn-3/2	0.092±0.002	0.391±0.008	0.207±0.016	2.481±0.043
Matérn-5/2	0.089±0.017	0.434±0.007	0.203±0.003	2.746±0.082
Gaussian	0.092±0.002	0.525±0.113	0.228±0.051	3.502±0.501
spectral mixture, four components	0.089±0.010	0.237±0.014	0.099±0.013	3.468±0.263
NNGP, ReLU	0.175±0.001	0.674±0.015	0.350±0.009	2.548±0.057
NTK, ReLU	0.263±0.002	1.055±0.015	0.621±0.014	2.321±0.048
additive Matérn-5/2	4.091±0.093	20.266±1.035	11.430±0.237	4.413±0.077
ARD by kernel flow	0.322±0.015	0.257±0.019	0.199±0.005	2.418±0.114
ARD by marginal likelihood	0.281±0.039	0.341±0.042	0.183±0.011	2.702±0.027
convex kernel sum	0.087±0.007	0.253±0.020	0.099±0.013	2.557±0.099
network	0.819±0.027	2.289±0.019	0.910±0.057	3.482±0.085
network + residual kernel	0.523±0.007	1.337±0.026	0.588±0.035	3.055±0.121
kernel on features, Matérn-5/2	0.105±0.006	0.211±0.016	0.114±0.008	2.303±0.216
kernel on features, Gaussian	0.126±0.007	0.262±0.018	0.138±0.014	2.426±0.214
kernel on features, NTK	0.152±0.005	0.307±0.015	0.156±0.007	2.150±0.059
kernel on features, ARD	0.116±0.007	0.214±0.014	0.111±0.004	2.217±0.207
coordinatewise selection	0.090±0.010	0.360±0.008	0.105±0.010	2.797±0.447
convex stack of the heads	0.066±0.003	0.156±0.010	0.077±0.004	1.956±0.037
ridge stack of the heads	0.087±0.005	0.288±0.031	0.087±0.005	2.310±0.071

Table 12: Kernel ridge regression on the six state inputs with ten kernels, one component at a time (rank-64 targets, error in physical units): the Matérn family and the Gaussian, a four-component spectral mixture [71], the arc-cosine kernels of an infinite-width ReLU network (NNGP [92, 90] and NTK [93]), a first-order additive Matérn [72], and two automatic-relevance-determination fits (kernel flow [57] and marginal likelihood [8]); a convex sum of the kernels with weights chosen on validation rows [74, 75]; a network of width 512 and depth 3 trained for 150 epochs, its residual correction and four kernels on its last hidden layer; and three combiners fitted on the validation rows and read once on the test rows. Every hyperparameter is chosen on the validation rows of the lane. Mean and standard deviation over the seeds listed in the file.

albedo; its promise is the bound (21), not a forward improvement, and the table shows the price: 0.22–0.23% radiance error for $\tau \leq 0.1$, 0.15% at $\tau = 0.3$ and 0.12% at $\tau \geq 1$, with constrained 95th percentiles between 0.10 and 0.17 and the 99th percentile at the clip $R = 0.7$ except at τ between 0.01 and 0.1, where it falls to 0.59 ± 0.22 . The threshold τ is fixed before the test read, with the validation coverage $F_t(\tau)$ it implies recorded in the file.

Table 12 asks whether the kernel on the state inputs is the limiting choice, and it is not. Among the ten kernels the spectral mixture and the Matérn-5/2 tie on Y_1 (0.089%) and the spectral mixture is best on Y_2 and Y_3 (0.237% and 0.099%); the Matérn-3/2 and the Gaussian are within a few thousandths of a point on Y_1 , the relevance-determination fits within two hundredths on Y_2 and a tenth on Y_3 ; the arc-cosine kernels are two to six times worse; and the additive kernel fails outright (4 to 20%), so no component is a sum of one-dimensional functions of the state. On the spherical albedo Y_4 every kernel but the additive one sits between 2.3 and 3.5%, the rough Matérn-1/2 and the NTK at the low end. The convex kernel sum puts all of its weight on one kernel in every lane, the kernel with the lowest validation error, and adds nothing. The network of this table is a modest one, and its residual correction improves it by a third or more on the first three components and by a tenth on Y_4 without reaching the best kernel, while the Matérn-5/2 kernel on its last hidden layer reaches 0.105%, 0.211%, 0.114% and 2.30%: ahead of every input kernel on Y_2 and Y_4 and within two hundredths of a point of the best on Y_1 and Y_3 . The convex stack of all heads is the best predictor of the table on every component (0.066%, 0.156%, 0.077% and 1.96%), 22 to 26% below the best single head on Y_1 to Y_3 and

predictor	alone	+ residual kernel	kernel on features
input-scaled kernel	0.941±0.109 / 0.284±0.078	–	–
network, control	2.266±0.031 / 0.809±0.029	1.903±0.055 / 0.617±0.033	0.955±0.038 / 0.153±0.052
network, weight decay	2.269±0.027 / 0.820±0.014	1.880±0.072 / 0.618±0.040	0.948±0.048 / 0.154±0.048
network, early stopping	2.716±0.445 / 1.293±0.373	1.949±0.073 / 0.789±0.127	0.986±0.032 / 0.147±0.054

Table 13: Regularisation of the network under a fixed budget (three runs per arm of 150 epochs, width 512, depth 3, the run with the lowest validation error kept): the control, a weight-decay grid [86] (10^{-5} , 10^{-4} , 10^{-3}) and an early-stopping grid [87] (patience 3, 5 and 10 epochs), each alone, with the residual Matérn-5/2 kernel added to its mean, and with the same kernel on its last hidden layer, beside the input-scaled kernel. Cells give the mean relative L^2 error over the four components and the radiance error on the test rows, mean and standard deviation over the seeds listed in the file.

nine percent below it on Y_4 ; the coordinatewise selection improves on no single head, and the ridge stack only on Y_3 . The spread over seeds is a few percent of the value for most heads and a tenth for the Gaussian kernel on Y_2 and Y_3 , whose length scale the validation rows select less stably.

Table 13 varies the regularisation of the network and applies the two kernel steps to each arm. Weight decay at the value the validation rows select (10^{-5} at two seeds, 10^{-4} at the third) leaves the network where the control is (2.27% against 2.27% in the component error, 0.82% against 0.81% in radiance), and early stopping at the selected patience, which halts training between epochs 37 and 97, makes it worse and less stable ($2.72 \pm 0.45\%$ and $1.29 \pm 0.37\%$). The residual kernel takes every arm to 1.88–1.95% in the components and 0.62–0.79% in radiance, and the kernel on its features to 0.95–0.99% and 0.15%: level with the input-scaled kernel in the components (0.94%) and at half its radiance error (0.28%). The order of the arms in the component error survives both kernel steps, with the differences between them compressed six times by the residual step and twelve times by the feature step, and in radiance the three feature kernels are indistinguishable. The last hidden layer of an under-trained network is as good a kernel input as that of the control.

Table 14 measures the perturbation bounds where the theorem allows them. The exact identity of Proposition 9(i) holds to rounding at every system. Every bound holds at every held-out row. The signed-spectrum bound is finite for a tenfold ridge increase, where $\|B\|$ is near nine, with the coefficient $s/(\lambda_{\min}(H) + s)$ of item (iv); the bound that requires $\|B\| < 1$ does not apply there. On the Nyström ladder of the input kernel the residual-action bound exceeds the median movement by a factor between 57 and 408 and the signed-spectrum bound by 1,500 to 4,000, so the first is tighter by a factor of six to seventy; on the feature kernels the same ratios are in the thousands and the tens of millions, the slack sitting in the product of the whitened residual energy and the whitened query norm rather than in the eigenvalues. Compressed inputs make B indefinite with a least eigenvalue near -1 , and there the residual-action bound, at 14 to 64 times the movement, is the only informative one.

Table 15 reads the correction along the ridge. At the smallest ridge the kernel step interpolates its own residual: the effective degrees of freedom are 3,841 of 4,000 rows for the input kernel on Y_1 , the noise factor is 0.92, and the correction fitted and read on the same half of the rows reports a 0.08% error against 2.5% when it is read on the other half. Raising the ridge to 10^{-4} brings the two readings within a tenth of each other (2.2% against 2.4%) while the kernel alone degrades from 0.18% to 0.98% and the corrected predictor only from 0.22% to 0.41%; at 10^{-2} the kernel alone is at 5.8% and the corrected predictor at 0.51%. The same pattern

holds for every component and for the feature kernel: the network mean carries the prediction when the kernel step is over-smoothed, and across the thirty-two rows of the table the corrected predictor is worse than the kernel alone only at the smallest ridge of the input kernel on Y_1 and at one intermediate ridge of the feature kernel on Y_2 , by four hundredths and one hundredth of a radiance point, which is what Proposition 4 says the correction can and cannot do.

11 Synthetic reference for the transfer theorem

Table 16 is a check of the mathematics on manufactured data, not a result about the corpora. The average-error exponents match $\beta/(\beta + 2)$ to three digits. The uniform-family slopes fall below their limits over the finite ladder, and the analytic curve on the same ladder gives the same slopes to three digits: at $\beta = 2$ the local slope is $2 - 2/(1 + 2\log(1/\varepsilon))$, which is 1.72 on the ladder used. The Monte Carlo agrees with the analytic curve of the profile family to at most five parts in ten thousand at every β , and with that of the uniform family to three parts in a thousand except at $\beta = 2$, where the difference is one part in a hundred; the pointwise inequality holds at every one of the ten million samples.

12 The matched benchmark and the transfer contrast

On the correction-coefficient corpus the comparison is made under the benchmark’s protocol and metrics, and the claim is improvement over the published values under that protocol and nothing wider. The residual-corrected network improves on the published pKANrtm values in every one of the four metrics on both splits. On the release’s own split with the state inputs alone it reaches an RMSE of 0.00581 ± 0.00004 against the published 0.00619, a MAE of 0.00151 against 0.00186, $R^2 = 0.9982$ against 0.9947 and a SMAPE of 3.54% against 3.91%; with the low-fidelity 6S coefficients as inputs, the benchmark paper’s own construction, the same head reaches 0.00578, 0.00148, 0.9982 and 3.19%. Out of distribution the margin widens: 0.00620 against the published 0.01022 in RMSE, 0.00252 against 0.00544 in MAE, 0.9980 against 0.9942 in R^2 and 3.37% against 5.40% in SMAPE with the state inputs, and 0.00596, 0.00232, 0.9982 and 2.92% with the 6S inputs. The ensemble without the correction is close behind the corrected head (RMSE 0.00591 and MAE 0.00170 on the release’s split, 0.00624 and 0.00258 out of distribution), the single network a little further back, and the kernel on the network’s features is the best of the kernels at 0.0087 (standard) and 0.0141 (OOD); the kernels on the raw state and the ridge readout are an order of magnitude worse, at 0.11 to 0.12, and the 6S inputs bring the raw-state kernel to 0.025 (standard) and 0.029 (OOD). The standard deviations are over three seeds on the release’s split and three initialisations on the out-of-distribution split.

On OCO-2 the reduced-coefficient and radiance metrics order the heads differently, and Table 18 measures by how much. The kernel on the network’s features trained under the flat loss is the best single head in the reduced coordinates (4.1%, 16.3% and 8.1% on the O_2 , weak CO_2 and strong CO_2 bands), but its radiance error (0.076%, 0.114%, 0.108%) is two to three times that of the same kernel on features trained under a radiance-weighted loss (0.031%, 0.048%, 0.064%), whose reduced errors are in turn four to six times worse (20–49%); the kernels on the raw state trail every feature kernel in both metrics (29–59% and 0.22–0.60%). The per-coordinate selection keeps the flat head’s reduced error and takes the radiance error to 0.028%, 0.051% and 0.060%; selecting on the radiance error rather than on the coefficients picks the same heads on two bands and changes one head on the third, to no measurable effect. The simplex

stack under the reduced metric holds the reduced error at the selection’s level and lowers the radiance error further, to 0.029%, 0.040% and 0.054%; the stack fitted on the exact radiance error, a convex quadratic in the weights solved to a KKT gap below 10^{-5} , gives the lowest radiance error on every band (0.026%, 0.039%, 0.051%), which is three to nine percent below the reduced-metric stack, and pays for it entirely in coordinates that carry little radiance: its reduced error is 18.6% on O_2 and 37.4% on weak CO_2 against 4.1% and 16.3%, and 8.9% against 8.1% on strong CO_2 . Every rule is fitted on the 2,000 validation rows and read once on the 2,000 test rows; the standard deviations are over the three seeds.

13 Comparison with OCO-2

The companion paper [11] applies the same construction to the OCO-2 radiative-transfer emulation problem of Lamminpää et al. [3]: per spectral band (O_2 , weak CO_2 , strong CO_2), a reduced atmospheric state of twenty or twenty-four coordinates maps to forty reduced radiance coefficients, with errors on the standardized coefficients and on the reconstructed radiance, under the ten-split protocol of Section 3. The corresponding results in [11] provide a comparison at a higher input dimension.

On this higher-dimensional problem, a kernel on the state trails the network by a factor of nine, four and five on the reduced coefficients: the isotropic Matérn reaches 40.6%, 58.6% and 43.4% on the three bands against 4.4%, 16.4% and 8.2% for the network. The residual correction adds no material gain in the reported OCO-2 comparisons. Little alignment with useful kernel directions is a possible explanation, but the public records do not measure the residual’s Gram-eigenvector coefficients. The design identity in Proposition 2 alone cannot establish that explanation on test data. The coordinate search that gives a kernel on the state its own length scales, a strong raw-input construction on the EMIT table, is not the strongest there either. Four learned metrics were run at the same ten splits on all three bands: a kernel-flow metric [57], a Mahalanobis metric under the same flow, the gradient outer product of the recursive feature machine [96], and per-input scales fitted by empirical Bayes. The last is the best of the four at every band, taking the isotropic kernel’s 40.6%, 58.6% and 44.3% on the coefficients to 18.7%, 25.3% and 16.9%, and on O_2 its radiance from 0.233% to 0.150%; the two flow metrics agree within half a point of each other at 26.9%, 29.8% and 23.9%, the recursive feature machine reads 28.2%, 47.9% and 31.3%, and an additive kernel 38.2%, 46.9% and 47.7%. The isotropic baseline agrees between the metric and baseline experiments on O_2 and weak CO_2 to two decimals; on strong CO_2 it reads 44.3% in the metric experiment and 43.4% in the baseline experiment. This band also has a spread of ± 3.4 because two of its ten splits land three points below the other eight. The learned metric reduces the isotropic error to between two fifths and a half of itself depending on the band. The feature kernel remains more accurate, with a band-dependent margin: the exact Matérn regression on the network’s frozen features reads 4.1%, 16.3% and 8.1% at the same splits, so the best state metric is 4.6 times that on O_2 but only 1.6 and 2.1 times it on the two CO_2 bands. The empirical-Bayes fit has substantial split variability: on O_2 the empirical-Bayes metric is bimodal across splits, eight of the ten landing near 17.2% and two at 24.7%, so its spread of ± 3.2 is thirty times the isotropic kernel’s ± 0.1 and its mean is not the value a single split is likely to show. The feature-kernel construction remains useful on both tasks: the exact Matérn regression on the network’s frozen features beats a ridge readout refitted on the same features at every split, on every band and under every training loss, and a per-coefficient choice among heads trained under different losses has lower mean error than the

release’s stored emulator under both reported metrics. In the archived loss campaign it wins the reduced-coefficient comparison at all ten splits on each band and the radiance comparison at ten, nine and ten splits on O₂, weak CO₂ and strong CO₂, respectively. These companion comparisons are not an independent EMIT validation; their tables are in [11].

The OCO-2 retrieval records overlap the release’s test set. The reference’s data release carries 976 retrieval records, each with an atmospheric state, the full-physics monochromatic radiance at that state, its rank-40 reconstruction, and the prediction of the release’s Gaussian-process emulator at the same forty coefficients. Every one of the 976 reduced states matches a stored test input within 6.67×10^{-16} ; the records map to 974 distinct test rows, and the nearest test-row index is the same in all three bands. The records are therefore not disjoint from the stored test set and cannot serve as an independent evaluation pool, though they do provide a different target, the full-physics radiance, at those states. The state reduction was replayed independently from the release’s two dimensional-reduction files and the `x_in` and `geo_in` fields of all 976 records; the replay code and its digests are archived with the per-run records.

14 Ten further emulation configurations

Table 19 extends the comparison to ten emulation configurations. Each family uses validation-based hyperparameter selection. The models and training budgets are specified below; these comparisons assess the selected family implementations rather than the best task-specific methods available for each corpus.

The corpora are the advection equation with discontinuous initial data from the operator suite of [102], which is also the source of the structural-mechanics problem of [11]; Burgers’ equation at three viscosities from PDEBench [103]; three one-step maps from The Well [104] — a two-dimensional turbulent radiative layer, an active-matter system, and a Helmholtz staircase; the ClimSim atmospheric physics emulation task [105]; and the paired correction-coefficient corpus of the physics-guided radiative-transfer work of [4], which is a third radiative-transfer emulation problem and the largest here at 487,790 rows. That last corpus is run in both of its configurations: from the atmospheric state alone, and with the low-fidelity 6S coefficients supplied as additional inputs, which is a different and easier problem.

The configurations are these. Advection is the 200-point input/output pair of the operator suite, used at full resolution with no reduction. Burgers is PDEBench’s `1D_Burgers_Sols_Nu/` release, the map from the initial condition to the solution at $t = 1$ on a grid subsampled by four to 256 cells, the input used raw and the output reduced to 128 principal coefficients. The three Well datasets are one-step maps of the full multi-channel state with 256 principal coefficients in and out: the turbulent radiative layer on its 128×384 grid (density, pressure, velocity), active matter on 256×256 (concentration, velocity, and the two tensor fields), and the Helmholtz staircase on its two pressure channels. ClimSim is the low-resolution release, 124 inputs to 128 tendencies, with 100,000 training rows drawn from the 10.1M available and validation and test taken from the release’s own validation and scoring files. The correction-coefficient corpus is the paired 6S/libRadtran Sentinel-2 set [13, 14], seven state variables and band wavelength mapping to the three libRadtran coefficients, split by atmospheric state so that all bands of a state travel together. The inputs are wavelength, solar zenith, view zenith, relative azimuth, aerosol optical depth, water vapor, ozone and elevation. The multi-fidelity configuration adds the three 6S coefficients, giving input dimensions eight and eleven, respectively. For these two configurations

the error is

$$E_{0.05} = \frac{1}{N} \sum_{i=1}^N \frac{\|Y_i - \hat{Y}_i\|_2}{\max\{\|Y_i\|_2, 0.05\}}. \quad (22)$$

The same denominator floor enters validation selection and the kernel-flow objective. It limits the weight of small-norm coefficient vectors; these two rows therefore use a floored relative error rather than an unmodified relative norm.

The families are those of Section 4 with two differences forced by scale. The reference family is a ridge regression linear in whatever inputs the models receive, not the cubic of Table 1: a cubic expansion of a 256-dimensional input is not tractable, and the linear reference is used consistently for the additional configurations. And the kernel solves are exact only when the training block has at most 20,000 rows; above that they use a Nyström approximation [59] on 6,000 uniformly drawn landmarks. That cap binds on ClimSim, on the Helmholtz staircase at 20,384 rows, and on the two correction-coefficient configurations. Those entries are what a capped fit gives; they are approximate rather than full dense solves. Their errors need not bound those of the full kernel, since approximation affects both fit and regularization. The convex stack and the per-coordinate selection are formed from the kernels, the networks, the residual correction and the feature kernel; the ridge is a reference row and is not one of their members.

What each row averages differs and is stated in the table. Advection, Burgers, ClimSim and the two correction-coefficient configurations have three to five random splits. The Well’s releases carry their own train, validation and test files, which are used unchanged, so the three turbulent-radiative-layer runs share one split and differ only in network initialisation; their spread, of order 0.01 points, measures initialisation and not sampling. The active-matter and Helmholtz-staircase rows are single runs.

The relative errors vary substantially with the problem. Where the map is smooth and the sample small — Burgers at $\nu = 0.1$ — the exact kernel is at 2.81% against the network’s 20.34%, a factor of seven, of the same kind as the five- to six-fold gap between the per-input kernel and the network on the EMIT table’s smooth components; as the viscosity falls and the solutions roughen the gap closes to 23.40% against 36.65% and then to 31.66% against 39.98%. Where the sample is large and the map rough — ClimSim, and the correction-coefficient corpus at nearly half a million rows — the network leads instead, 25.10% against 27.20% and 3.75% against 62.93%, which is the OCO-2 regime. Where the operator is linear, on advection, the linear ridge is the lowest row and every other family lands within 0.4 points of it, consistent with the linear form of the operator.

Learned metrics give particularly large improvements on the correction-coefficient corpus. On the correction-coefficient corpus from the state alone the isotropic kernel is at 62.93 ± 0.51 and a kernel-flow metric [57] takes it to 3.97 ± 0.42 , a factor of sixteen in the mean, which brings a kernel on the state to within six percent of the network on a corpus where the plain kernel was seventeen times worse. With the low-fidelity coefficients supplied as inputs the same pair reads 15.76 and 4.60, a factor of three: the easier the problem is made for the plain kernel, the less the learned metric adds. This comparison, like the EMIT and OCO-2 results, shows that the choice of input metric can substantially change a kernel baseline.

The convex stack is at or near the best member of its candidate pool on each configuration. The convex stack is the lowest row on eight of the ten configurations and the second lowest on the other two; the residual-corrected network is in the lowest two on five. Both exceptions are the linear problems, and on both the row that beats the stack is the ridge, which the stack

does not contain. On advection the ridge is lower by 0.07 points while the stack, at 11.36%, matches the best member it has — the isotropic kernel, on which it puts weight 0.87 — to two decimals. On the Helmholtz staircase the ridge reaches 0.004%, sixteen times below the next row of the table and about a hundred times below the weakest member of the pool, and the stack reaches 0.066%, again a shade below its own best member, the learned-metric kernel at 0.067%, on which it puts weight 0.94. In both cases the stack does what it is asked to do inside its pool; what it cannot do is recover a family the pool excludes. The statement this supports is correspondingly narrow: a combination is a good default when no member of its own pool dominates it, and a combination can only be read against the whole comparison when the pool contains the families the comparison includes.

The comparisons have several limitations. These rows use one training budget per corpus and one architecture family, and several of these corpora have published results from methods tuned for them specifically; the comparison here is between families under a common protocol, not against the state of the art on each problem. With three to five runs, and one on two rows, the spreads support the large orderings above and not fine distinctions between adjacent families. And every kernel column here inherits the tuning protocol Section 6.1 measures: the length scale is chosen on a 6,000-row subsample and the winner refit on every row, and this difference in sample size can affect the selected hyperparameters. On the EMIT table that costs the isotropic kernel a factor 0.882; on ClimSim, where the subsample is a sixteenth of the solve, re-choosing the scale against the solve that actually runs halves it and lowers the kernel’s error by 3.6%; on a 20,000-row corpus it changes nothing. These examples show that tuning can affect the comparison. They do not imply that every reported kernel error is pessimistic or that a reported gap is an upper bound: validation retuning can also increase test error. The columns of Table 19 describe the specified training and selection procedures, not the optimum attainable by either family.

15 Discussion

Exact kernel regression is effective on this six-dimensional radiative-transfer table. Input-specific length scales substantially reduce radiance error relative to the isotropic kernel, cubic regression and width-512 network. Residual fitting improves the network under the implemented PCA projection. The width-2000 network has larger standalone error under the tested training configuration, yet its features support more accurate kernel predictions at all ten splits. Since width and training duration change together, this comparison concerns the fitted representations rather than an intrinsic effect of width.

The central result is the difference between forward and inverse rankings. The input-scaled kernel has the smallest median reflectance error in the main comparison, whereas the feature kernels have smaller 95th-percentile inverse errors. Both narrow and wide convex stacks increase those tails at every split relative to their feature kernels. For the wide model, the mean radiance gain is only 0.0006 percentage points, while the mean within-split 95th-percentile reflectance error increases from 7.57 to 22.71 points. A ranking based only on mean radiance would obscure this trade-off. The finite-perturbation identities show that inversion also requires control of the predicted denominator relative to component errors. Retrieval validation should specify a physical passband, report retained coverage and invalid inversions, and evaluate a range of surface reflectances.

The residual-error identity yields an exact alignment criterion for improvement in a quadratic

metric. The PCA decomposition adds the truncation error explicitly. Applying a kernel learning-rate bound to the neural residual requires a uniform bound and a separate source-norm estimate; the observed learning slopes describe only the measured sample-size range. The OCO-2 and additional-corpus comparisons likewise support task-dependent choices of representation and kernel metric, rather than a single preferred model family.

Theorem 6 turns the forward-inverse distinction into a quantitative statement. The reflectance error of the constrained retrieval is at most the transmission-weighted component error divided by the transmission, with constant one, and the mean-square error over an evaluation distribution is bounded by the tail mass of the transmission plus the mean of that ratio where the transmission is not small. On the EMIT table the absolute component errors of every head are not concentrated in the opaque bands; the retrieval tails come from dividing ordinary errors by a vanishing transmission, and the observed profile exponents are far from the value at which Theorem 8 gives a rate without loss. Two consequences follow for retrieval validation. On the all-band distribution the reflectance tails measure the opaque bands and rank the heads by where their small errors happen to sit; on a passband with transmission bounded below, the component rate transfers unchanged, and the heads are compared there. Proposition 9 adds the other half of the picture: a Nyström approximation [59, 64], a compressed representation or a changed ridge moves the corrected predictor through its action on the fitted residual, so the bound to check before trusting an approximation is the residual-action bound, which holds at every held-out row in the experiments and is finite where a bound in the norm of the Gram perturbation is not.

The replication at further splits, the ridge path and the two external comparisons sharpen the picture in three ways. The concatenated-feature kernel’s gain over the single-member kernel survives four new splits in the mean and, at these splits, in the reflectance tail as well, so the tail cost belongs to the stacking step and not to the wider representation. The correction degrades gracefully where its kernel step does not: along the ridge the kernel alone loses an order of magnitude before the corrected predictor loses a factor of two, because the network mean carries the prediction once the kernel is over-smoothed, which is the content of Proposition 4 read on data. The kernel on the state inputs is not the limiting choice either: across ten of them the best is within two hundredths of a point of the kernel on a modest network’s features, and what moves the error is combination, the convex stack of all heads standing roughly a quarter below the best single head on three of the four components; regularising the network changes its error little, and both kernel steps compress the differences between its arms by an order of magnitude, the correction doing the work the regulariser could not. And the metric decides the model on the two benchmarks in the same way it decides it here: on the correction-coefficient corpus the residual-corrected network improves on the published values in all four metrics on both splits, with the widest margin out of distribution, while on OCO-2 the heads that win in the reduced coordinates lose by a factor of two to three in radiance and vice versa, and the rule that aligns the combination with the radiance metric buys its improvement entirely in coordinates that carry little radiance.

Accuracy must also be considered alongside cost. Dense kernel regression stores order n^2 entries and requires order n^3 operations for factorization. One $18,884 \times 18,884$ double-precision matrix occupies approximately 2.85 GB before factorization workspaces and other arrays. Feature kernels evaluate both the network and the kernel against the training design; stacks evaluate each active member. Per-model throughput was not measured under matched conditions, so these results do not establish a speed advantage over operational emulators such as sRTMnet

[1].

The ten partitions quantify within-table split sensitivity. Model development used the same table, and the partitions do not provide an untouched confirmation set. The two principal EMIT comparisons have complete per-split summaries; individual records are incomplete for some secondary ablations. Access to the raw arrays and simulator configuration is required to reproduce the training campaign. Within these limits, the results show that neural features and exact kernels can complement one another, while a small forward-error gain can coincide with a substantially worse inverse-error tail.

Acknowledgements

The research presented in this paper was supported by the European Research Council (ERC) under the European Union’s Horizon Europe research and innovation programme (grant agreement No. 101041711), by the Simons Foundation as part of the Collaboration on the Mathematical and Scientific Foundations of Deep Learning, by Heights Labs, by Convex Nexus Capital, by the Israel Science Foundation (grant number 2258/19), by the Israel Science Foundation (ISF Grant 4101/25), and by Google Cloud through its education credits program (credit 530550103).

Declaration of competing interest

The author holds equity in Heights Labs and in Convex Nexus Capital and has a management role at Convex Nexus Capital; both entities supported this work and are named above. These entities and the public funders had no role in the study design, analysis, manuscript preparation, or decision to publish.

AI-assisted contributions

Claude Code (Anthropic) and Codex (OpenAI) implemented and ran the data pipeline, model comparisons, scaling experiments, diagnostics and figure generation, and drafted the manuscript from the author’s data and the methodology of [11]. ChatGPT (OpenAI) developed the residual-error and inversion-stability arguments, checked the finite identities, analysed the stored per-split results, implemented the conditioning diagnostics and verification tests, and edited the manuscript. Claude Code implemented and ran the later experiments of Sections 9–12: the transmission-conditioned retrieval diagnostics, the ten-kernel comparison and its combiners, the regularisation suite, the ridge-path and perturbation measurements, the synthetic reference, the matched benchmark rescoring, and the OCO-2 selection through the Gram matrix of its design; it also generated every table of those sections from the stored run records and wrote their text. The author supplied the data, set the questions and reviewed the results.

Data and code availability

The manuscript, model drivers, run summaries and table-generation scripts are available at <https://github.com/yspennstate/emit-oco2-kernel-emulators>.

The width-512 and width-2000 EMIT experiments each contain ten per-split records with array digests, selected hyperparameters, split sizes and evaluation metrics. These twenty records

reproduce Tables 1, 7, 2 and 3. Reproduction commands and the file inventory are given in the repository documentation.

Secondary learning-curve, width and tuning summaries contain ten-run aggregates. The public individual archive lacks one isotropic and one input-scaled record at 8,000 rows; it contains three width-1000 records and five full-data-tuning records. The three-member wide-feature comparison uses three splits, and width 4,000 uses two. The corresponding aggregate tables and individual records are distinguished in the reproduction manifest.

The EMIT arrays, identified as `data_EMIT_24k` and attributed to the Jet Propulsion Laboratory, are not redistributed. Raw arrays, trained weights and prediction files are available from the author on request. The public archive does not specify the simulator version, generation configuration or redistribution permissions. Scoring uses float64; prediction archives stored as float32 cannot reproduce cancellation-sensitive inverse errors at the same precision. The conditioning utility requires explicit approval of precision-altered inputs. Its tests use synthetic arrays; the EMIT results reported here are all-band statistics rather than conditioned evaluations.

The OCO-2 arrays and reference predictions are from the release accompanying [3]. The additional configurations use the cited public releases.

References

- [1] P. G. Brodrick, D. R. Thompson, J. E. Fahlen, M. L. Eastwood, C. M. Sarture, S. R. Lundeen, W. Olson-Duvall, N. Carmon, and R. O. Green. Generalized radiative transfer emulation for imaging spectroscopy reflectance retrievals. *Remote Sensing of Environment*, 261:112476, 2021.
- [2] D. R. Thompson, V. Natraj, R. O. Green, M. C. Helmlinger, B.-C. Gao, and M. L. Eastwood. Optimal estimation for imaging spectrometer atmospheric correction. *Remote Sensing of Environment*, 216:355–373, 2018.
- [3] O. Lamminpää, J. Susiluoto, J. Hobbs, J. McDuffie, A. Braverman, and H. Owhadi. Forward model emulator for atmospheric radiative transfer using Gaussian processes and cross validation. *Atmospheric Measurement Techniques*, 18:673–694, 2025.
- [4] M. A. A. Mazid and N. Rishe. Multi-fidelity emulation of atmospheric correction coefficients with physics-guided Kolmogorov–Arnold networks. *Remote Sensing*, 18(11):1826, 2026.
- [5] P. Batlle, M. Darcy, B. Hosseini, and H. Owhadi. Kernel methods are competitive for operator learning. *Journal of Computational Physics*, 496:112549, 2024.
- [6] L. Lu, P. Jin, G. Pang, Z. Zhang, and G. E. Karniadakis. Learning nonlinear operators via DeepONet based on the universal approximation theorem of operators. *Nature Machine Intelligence*, 3:218–229, 2021.
- [7] Z. Li, N. Kovachki, K. Azizzadenesheli, B. Liu, K. Bhattacharya, A. Stuart, and A. Anandkumar. Fourier neural operator for parametric partial differential equations. In *International Conference on Learning Representations*, 2021.
- [8] C. E. Rasmussen and C. K. I. Williams. *Gaussian Processes for Machine Learning*. MIT Press, 2006.

- [9] A. G. Wilson, Z. Hu, R. Salakhutdinov, and E. P. Xing. Deep kernel learning. In *Proceedings of the 19th International Conference on Artificial Intelligence and Statistics*, pages 370–378, 2016.
- [10] G. R. Yoo and H. Owhadi. Deep regularization and direct training of the inner layers of neural networks with kernel flows. *Physica D: Nonlinear Phenomena*, 426:132952, 2021.
- [11] Y. Shmalo. Neural means and kernel corrections for operator learning. Preprint, ResearchGate, 2026. Code and per-run summaries at <https://github.com/yspennstate/neural-means-kernel-corrections>.
- [12] R. O. Green et al. The Earth surface mineral dust source investigation: an Earth science imaging spectroscopy mission. In *IEEE Aerospace Conference*, 2020.
- [13] E. F. Vermote, D. Tanré, J. L. Deuzé, M. Herman, and J.-J. Morcrette. Second simulation of the satellite signal in the solar spectrum, 6S: an overview. *IEEE Transactions on Geoscience and Remote Sensing*, 35(3):675–686, 1997.
- [14] B. Mayer and A. Kylling. Technical note: the libRadtran software package for radiative transfer calculations — description and examples of use. *Atmospheric Chemistry and Physics*, 5:1855–1877, 2005.
- [15] C. Emde, R. Buras-Schnell, A. Kylling, B. Mayer, J. Gasteiger, U. Hamann, J. Kylling, B. Richter, C. Pause, T. Dowling, and L. Bugliaro. The libRadtran software package for radiative transfer calculations (version 2.0.1). *Geoscientific Model Development*, 9:1647–1672, 2016.
- [16] A. Berk, P. Conforti, R. Kennett, T. Perkins, F. Hawes, and J. van den Bosch. MODTRAN6: a major upgrade of the MODTRAN radiative transfer code. In *Proceedings of SPIE 9088, Algorithms and Technologies for Multispectral, Hyperspectral, and Ultraspectral Imagery XX*, 2014.
- [17] C. D. Rodgers. *Inverse Methods for Atmospheric Sounding: Theory and Practice*. World Scientific, 2000.
- [18] A. Eldering, C. W. O’Dell, P. O. Wennberg, D. Crisp, M. R. Gunson, et al. The Orbiting Carbon Observatory-2: first 18 months of science data products. *Atmospheric Measurement Techniques*, 10:549–563, 2017.
- [19] D. Crisp, H. R. Pollock, R. Rosenberg, L. Chapsky, R. A. M. Lee, et al. The on-orbit performance of the Orbiting Carbon Observatory-2 instrument and its radiometrically calibrated products. *Atmospheric Measurement Techniques*, 10:59–81, 2017.
- [20] B. J. Connor, H. Boesch, G. Toon, B. Sen, C. Miller, and D. Crisp. Orbiting Carbon Observatory: inverse method and prospective error analysis. *Journal of Geophysical Research: Atmospheres*, 113:D05305, 2008.
- [21] C. W. O’Dell, B. Connor, H. Bosch, D. O’Brien, C. Frankenberg, et al. The ACOS CO₂ retrieval algorithm — part 1: description and validation against synthetic observations. *Atmospheric Measurement Techniques*, 5:99–121, 2012.

- [22] F. Chevallier, F. Chérut, N. A. Scott, and A. Chédin. A neural network approach for a fast and accurate computation of a longwave radiative budget. *Journal of Applied Meteorology*, 37:1385–1397, 1998.
- [23] V. M. Krasnopolsky, M. S. Fox-Rabinovitz, and D. V. Chalikov. New approach to calculation of atmospheric model physics: accurate and fast neural network emulation of longwave radiation in a climate model. *Monthly Weather Review*, 133:1370–1383, 2005.
- [24] M. A. Veerman, R. Pincus, R. Stoffer, C. M. van Leeuwen, D. Podareanu, and C. C. van Heerwaarden. Predicting atmospheric optical properties for radiative transfer computations using neural networks. *Philosophical Transactions of the Royal Society A*, 379:20200095, 2021.
- [25] P. Ukkonen. Exploring pathways to more accurate machine learning emulation of atmospheric radiative transfer. *Journal of Advances in Modeling Earth Systems*, 14:e2021MS002875, 2022.
- [26] S. Rasp, M. S. Pritchard, and P. Gentine. Deep learning to represent subgrid processes in climate models. *Proceedings of the National Academy of Sciences*, 115(39):9684–9689, 2018.
- [27] N. D. Brenowitz and C. S. Bretherton. Prognostic validation of a neural network unified physics parameterization. *Geophysical Research Letters*, 45:6289–6298, 2018.
- [28] J. P. Rivera, J. Verrelst, J. Gómez-Dans, J. Muñoz-Marí, J. Moreno, and G. Camps-Valls. An emulator toolbox to approximate radiative transfer models with statistical learning. *Remote Sensing*, 7(7):9347–9370, 2015.
- [29] J. Verrelst, N. Sabater, J. P. Rivera, J. Muñoz-Marí, J. Vicent, G. Camps-Valls, and J. Moreno. Emulation of leaf, canopy and atmosphere radiative transfer models for fast global sensitivity analysis. *Remote Sensing*, 8(8):673, 2016.
- [30] J. L. Gómez-Dans, P. E. Lewis, and M. Disney. Efficient emulation of radiative transfer codes using Gaussian processes and application to land surface parameter inferences. *Remote Sensing*, 8(2):119, 2016.
- [31] G. Camps-Valls, J. Verrelst, J. Muñoz-Marí, V. Laparra, F. Mateo-Jiménez, and J. Gómez-Dans. A survey on Gaussian processes for earth-observation data analysis. *IEEE Geoscience and Remote Sensing Magazine*, 4(2):58–78, 2016.
- [32] J. Sacks, W. J. Welch, T. J. Mitchell, and H. P. Wynn. Design and analysis of computer experiments. *Statistical Science*, 4(4):409–423, 1989.
- [33] M. C. Kennedy and A. O’Hagan. Predicting the output from a complex computer code when fast approximations are available. *Biometrika*, 87(1):1–13, 2000.
- [34] M. C. Kennedy and A. O’Hagan. Bayesian calibration of computer models. *Journal of the Royal Statistical Society: Series B*, 63(3):425–464, 2001.
- [35] P. Perdikaris, M. Raissi, A. Damianou, N. D. Lawrence, and G. E. Karniadakis. Nonlinear information fusion algorithms for data-efficient multi-fidelity modelling. *Proceedings of the Royal Society A*, 473:20160751, 2017.

- [36] B. Peherstorfer, K. Willcox, and M. Gunzburger. Survey of multifidelity methods in uncertainty propagation, inference, and optimization. *SIAM Review*, 60(3):550–591, 2018.
- [37] K. Cawse-Nicholson, P. A. Townsend, D. Schimel, A. M. Assiri, P. L. Blake, et al. NASA surface biology and geology designated observable: a perspective on surface imaging algorithms. *Remote Sensing of Environment*, 257:112349, 2021.
- [38] N. Aronszajn. Theory of reproducing kernels. *Transactions of the American Mathematical Society*, 68:337–404, 1950.
- [39] G. Kimeldorf and G. Wahba. Some results on Tchebycheffian spline functions. *Journal of Mathematical Analysis and Applications*, 33(1):82–95, 1971.
- [40] G. Wahba. *Spline Models for Observational Data*. SIAM, 1990.
- [41] B. Schölkopf, R. Herbrich, and A. J. Smola. A generalized representer theorem. In *Proceedings of the 14th Annual Conference on Computational Learning Theory*, pages 416–426, 2001.
- [42] A. E. Hoerl and R. W. Kennard. Ridge regression: biased estimation for nonorthogonal problems. *Technometrics*, 12(1):55–67, 1970.
- [43] H. W. Engl, M. Hanke, and A. Neubauer. *Regularization of Inverse Problems*. Kluwer, 1996.
- [44] M. L. Stein. *Interpolation of Spatial Data: Some Theory for Kriging*. Springer, 1999.
- [45] A. Caponnetto and E. De Vito. Optimal rates for the regularized least-squares algorithm. *Foundations of Computational Mathematics*, 7(3):331–368, 2007.
- [46] D. J. Sutherland. Fixing an error in Caponnetto and de Vito (2007). arXiv:1702.02982v2, 2021.
- [47] S. Fischer and I. Steinwart. Sobolev norm learning rates for regularized least-squares algorithms. *Journal of Machine Learning Research*, 21(205):1–38, 2020.
- [48] S. Smale and D.-X. Zhou. Learning theory estimates via integral operators and their approximations. *Constructive Approximation*, 26:153–172, 2007.
- [49] B. Schölkopf and A. J. Smola. *Learning with Kernels*. MIT Press, 2002.
- [50] I. Steinwart and A. Christmann. *Support Vector Machines*. Springer, 2008.
- [51] T. Hastie, R. Tibshirani, and J. Friedman. *The Elements of Statistical Learning*. Springer, second edition, 2009.
- [52] C. A. Micchelli and M. Pontil. On learning vector-valued functions. *Neural Computation*, 17(1):177–204, 2005.
- [53] M. A. Álvarez, L. Rosasco, and N. D. Lawrence. Kernels for vector-valued functions: a review. *Foundations and Trends in Machine Learning*, 4(3):195–266, 2012.

- [54] H. Owhadi. Bayesian numerical homogenization. *Multiscale Modeling & Simulation*, 13(3):812–828, 2015.
- [55] H. Owhadi and C. Scovel. *Operator-Adapted Wavelets, Fast Solvers, and Numerical Homogenization*. Cambridge University Press, 2019.
- [56] Y. Chen, B. Hosseini, H. Owhadi, and A. M. Stuart. Solving and learning nonlinear PDEs with Gaussian processes. *Journal of Computational Physics*, 447:110668, 2021.
- [57] H. Owhadi and G. R. Yoo. Kernel flows: from learning kernels from data into the abyss. *Journal of Computational Physics*, 389:22–47, 2019.
- [58] B. Hamzi and H. Owhadi. Learning dynamical systems from data: a simple cross-validation perspective, part I: parametric kernel flows. *Physica D: Nonlinear Phenomena*, 421:132817, 2021.
- [59] C. K. I. Williams and M. Seeger. Using the Nyström method to speed up kernel machines. In *Advances in Neural Information Processing Systems 13*, pages 682–688, 2001.
- [60] P. Drineas and M. W. Mahoney. On the Nyström method for approximating a Gram matrix for improved kernel-based learning. *Journal of Machine Learning Research*, 6:2153–2175, 2005.
- [61] F. Bach. Sharp analysis of low-rank kernel matrix approximations. In *Proceedings of the 26th Annual Conference on Learning Theory*, pages 185–209, 2013.
- [62] A. Rahimi and B. Recht. Random features for large-scale kernel machines. In *Advances in Neural Information Processing Systems 20*, pages 1177–1184, 2008.
- [63] N. H. Nelsen and A. M. Stuart. The random feature model for input-output maps between Banach spaces. *SIAM Journal on Scientific Computing*, 43(5):A3212–A3243, 2021.
- [64] A. Rudi, L. Carratino, and L. Rosasco. FALKON: an optimal large scale kernel method. In *Advances in Neural Information Processing Systems 30*, 2017.
- [65] G. Meanti, L. Carratino, L. Rosasco, and A. Rudi. Kernel methods through the roof: handling billions of points efficiently. In *Advances in Neural Information Processing Systems 33*, 2020.
- [66] J. Quiñonero-Candela and C. E. Rasmussen. A unifying view of sparse approximate Gaussian process regression. *Journal of Machine Learning Research*, 6:1939–1959, 2005.
- [67] M. K. Titsias. Variational learning of inducing variables in sparse Gaussian processes. In *Proceedings of the 12th International Conference on Artificial Intelligence and Statistics*, pages 567–574, 2009.
- [68] J. Hensman, N. Fusi, and N. D. Lawrence. Gaussian processes for big data. In *Proceedings of the 29th Conference on Uncertainty in Artificial Intelligence*, pages 282–290, 2013.
- [69] A. G. Wilson and H. Nickisch. Kernel interpolation for scalable structured Gaussian processes. In *Proceedings of the 32nd International Conference on Machine Learning*, pages 1775–1784, 2015.

- [70] J. R. Gardner, G. Pleiss, D. Bindel, K. Q. Weinberger, and A. G. Wilson. GPyTorch: blackbox matrix-matrix Gaussian process inference with GPU acceleration. In *Advances in Neural Information Processing Systems 31*, 2018.
- [71] A. G. Wilson and R. P. Adams. Gaussian process kernels for pattern discovery and extrapolation. In *Proceedings of the 30th International Conference on Machine Learning*, pages 1067–1075, 2013.
- [72] D. Duvenaud, H. Nickisch, and C. E. Rasmussen. Additive Gaussian processes. In *Advances in Neural Information Processing Systems 24*, 2011.
- [73] D. Duvenaud, J. R. Lloyd, R. Grosse, J. B. Tenenbaum, and Z. Ghahramani. Structure discovery in nonparametric regression through compositional kernel search. In *Proceedings of the 30th International Conference on Machine Learning*, pages 1166–1174, 2013.
- [74] G. R. G. Lanckriet, N. Cristianini, P. Bartlett, L. El Ghaoui, and M. I. Jordan. Learning the kernel matrix with semidefinite programming. *Journal of Machine Learning Research*, 5:27–72, 2004.
- [75] M. Gönen and E. Alpaydm. Multiple kernel learning algorithms. *Journal of Machine Learning Research*, 12:2211–2268, 2011.
- [76] N. Cristianini, J. Shawe-Taylor, A. Elisseeff, and J. Kandola. On kernel-target alignment. In *Advances in Neural Information Processing Systems 14*, pages 367–373, 2002.
- [77] C. Cortes, M. Mohri, and A. Rostamizadeh. Algorithms for learning kernels based on centered alignment. *Journal of Machine Learning Research*, 13:795–828, 2012.
- [78] J. M. Bates and C. W. J. Granger. The combination of forecasts. *Operational Research Quarterly*, 20(4):451–468, 1969.
- [79] D. H. Wolpert. Stacked generalization. *Neural Networks*, 5(2):241–259, 1992.
- [80] L. Breiman. Stacked regressions. *Machine Learning*, 24:49–64, 1996.
- [81] M. Stone. Cross-validated choice and assessment of statistical predictions. *Journal of the Royal Statistical Society: Series B*, 36(2):111–147, 1974.
- [82] G. H. Golub, M. Heath, and G. Wahba. Generalized cross-validation as a method for choosing a good ridge parameter. *Technometrics*, 21(2):215–223, 1979.
- [83] S. Arlot and A. Celisse. A survey of cross-validation procedures for model selection. *Statistics Surveys*, 4:40–79, 2010.
- [84] K. Hornik, M. Stinchcombe, and H. White. Multilayer feedforward networks are universal approximators. *Neural Networks*, 2(5):359–366, 1989.
- [85] D. P. Kingma and J. Ba. Adam: a method for stochastic optimization. In *International Conference on Learning Representations*, 2015.
- [86] I. Loshchilov and F. Hutter. Decoupled weight decay regularization. In *International Conference on Learning Representations*, 2019.

- [87] L. Prechelt. Early stopping — but when? In *Neural Networks: Tricks of the Trade*, volume 1524 of *Lecture Notes in Computer Science*, pages 55–69. Springer, 1998.
- [88] N. Srivastava, G. Hinton, A. Krizhevsky, I. Sutskever, and R. Salakhutdinov. Dropout: a simple way to prevent neural networks from overfitting. *Journal of Machine Learning Research*, 15:1929–1958, 2014.
- [89] R. M. Neal. *Bayesian Learning for Neural Networks*. Springer, 1996.
- [90] J. Lee, Y. Bahri, R. Novak, S. S. Schoenholz, J. Pennington, and J. Sohl-Dickstein. Deep neural networks as Gaussian processes. In *International Conference on Learning Representations*, 2018.
- [91] A. G. de G. Matthews, M. Rowland, J. Hron, R. E. Turner, and Z. Ghahramani. Gaussian process behaviour in wide deep neural networks. In *International Conference on Learning Representations*, 2018.
- [92] Y. Cho and L. K. Saul. Kernel methods for deep learning. In *Advances in Neural Information Processing Systems 22*, pages 342–350, 2009.
- [93] A. Jacot, F. Gabriel, and C. Hongler. Neural tangent kernel: convergence and generalization in neural networks. In *Advances in Neural Information Processing Systems 31*, 2018.
- [94] S. Arora, S. S. Du, W. Hu, Z. Li, R. Salakhutdinov, and R. Wang. On exact computation with an infinitely wide neural net. In *Advances in Neural Information Processing Systems 32*, 2019.
- [95] M. Belkin, S. Ma, and S. Mandal. To understand deep learning we need to understand kernel learning. In *Proceedings of the 35th International Conference on Machine Learning*, pages 541–549, 2018.
- [96] A. Radhakrishnan, D. Beaglehole, P. Pandit, and M. Belkin. Mechanism for feature learning in neural networks and backpropagation-free machine learning models. *Science*, 383(6690):1461–1467, 2024.
- [97] A. N. Kolmogorov. On the representation of continuous functions of several variables by superpositions of continuous functions of one variable and addition. *Doklady Akademii Nauk SSSR*, 114:953–956, 1957.
- [98] Z. Liu, Y. Wang, S. Vaidya, F. Ruehle, J. Halverson, M. Soljačić, T. Y. Hou, and M. Tegmark. KAN: Kolmogorov–Arnold networks. arXiv:2404.19756, 2024.
- [99] N. Kovachki, Z. Li, B. Liu, K. Azizzadenesheli, K. Bhattacharya, A. Stuart, and A. Anandkumar. Neural operator: learning maps between function spaces with applications to PDEs. *Journal of Machine Learning Research*, 24(89):1–97, 2023.
- [100] L. Lu, X. Meng, S. Cai, Z. Mao, S. Goswami, Z. Zhang, and G. E. Karniadakis. A comprehensive and fair comparison of two neural operators (with practical extensions) based on FAIR data. *Computer Methods in Applied Mechanics and Engineering*, 393:114778, 2022.

- [101] M. Raissi, P. Perdikaris, and G. E. Karniadakis. Physics-informed neural networks: a deep learning framework for solving forward and inverse problems involving nonlinear partial differential equations. *Journal of Computational Physics*, 378:686–707, 2019.
- [102] M. V. de Hoop, D. Z. Huang, E. Qian, and A. M. Stuart. The cost-accuracy trade-off in operator learning with neural networks. *Journal of Machine Learning*, 1(3):299–341, 2022.
- [103] M. Takamoto, T. Praditia, R. Leiteritz, D. MacKinlay, F. Alesiani, D. Pflüger, and M. Niepert. PDEBench: an extensive benchmark for scientific machine learning. In *Advances in Neural Information Processing Systems 35*, 2022.
- [104] R. Ohana, M. McCabe, L. Meyer, R. Morel, F. Agocs, M. Beneitez, M. Berger, B. Burkhart, S. Dalziel, D. Fielding, et al. The Well: a large-scale collection of diverse physics simulations for machine learning. In *Advances in Neural Information Processing Systems 37*, pages 44989–45037, 2024.
- [105] S. Yu, W. M. Hannah, L. Peng, J. Lin, M. A. Bhouiri, R. Gupta, B. Lütjens, J. C. Will, G. Behrens, et al. ClimSim: a large multi-scale dataset for hybrid physics-ML climate emulation. In *Advances in Neural Information Processing Systems 36*, 2023.
- [106] G. H. Golub and C. F. Van Loan. *Matrix Computations*. Johns Hopkins University Press, fourth edition, 2013.
- [107] W. Hoeffding. Probability inequalities for sums of bounded random variables. *Journal of the American Statistical Association*, 58(301):13–30, 1963.
- [108] G. W. Stewart and J.-G. Sun. *Matrix Perturbation Theory*. Academic Press, 1990.
- [109] N. J. Higham. *Accuracy and Stability of Numerical Algorithms*. SIAM, second edition, 2002.

kernel	perturbed system	$\lambda_{\min}(B)$	$\ B\ $	median movement	residual-action / movement	signed-spectrum / movement	$\delta < 1$ bound / movement	ridge bound / movement
state inputs	nystrom_m50	-0.941	0.941	2.01	57.3	4.05e+03	4.05e+03	-
state inputs	nystrom_m100	-0.872	0.872	1.05	82.1	3.05e+03	3.05e+03	-
state inputs	nystrom_m200	-0.642	0.642	0.498	136	1.65e+03	1.65e+03	-
state inputs	nystrom_m400	-0.411	0.411	0.2	178	1.54e+03	1.54e+03	-
state inputs	nystrom_m800	-0.193	0.193	0.0702	264	1.47e+03	1.47e+03	-
state inputs	nystrom_m1600	-0.078	0.078	0.0103	408	3.51e+03	3.51e+03	-
state inputs	gamma_x0.5	-0.500	0.500	0.304	517	1.34e+03	1.34e+03	-
state inputs	gamma_x2.0	+0.000	0.999	0.372	835	550	1.32e+06	550
state inputs	gamma_x10.0	+0.002	8.993	1.36	1.98e+03	271	-	271
state inputs	gamma_x100.0	+0.018	98.918	3.67	7.38e+03	111	-	111
state inputs	pca_inputs_d2	-0.998	3.600	10.5	14.3	4.81e+04	-	-
state inputs	pca_inputs_d3	-0.998	2.523	9.8	14.2	4.49e+04	-	-
state inputs	pca_inputs_d4	-0.998	1.581	5.24	25.3	6.52e+04	-	-
state inputs	pca_inputs_d5	-0.997	0.997	1.9	64.3	1.39e+05	1.39e+05	-
features Y1	pca_features_d1000	-0.192	0.192	0.00383	4.01e+03	3.8e+04	3.8e+04	-
features Y1	pca_features_d500	-0.503	0.503	0.0107	4.41e+03	5.82e+04	5.82e+04	-
features Y1	pca_features_d250	-0.764	0.764	0.0185	4.3e+03	1.1e+05	1.1e+05	-
features Y1	pca_features_d125	-0.917	0.917	0.0279	3.88e+03	2.6e+05	2.6e+05	-
features Y1	pca_features_d64	-0.965	0.965	0.0407	3.55e+03	4.78e+05	4.78e+05	-
features Y1	pca_features_d32	-0.991	0.991	0.0595	3.13e+03	1.59e+06	1.59e+06	-
features Y1	pca_features_d16	-0.999	0.999	0.0923	2.48e+03	1.3e+07	1.3e+07	-
features Y1	nystrom_m100	-1.000	1.000	0.465	637	8.85e+07	8.85e+07	-
features Y1	nystrom_m400	-1.000	1.000	0.223	1.27e+03	7.11e+07	7.11e+07	-
features Y1	nystrom_m1600	-1.000	1.000	0.0746	2.58e+03	2.74e+07	2.74e+07	-
features Y1	nystrom_m3200	-0.997	0.997	0.0196	5.22e+03	9.75e+06	9.75e+06	-
features Y1	gamma_x0.5	-0.490	0.490	0.0102	9.61e+03	5.74e+04	5.74e+04	-
features Y1	gamma_x2.0	+0.000	0.979	0.0115	1.7e+04	2.62e+04	2.5e+06	2.62e+04
features Y1	gamma_x10.0	+0.000	8.813	0.0407	4.32e+04	1.34e+04	-	1.34e+04
features Y1	gamma_x100.0	+0.000	96.942	0.0889	2.17e+05	6.78e+03	-	6.78e+03
features Y2	pca_features_d1000	-0.122	0.122	0.00143	6.41e+03	4.95e+04	4.95e+04	-
features Y2	pca_features_d500	-0.418	0.418	0.00425	6.8e+03	8.59e+04	8.59e+04	-
features Y2	pca_features_d250	-0.726	0.726	0.00802	6.45e+03	1.7e+05	1.7e+05	-
features Y2	pca_features_d125	-0.868	0.868	0.0117	6.33e+03	2.99e+05	2.99e+05	-
features Y2	pca_features_d64	-0.977	0.977	0.0184	5.76e+03	1.3e+06	1.3e+06	-
features Y2	pca_features_d32	-0.989	0.989	0.0267	5.78e+03	2.3e+06	2.3e+06	-
features Y2	pca_features_d16	-0.999	0.999	0.0419	4.83e+03	1.49e+07	1.49e+07	-
features Y2	nystrom_m100	-1.000	1.000	0.345	731	1.01e+08	1.01e+08	-
features Y2	nystrom_m400	-1.000	1.000	0.122	1.97e+03	8.48e+07	8.48e+07	-
features Y2	nystrom_m1600	-1.000	1.000	0.0823	2.38e+03	5.34e+07	5.34e+07	-
features Y2	nystrom_m3200	-0.998	0.998	0.0161	7.05e+03	1.86e+07	1.86e+07	-
features Y2	gamma_x0.5	-0.491	0.491	0.00632	1.71e+04	7.71e+04	7.71e+04	-
features Y2	gamma_x2.0	+0.000	0.982	0.00687	3.14e+04	3.64e+04	3.96e+06	3.64e+04
features Y2	gamma_x10.0	+0.000	8.836	0.0249	7.8e+04	1.82e+04	-	1.82e+04
features Y2	gamma_x100.0	+0.000	97.199	0.0545	3.92e+05	9.17e+03	-	9.17e+03
features Y3	pca_features_d1000	-0.170	0.170	0.00292	6.17e+03	4.56e+04	4.56e+04	-
features Y3	pca_features_d500	-0.475	0.475	0.00766	7.02e+03	7.73e+04	7.73e+04	-
features Y3	pca_features_d250	-0.738	0.738	0.0136	6.44e+03	1.38e+05	1.38e+05	-
features Y3	pca_features_d125	-0.898	0.898	0.0211	5.61e+03	2.87e+05	2.87e+05	-
features Y3	pca_features_d64	-0.964	0.964	0.028	5.45e+03	7.06e+05	7.06e+05	-
features Y3	pca_features_d32	-0.990	0.990	0.0406	4.9e+03	2.09e+06	2.09e+06	-
features Y3	pca_features_d16	-0.999	0.999	0.0725	3.51e+03	1.63e+07	1.63e+07	-
features Y3	nystrom_m100	-1.000	1.000	0.35	876	4.97e+07	4.97e+07	-
features Y3	nystrom_m400	-1.000	1.000	0.191	1.55e+03	4.53e+07	4.53e+07	-
features Y3	nystrom_m1600	-0.999	0.999	0.0709	3.33e+03	1.45e+07	1.45e+07	-
features Y3	nystrom_m3200	-0.997	0.997	0.0331	4.54e+03	6.39e+06	6.39e+06	-
features Y3	gamma_x0.5	-0.492	0.492	0.00865	1.65e+04	7.26e+04	7.26e+04	-
features Y3	gamma_x2.0	+0.000	0.985	0.00949	3e+04	3.38e+04	4.43e+06	3.38e+04
features Y3	gamma_x10.0	+0.000	8.864	0.0357	7.18e+04	1.63e+04	-	1.63e+04
features Y3	gamma_x100.0	+0.000	97.500	0.0821	3.44e+05	7.81e+03	-	7.81e+03
features Y4	pca_features_d1000	-0.006	0.006	0.00138	538	4.64e+03	4.64e+03	-
features Y4	pca_features_d500	-0.032	0.032	0.00463	712	7.83e+03	7.83e+03	-
features Y4	pca_features_d250	-0.114	0.114	0.0104	826	1.35e+04	1.35e+04	-
features Y4	pca_features_d125	-0.269	0.269	0.0191	897	2.12e+04	2.12e+04	-
features Y4	pca_features_d64	-0.594	0.594	0.0286	963	5.8e+04	5.8e+04	-
features Y4	pca_features_d32	-0.741	0.741	0.0436	983	8.06e+04	8.06e+04	-
features Y4	pca_features_d16	-0.902	0.902	0.0769	987	1.65e+05	1.65e+05	-
features Y4	nystrom_m100	-0.998	0.998	0.603	304	1.2e+06	1.2e+06	-
features Y4	nystrom_m400	-0.985	0.985	0.35	500	2.11e+05	2.11e+05	-
features Y4	nystrom_m1600	-0.860	0.860	0.135	916	4.86e+04	4.86e+04	-
features Y4	nystrom_m3200	-0.762	0.762	0.0397	1.71e+03	8.46e+04	8.46e+04	-
features Y4	gamma_x0.5	-0.494	0.494	0.063	3.15e+03	1.61e+04	1.61e+04	-
features Y4	gamma_x2.0	+0.000	0.989	0.0599	6.6e+03	8.63e+03	1.56e+06	8.63e+03
features Y4	gamma_x10.0	+0.000	8.901	0.191	1.86e+04	4.89e+03	-	4.89e+03
features Y4	gamma_x100.0	+0.000	97.908	0.36	1.07e+05	2.86e+03	-	2.86e+03
features Y1	pca_features_d1000	-0.184	0.184	0.00397	4.68e+03	3.85e+04	3.85e+04	-
features Y1	pca_features_d500	-0.471	0.471	0.0104	5.01e+03	5.82e+04	5.82e+04	-
features Y1	pca_features_d250	-0.724	0.724	0.0175	4.8e+03	1.04e+05	1.04e+05	-
features Y1	pca_features_d125	-0.901	0.901	0.0267	4.25e+03	2.45e+05	2.45e+05	-
features Y1	pca_features_d64	-0.959	0.959	0.0373	3.87e+03	4.76e+05	4.76e+05	-
features Y1	pca_features_d32	-0.985	0.985	0.055	3.43e+03	1.07e+06	1.07e+06	-
features Y1	pca_features_d16	-0.997	0.997	0.0893	2.61e+03	3.51e+06	3.51e+06	-
features Y1	concat_all_members	-0.749	76.009	0.0427	2.3e+05	1.99e+05	-	-
features Y1	member1_alone	-0.973	175.948	0.049	2.34e+05	2.24e+06	-	-
features Y1	member2_alone	-0.967	126.717	0.0541	3.81e+05	2.02e+06	-	-
features Y1	nystrom_m100	-1.000	1.000	0.861	379	9.61e+07	9.61e+07	-
features Y1	nystrom_m400	-1.000	1.000	0.247	1.25e+03	6.69e+07	6.69e+07	-
features Y1	nystrom_m1600	-1.000	1.000	0.112	2.01e+03	4.7e+07	4.7e+07	-
features Y1	nystrom_m3200	-1.000	1.000	0.0249	4.43e+03	5.86e+07	5.86e+07	-
features Y1	gamma_x0.5	-0.493	0.493	0.011	1.06e+04	5.96e+04	5.96e+04	-
features Y1	gamma_x2.0	+0.000	0.986	0.0124	1.89e+04	2.7e+04	3.93e+06	2.7e+04
features Y1	gamma_x10.0	+0.000	8.877	0.0445	4.73e+04	1.36e+04	-	1.36e+04
features Y1	gamma_x100.0	+0.000	97.646	0.102	2.27e+05	6.54e+03	-	6.54e+03
features Y2	pca_features_d1000	-0.108	0.108	0.00163	7.72e+03	5.02e+04	5.02e+04	-
features Y2	pca_features_d500	-0.370	0.370	0.00553	8.84e+03	7.23e+04	7.23e+04	-
features Y2	pca_features_d250	-0.668	0.668	0.0107	8.77e+03	1.29e+05	1.29e+05	-
features Y2	pca_features_d125	-0.834	0.834	0.0166	7.71e+03	2.12e+05	2.12e+05	-
features Y2	pca_features_d64	-0.955	0.955	0.0234	6.59e+03	6.62e+05	6.62e+05	-
features Y2	pca_features_d32	-0.983	0.983	0.0326	6.12e+03	1.46e+06	1.46e+06	-
features Y2	pca_features_d16	-0.998	0.998	0.0503	5.04e+03	1.05e+07	1.05e+07	-
features Y2	concat_all_members	-0.734	96.255	0.0279	6.7e+05	3.68e+05	-	-
features Y2	member1_alone	-0.974	200.331	0.0328	8.4e+05	5.59e+06	-	-
features Y2	member2_alone	-0.968	266.674	0.035	9.16e+05	3.4e+06	-	-
features Y2	nystrom_m100	-1.000	1.000	0.713	451	1.85e+08	1.85e+08	-
features Y2	nystrom_m400	-1.000	1.000	0.119	2.5e+03	4.07e+07	4.07e+07	-
features Y2	nystrom_m1600	-0.999	0.999	0.0322	6.59e+03	1.8e+07	1.8e+07	-
features Y2	nystrom_m3200	-0.992	0.992	0.0104	1.25e+04	8.1e+06	8.1e+06	-

component	kernel	γ	$\text{tr } S_\lambda$	bias term	noise factor	median $\ a_\lambda(x)\ ^2$	median $\hat{P}_\lambda$	corrected test [%]	kernel alone [%]	half-sample, same / split fit [%]
Y1	input	1e-08	3841	0.000226	0.924	0.87	0.0324	0.2168	0.1754	0.079 / 2.465
Y1	input	1e-06	1624	0.0712	0.26	0.226	0.0395	0.2376	0.2420	1.145 / 2.214
Y1	input	0.0001	212	0.282	0.0269	0.0237	0.0997	0.4132	0.9768	2.221 / 2.395
Y1	input	0.01	19	0.413	0.00194	0.00197	0.315	0.5095	5.7578	2.716 / 2.728
Y1	feature	1e-08	2541	0.0012	0.51	0.481	0.00637	0.2027	0.2111	0.280 / 1.171
Y1	feature	1e-06	689	0.0225	0.111	0.0668	0.0151	0.3130	0.3275	1.085 / 1.464
Y1	feature	0.0001	103	0.156	0.0137	0.00825	0.0529	0.4546	0.6396	2.105 / 2.187
Y1	feature	0.01	9	0.418	0.001	0.000845	0.185	0.5326	5.9929	2.783 / 2.786
Y2	input	1e-08	3884	0.000107	0.943	0.523	0.0369	0.4559	0.6885	0.043 / 1.784
Y2	input	1e-06	1554	0.061	0.219	0.158	0.0416	0.4849	0.9183	0.886 / 1.632
Y2	input	0.0001	132	0.211	0.0143	0.0127	0.0845	0.6325	4.4719	1.516 / 1.580
Y2	input	0.01	10	0.273	0.00115	0.00117	0.227	0.8581	26.9965	1.773 / 1.777
Y2	feature	1e-08	2339	0.000906	0.451	0.44	0.00562	0.1431	0.1439	0.184 / 0.738
Y2	feature	1e-06	564	0.0159	0.0893	0.0618	0.0143	0.2768	0.2696	0.662 / 0.899
Y2	feature	0.0001	86	0.134	0.0118	0.00829	0.0512	0.5764	1.0187	1.275 / 1.326
Y2	feature	0.01	9	0.27	0.000998	0.000909	0.188	0.8905	15.5377	1.738 / 1.741
Y3	input	1e-08	3308	0.00797	0.7	0.412	0.0132	0.2382	0.3731	0.324 / 2.147
Y3	input	1e-06	558	0.283	0.0644	0.0523	0.0192	0.2627	0.9135	1.656 / 1.951
Y3	input	0.0001	45	0.455	0.00546	0.00501	0.0457	0.3432	5.4172	2.173 / 2.155
Y3	input	0.01	5	0.554	0.000642	0.000643	0.152	0.4914	19.3748	2.715 / 2.702
Y3	feature	1e-08	2371	0.00158	0.461	0.458	0.00569	0.0999	0.1197	0.221 / 0.903
Y3	feature	1e-06	590	0.0255	0.094	0.063	0.0145	0.1650	0.1866	0.829 / 1.136
Y3	feature	0.0001	90	0.267	0.0123	0.00842	0.0513	0.3145	0.6620	1.790 / 1.882
Y3	feature	0.01	9	0.525	0.001	0.000885	0.189	0.4913	7.4936	2.612 / 2.624
Y4	input	1e-08	3998	9.17e-07	0.999	0.473	0.269	2.1826	2.3706	0.007 / 12.844
Y4	input	1e-06	3792	0.00736	0.9	0.437	0.27	2.1557	2.3495	0.629 / 12.779
Y4	input	0.0001	1074	1.16	0.125	0.0912	0.319	1.9633	2.4907	7.986 / 11.573
Y4	input	0.01	51	2.3	0.00312	0.00306	0.619	2.3401	8.9051	12.340 / 12.554
Y4	feature	1e-08	3823	0.00807	0.917	0.975	0.031	2.4285	2.4310	0.626 / 14.459
Y4	feature	1e-06	1652	0.794	0.272	0.24	0.0392	1.8004	1.8078	6.037 / 11.377
Y4	feature	0.0001	229	1.83	0.0293	0.0257	0.103	1.8454	1.9424	9.636 / 10.497
Y4	feature	0.01	19	2.39	0.0019	0.00181	0.323	2.4706	4.7750	12.907 / 12.992

Table 15: Ridge-path diagnostics of the residual correction on a 4,000-row training subsample (Proposition 4): effective degrees of freedom, the bias and noise terms of item (iii), the out-of-sample noise factor and the power function on held-out rows, the corrected and kernel-alone test errors along the ridge, and the same-sample against split-sample fit of the correction read on the held-out half. The test-error columns are a curve read once at every ridge value, not a selected model.

β	slope, MSE vs ε_R^2	$\beta/(\beta+2)$	slope, uniform family	analytic slope, same ladder	MC vs analytic (uniform)	MC vs analytic (profile)	violations
0.5	0.200	0.200	0.494	0.494	5.7e-04	4.4e-04	0
1	0.333	0.333	0.973	0.973	1.7e-03	3.9e-04	0
2	0.498	0.500	1.721	1.724	1.0e-02	3.1e-04	0
3	0.597	0.600	1.963	1.963	1.4e-03	5.2e-04	0
4	0.664	0.667	1.995	1.995	2.8e-04	2.8e-04	0

Table 16: The model of Proposition 7 with $F_t(u) = u^\beta$, two million samples per point, the constrained retrieval run as specified: the measured exponent of the mean-square reflectance error against ε_R^2 for the average-error family against $\beta/(\beta+2)$, the measured and analytic slopes of the uniform family on the same threshold ladder, the largest relative difference between the Monte Carlo and the analytic curves of the uniform and profile families, and the number of entries at which the pointwise inequality fails.

split	inputs	head	RMSE	MAE	R^2	SMAPE [%]
OOD	state only	dkr	0.01408±0.00177	0.00709±0.00083	0.98973±0.00260	5.714±0.441
OOD	state only	krr	0.11283±0.00069	0.05430±0.00197	0.34750±0.00794	30.762±1.160
OOD	state only	mlp	0.00750±0.00044	0.00355±0.00036	0.99711±0.00034	4.500±0.249
OOD	state only	mlp_ens	0.00624±0.00015	0.00258±0.00017	0.99800±0.00010	3.470±0.370
OOD	state only	mlp_resid	0.00620±0.00015	0.00252±0.00015	0.99803±0.00009	3.372±0.308
OOD	state only	ridge	0.11944±0.00000	0.06195±0.00000	0.26894±0.00000	37.228±0.000
OOD	with 6S inputs	dkr	0.00735±0.00036	0.00336±0.00033	0.99723±0.00027	4.101±0.474
OOD	with 6S inputs	krr	0.02904±0.00015	0.01432±0.00005	0.95678±0.00045	11.238±0.033
OOD	with 6S inputs	lowfi_resid	0.02901±0.00016	0.01430±0.00006	0.95687±0.00046	11.416±0.047
OOD	with 6S inputs	mlp	0.00741±0.00080	0.00347±0.00063	0.99716±0.00062	4.406±0.366
OOD	with 6S inputs	mlp_ens	0.00611±0.00030	0.00246±0.00023	0.99808±0.00019	3.009±0.093
OOD	with 6S inputs	mlp_resid	0.00596±0.00031	0.00232±0.00022	0.99818±0.00019	2.921±0.067
OOD	with 6S inputs	ridge	0.08048±0.00000	0.04396±0.00000	0.66806±0.00000	29.064±0.000
standard	state only	dkr	0.00874±0.00121	0.00387±0.00071	0.99593±0.00116	5.342±0.323
standard	state only	krr	0.11288±0.00001	0.05157±0.00002	0.32981±0.00008	31.178±0.032
standard	state only	mlp	0.00657±0.00026	0.00242±0.00030	0.99773±0.00018	5.136±0.550
standard	state only	mlp_ens	0.00591±0.00007	0.00170±0.00005	0.99816±0.00005	3.873±0.132
standard	state only	mlp_resid	0.00581±0.00004	0.00151±0.00004	0.99822±0.00002	3.543±0.066
standard	state only	ridge	0.11959±0.00000	0.06067±0.00000	0.24772±0.00000	41.647±0.000
standard	with 6S inputs	dkr	0.00627±0.00025	0.00218±0.00021	0.99793±0.00016	4.429±0.169
standard	with 6S inputs	krr	0.02514±0.00010	0.01155±0.00003	0.96674±0.00026	10.184±0.037
standard	with 6S inputs	lowfi_resid	0.02513±0.00010	0.01155±0.00003	0.96679±0.00026	10.178±0.041
standard	with 6S inputs	mlp	0.00632±0.00018	0.00224±0.00017	0.99790±0.00012	4.496±0.160
standard	with 6S inputs	mlp_ens	0.00587±0.00014	0.00163±0.00012	0.99819±0.00008	3.407±0.130
standard	with 6S inputs	mlp_resid	0.00578±0.00012	0.00148±0.00008	0.99824±0.00007	3.185±0.097
standard	with 6S inputs	ridge	0.06981±0.00000	0.03631±0.00000	0.74366±0.00000	28.300±0.000
standard	published	pKANrtm	0.00619	0.00186	0.99472	3.912
OOD	published	pKANrtm	0.01022	0.00544	0.99418	5.399

Table 17: The correction-coefficient corpus of [4] under its own protocol: the release’s state split (35,000 / 7,500 / 7,500 states) and the paper’s out-of-distribution split, whose test states lie above the 0.85 quantile of aerosol optical depth or of column water vapour, with the remaining states split 85/15; the paper’s categorical inputs, with and without the 6S [13] coefficients as inputs; pooled RMSE, MAE, R^2 and SMAPE over the three coefficients and thirteen bands, mean and standard deviation over three seeds (standard) or three initialisations (OOD); the published values on the last rows.

band	rule	reduced rel. L^2 [%]	radiance rel. L^2 [%]
o2	select_radiance	4.08±0.05	0.0283±0.0017
o2	select_reduced	4.09±0.06	0.0283±0.0017
o2	single dkr_flat	4.07±0.06	0.0757±0.0012
o2	single dkr_radx	20.14±0.07	0.0351±0.0006
o2	single dkr_wnum	21.70±0.60	0.0309±0.0017
o2	single kernel_ard	30.02±0.03	0.2706±0.1030
o2	single kernel_raw	40.53±0.09	0.2353±0.0047
o2	single mean_flat	4.40±0.08	0.1756±0.0111
o2	single mean_radx	58.72±0.60	0.0510±0.0049
o2	single mean_wnum	49.22±0.34	0.0437±0.0002
o2	single ridge_flat	4.67±0.09	0.2205±0.0124
o2	single ridge_radx	44.83±0.43	0.0538±0.0044
o2	single ridge_wnum	45.55±0.73	0.0527±0.0016
o2	stack_radiance	18.59±0.25	0.0264±0.0007
o2	stack_reduced	4.10±0.07	0.0291±0.0029
sco2	select_radiance	8.09±0.02	0.0598±0.0013
sco2	select_reduced	8.09±0.02	0.0598±0.0013
sco2	single dkr_flat	8.09±0.02	0.1083±0.0114
sco2	single dkr_radx	15.27±0.25	0.0780±0.0094
sco2	single dkr_wnum	15.48±0.31	0.0638±0.0022
sco2	single kernel_ard	28.72±0.02	0.6023±0.0097
sco2	single kernel_raw	41.01±0.09	0.5229±0.0081
sco2	single mean_flat	8.23±0.01	0.2107±0.0194
sco2	single mean_radx	46.08±0.91	0.0936±0.0082
sco2	single mean_wnum	36.64±0.38	0.0750±0.0020
sco2	single ridge_flat	8.34±0.02	0.2568±0.0244
sco2	single ridge_radx	41.14±1.78	0.0967±0.0081
sco2	single ridge_wnum	40.25±0.24	0.0887±0.0034
sco2	stack_radiance	8.87±0.15	0.0511±0.0016
sco2	stack_reduced	8.09±0.01	0.0536±0.0020
wco2	select_radiance	16.26±0.07	0.0510±0.0086
wco2	select_reduced	16.26±0.07	0.0510±0.0086
wco2	single dkr_flat	16.26±0.07	0.1135±0.0069
wco2	single dkr_radx	48.28±0.28	0.0475±0.0055
wco2	single dkr_wnum	48.60±0.38	0.0516±0.0086
wco2	single kernel_ard	52.45±0.07	0.2231±0.0014
wco2	single kernel_raw	58.63±0.20	0.4431±0.0063
wco2	single mean_flat	16.42±0.06	0.2254±0.0161
wco2	single mean_radx	67.92±0.46	0.0605±0.0075
wco2	single mean_wnum	62.39±0.42	0.0700±0.0087
wco2	single ridge_flat	16.49±0.07	0.2571±0.0234
wco2	single ridge_radx	61.03±0.31	0.0663±0.0080
wco2	single ridge_wnum	61.33±0.05	0.0778±0.0095
wco2	stack_radiance	37.43±0.20	0.0391±0.0026
wco2	stack_reduced	16.28±0.07	0.0404±0.0034

Table 18: OCO-2 [18]: the heads of the seeded campaign (three training losses, the kernel and the ridge readout on each network’s features, the matched raw-input kernels) and four combination rules fitted on validation and read once on test: the per-coordinate selection on the reduced

corpus	n	d	q	runs	ridge (linear)	KRR	KRR, learned metric	network	network ensemble	network + residual KRR	kernel on features	convex stack
Advection	19,000	200	200	5	11.29±0.04	11.36±0.04	11.36±0.04	11.69±0.07	11.50±0.05	11.50±0.05	11.64±0.07	11.36±0.04
Burgers $\nu = 0.1$	7,200	256	128	3	13.89±0.53	2.81±0.28	2.90±0.33	20.34±0.61	19.28±0.09	2.78±0.29	4.92±0.30	2.78±0.29
Burgers $\nu = 0.01$	7,200	256	128	3	33.68±0.28	23.40±0.54	23.58±0.50	36.65±1.95	35.31±0.70	23.12±0.56	29.20±0.59	23.01±0.54
Burgers $\nu = 0.001$	7,200	256	128	3	38.38±0.12	31.66±0.16	33.23±1.73	39.98±0.65	38.89±0.11	31.82±0.11	35.59±0.48	31.34±0.14
Turbulent radiative layer	7,200	256	256	3	31.73	30.31	30.26±0.01	32.17±0.02	31.58±0.02	30.40±0.01	31.08±0.01	30.12
Active matter	14,000	256	256	1	37.24	35.81	35.18	37.31	36.93	35.77	36.22	35.05
Helmholtz staircase	20,384	256	256	1	0.00	0.26	0.07	0.40	0.24	0.23	0.45	0.07
ClimSim	100,000	124	128	5	32.24±0.06	27.20±0.05	26.31±0.04	25.10±0.18	24.43±0.17	24.33±0.13	24.71±0.25	24.25±0.17
RT corrections, state inputs	487,773–487,823	8	3	5	71.54±0.59	62.93±0.51	3.97±0.42	3.75±0.25	3.24±0.11	3.13±0.05	3.52±0.07	3.09±0.05
RT corrections, multi-fidelity inputs	487,773–487,823	11	3	5	33.83±0.26	15.76±1.98	4.60±0.31	2.79±0.14	2.40±0.11	2.34±0.09	2.80±0.09	2.32±0.08

Table 19: Test error [%] on ten further emulation configurations, mean $\pm$ standard deviation over the available runs. Errors are mean relative L^2 norms, except the two correction-coefficient rows, which use the denominator floor in (22). n is the number of examples — a range where the state-level split of the correction-coefficient corpus makes it differ by run — d the input dimension the models receive and q the output dimension they regress. The runs are random splits except for the turbulent radiative layer, whose three runs share The Well’s own split and differ in network initialisation, and the two single-run rows. The family set and validation-based selection rule are common across rows; the network budget is not, and is 200 epochs at width 512 except on ClimSim (60 at 512) and the two correction-coefficient configurations (120 at 384). The ridge column is linear in the inputs the models receive, unlike the cubic reference of Table 1, and is not a member of the convex stack; its Helmholtz-staircase entry is 0.004, below the two decimals shown. Kernel solves are exact when the training block has at most 20,000 rows and otherwise use 6,000 Nyström landmarks, which is the case on ClimSim, the Helmholtz staircase and the two correction-coefficient rows.