

Neural means and kernel corrections for an EMIT atmospheric radiative-transfer emulator

Yitzchak Shmalo*

September 2026

Abstract

We compare neural networks, exact Matérn kernel regression, and their combinations on a six-dimensional radiative-transfer table over the 285-band EMIT grid. Across ten random splits with validation-based selection and 18,884 training samples, input-scaled kernels achieve 0.095% mean relative radiance error, compared with 0.76% for cubic regression and 0.39% for a neural network. Residual correction reduces the network error to 0.14%; kernels on wider-network features and their convex stacks reach 0.079% and 0.078%. This small stacking gain does not extend to reflectance tails: the wide stack has a larger 95th-percentile inversion error at every split. Learning-curve, output-rank, and width experiments expose sensitivity to the representation and tuning protocol. We give exact residual and error-comparison identities and a conditional stability bound for reflectance inversion, without asserting universal learning rates. Contrasting OCO-2 results delimit the coupling’s scope. The findings concern interpolation of tabulated physics, not validated operational retrieval.

1 Introduction

Imaging spectroscopy missions such as EMIT [13], SBG, and AVIRIS-NG invert measured radiances for surface and atmospheric quantities, and the inversion loop calls an atmospheric radiative-transfer model many times per spectrum. At mission data rates the physics code itself (MODTRAN in the pipelines relevant here [1]) is far too slow, so operational retrievals interpolate precomputed lookup tables, and a body of work replaces table and physics alike with learned emulators: Gaussian-process surrogates [5], the neural-network MODTRAN emulator that serves the EMIT mission’s operational atmospheric correction [3, 2], and, recently, physics-guided and multi-fidelity architectures for correction coefficients [4]; network training itself can draw on a growing toolbox of refinements [7, 8]. Both families have a record on emulation tasks of this kind: networks learn their own representations and scale easily to wide spectral outputs [10, 11], while kernel and Gaussian-process methods come with exact solves and a strong record both on operator-learning benchmarks [9] and on smooth, low-dimensional regression [14]. Which family is preferable for a given emulation problem, and whether they can be combined so that neither’s advantage is given up, are practical questions that the structure of the data itself should decide.

This paper studies those questions on a tabulated radiative-transfer emulation problem over the EMIT wavelength grid, using the coupling construction of [6]: train a neural network on

*Einstein Institute of Mathematics, The Hebrew University of Jerusalem, Jerusalem, Israel.
yitzchak.shmalo@gmail.com.

standardized output bands as the mean, correct it with exact Matérn kernel regressions — of its residuals on the input space, and of the target on the network’s own learned features — and let a per-coordinate selection on held-out data choose, coefficient by coefficient, among the network, the kernels, and the corrected predictors. We ask what structure the dataset has, which model families that structure favors, and whether the kernel correction of a neural mean improves on both of its parents here.

The dataset answers the first question quickly (Section 2). The input is six-dimensional, and two coordinates carry most of the variance of the output; the output spectra are smooth curves whose empirical rank is a few dozen out of 285 bands; and the map from state to spectrum is smooth enough that cubic polynomial regression already lands near two percent relative error. In the taxonomy of [6] this is the regime that favors exact kernel methods, in contrast to emulation problems with higher-dimensional reduced states where representation learning carries more of the weight. The interest of the comparison is therefore not whether some exotic architecture wins, but how the familiar families divide the problem between them when the geometry is easy, and whether the coupling machinery preserves the best of each.

The contribution is a controlled empirical comparison, an analysis of metric-dependent model rankings, and exact algebraic diagnostics for the fitted constructions; it is not a new architecture, a universal rate theorem, or a validated operational retrieval system.

Sections 3 and 4 give the protocol and the measurements. The results are these. With every training row in the solve and a length scale per input, the input-scaled kernel is the strongest raw-input kernel and narrowly leads the width-512 feature kernel on mean radiance error; its learned metric is consistent with the expected anisotropy. The kernel regression of the network’s residuals does what it was built to do, handing the network the kernel’s accuracy on the smooth components, but it is no longer the lowest row. A per-component convex stack of the heads is, by about one standard deviation across the ten splits, while per-coordinate selection does not transfer from validation to test. The reflectance metric has a difficulty of its own, set by the atmospheric absorption windows in which the transmitted flux approaches zero: there the sensitivity of the inversion to the path radiance and the two flux components is unbounded, and at exactly zero flux the physical reflectance is not identifiable at all. We measure what the metric returns for exact inputs and report reflectance medians together with all-band tails and a finite-perturbation stability bound. The tails reveal a ranking reversal not visible in the radiance mean. Section 5 adds what a single training size cannot show: learning curves from 500 rows to the full block, refits at other output ranks, a ladder of network widths, and the two questions those curves raise — where the kernel’s hyperparameters should be chosen, since the protocol used here turns out to cost it a measurable amount, and whether its metric can be read off the data instead of searched for. Section 6 gives the identity behind the residual correction, which separates exact residual identities and error comparisons from unproved learning-rate interpretations. Section 7 records where the constructions that win here do not transfer, on the OCO-2 problem of the companion paper [6], and Section 8 applies the same family set to ten further emulation configurations.

2 Data and protocol

The data are $n = 23,313$ tabulated radiative-transfer evaluations. Each input $x \in \mathbb{R}^6$ collects the atmospheric and viewing state; the coordinate ranges are consistent with the standard lookup-table axes for imaging-spectroscopy atmospheric correction [2] — an aerosol optical depth in

[0.05, 1.2], a surface elevation in $[-0.5, 6]$ km, a water-vapor column in $[0.05, 7]$ cm, a relative azimuth in $[0, 180]^\circ$, a solar zenith angle in $[0, 65]^\circ$, and a view zenith angle within 30° of nadir. Each output is a set of four spectral components on the same 285-point EMIT wavelength grid spanning 381–2493 nm: path radiance Y_1 , direct flux Y_2 , diffuse flux Y_3 , and spherical albedo Y_4 (Figure 1). The quantities retrieved downstream are the radiance at a test surface reflectance ρ ,

$$L(\rho) = Y_1 + \rho \frac{Y_2 + Y_3}{1 - \rho Y_4}, \quad (1)$$

and the reflectance recovered by inverting (1) with predicted components at the true radiance,

$$\hat{\rho} = \frac{L - \hat{Y}_1}{\hat{Y}_2 + \hat{Y}_3 + \hat{Y}_4(L - \hat{Y}_1)}, \quad (2)$$

both evaluated at $\rho = 0.7$.

The protocol is the same for every family. The table is split into a test block of 2,331 rows and a remaining block of 20,982; 2,098 of the latter are held out for validation, leaving 18,884 for fitting. The procedure is repeated with seeds 101–110. Standardization and principal components are fitted only on each split’s training rows. Hyperparameters, checkpoints, coordinatewise selections and combination weights use that split’s validation rows. Each fitted family is then evaluated on that split’s test rows. Unless stated otherwise, tables report the mean and sample standard deviation over the ten splits.

These splits assess interpolation within one fixed lookup table. An earlier exploratory version consulted a test split of the same table. Redrawing the split does not erase that information or create an untouched external dataset: rows can reappear across split roles, and the study’s choice of families and subsequent ablations was informed by earlier results. The repeated-split evidence is therefore descriptive rather than a confirmatory test on ten independent datasets. The main and wide-pipeline tables have complete individual records; some secondary ablations are retained only partly at per-seed level, as specified under data availability. Paired differences on common splits are reported in Section 4.1, without treating their SD as a confidence interval. The same validation rows serve early stopping, tuning and combination, so they are not an independent calibration sample and support no coverage guarantee. Evaluation at a single surface reflectance, $\rho = 0.7$, does not establish performance for variable surface spectra.

Three empirical structural diagnostics motivate the model comparison.

The outputs are low rank. On standardized training outputs, 99.99% of the variance of every component lies in the first 11–25 principal components, and 99.9999% in at most 44 (Figure 2, left); the minimum adjacent-band correlation across the four components is 0.994952 on the training rows. All regressors below act on the first 64 principal coefficients, whose training unexplained-variance fraction is below 10^{-6} , except the networks, which are trained directly on the 285 standardized bands. A small aggregate variance tail is not a uniform relative-error guarantee in physical units; Section 6 states the projection error explicitly.

The measured input relevance is strongly anisotropic. Shuffling the aerosol coordinate destroys a linear model’s validation R^2 (a drop of 0.89); the solar zenith follows at 0.54; elevation and water vapor sit near 0.1; the two remaining angles barely register (Figure 2, right). The values reported for these probes, and for the other structure measurements used to motivate architectures, are their validation-split rescaling; the initial exploratory versions were test-scored as disclosed above. Six nominal dimensions have four coordinates with measurable first-order relevance, two of them dominant. This finite-sample diagnostic helps explain why smooth methods work, but a selected scalar length scale does not remove anisotropy. Any theorem attributing




Figure 1: Example output curves across the EMIT wavelength grid, five samples per panel. The curves are smooth, share the deep water-vapor absorption features near 1380 and 1900 nm, and differ mainly by smooth rescalings along the state axes.

generalization to anisotropic metric adaptation would also have to control regularization bias, selection of scale and nugget, and the empirical output-rank truncation. We use anisotropy as a measured property of this split, not as such a guarantee.

Low-degree surrogates are accurate. A linear model in the standardized state reaches $R^2 = 0.83$; distance-weighted 5-nearest-neighbors does worse (0.74), so this particular local baseline underperforms the linear one; and a cubic polynomial ridge reaches mean relative L^2 errors near two percent (Table 1). Smooth, low-dimensional, low-rank: in the taxonomy of the companion paper this is the regime of the Helmholtz and Navier–Stokes benchmarks, where the kernel stage won the internal selection, not the regime of its OCO-2 problem, where representation learning carried the problem.

One more property of the task matters for reading any reflectance number. In the deep absorption windows the transmitted flux $Y_2 + Y_3$ collapses toward zero — in this table, to values as small as 10^{-21} — and the inversion (2) is undefined at a fully opaque band (0/0): the physical reflectance is non-identifiable there. The protocol maps a non-finite inverse to zero, but that branch is not what creates the measured true-component error. The exact components, put through the same evaluation, give a reflectance RMSE of 0.0054 against the true $\rho = 0.7$, produced entirely by 0.25% of (sample, band) entries — 1,652 of 664,620, measured on the earlier single split, whose test block held one row more than the ten drawn here — whose denominators cancel to below 2×10^{-13} ; every value returned is finite (the protocol’s zero-fill for non-finite inversions is never triggered by the true components), and the effect is invisible in component




Figure 2: Left: explained-variance spectra of the standardized outputs. Every component concentrates 99.99% of its variance in 11–25 principal components out of 285 bands. Right: drop in linear validation R^2 when a single standardized input coordinate is shuffled. Aerosol optical depth and solar zenith angle dominate; the azimuth and view-zenith coordinates are almost inert at first order.

errors. What this measures is the conditioning of the inversion at those entries and not the accuracy of any emulator: it is what the evaluation returns when the components handed to it are exact. It is not a lower bound on an emulator’s reflectance error. A predictor whose component errors offset the cancellation would score below it, and a predictor with generic errors of its own can score far above it; what the measurement establishes is that a mean-square reflectance summary on this table reports the conditioning of these entries alongside the emulator’s errors. Median error summarizes the center of the distribution but cannot assess its tail; both are reported below.

Differentiating (2) separates the components. Write $D = \hat{Y}_2 + \hat{Y}_3 + \hat{Y}_4(L - \hat{Y}_1)$ for its denominator. Then

$$\frac{\partial \hat{\rho}}{\partial \hat{Y}_4} = -\hat{\rho}^2, \quad \frac{\partial \hat{\rho}}{\partial \hat{Y}_2} = \frac{\partial \hat{\rho}}{\partial \hat{Y}_3} = -\frac{\hat{\rho}}{D}, \quad \frac{\partial \hat{\rho}}{\partial \hat{Y}_1} = -\frac{\hat{Y}_2 + \hat{Y}_3}{D^2}. \quad (3)$$

At exact, identifiable components, $\hat{\rho} = \rho = 0.7$ and the albedo derivative is -0.49 , independent of the positive transmitted flux. This is a local statement at the truth, not a global bound for arbitrary predictions: when an inaccurate inverse has large $|\hat{\rho}|$, its albedo derivative can also be large. The explicit inverse-denominator sensitivity of the path radiance and flux components explains why errors near opaque bands require separate examination. In the earlier single-split diagnostic, 7.4% of entries had absolute reflectance error above 0.5, with 95% of those at flux below one percent of its median. That earlier diagnostic is not part of the ten-split model comparison. The current all-band tails are reported directly in Section 4.1.

2.1 Finite-perturbation control of the inversion

The derivatives in (3) describe local conditioning. The following elementary identities give a finite-error test and distinguish forward accuracy from inversion stability. For one sample and band, write $a = Y_1$, $t = Y_2 + Y_3$, $s = Y_4$ and $L = a + \rho t / (1 - \rho s)$. Set $e_a = \hat{a} - a$, $e_t = \hat{t} - t$, $e_s = \hat{s} - s$.

Proposition 1 (Conditional radiance and reflectance stability). *Suppose $0 \leq \rho \leq \rho_{\max}$, $t > 0$ and $q = 1 - \rho s > 0$. If also $\hat{q} = 1 - \rho \hat{s} > 0$, the forward error satisfies the exact identity*

$$\hat{L} - L = e_a + \frac{\rho e_t}{\hat{q}} + \frac{\rho^2 t e_s}{q \hat{q}}. \quad (4)$$

In particular, if $q, \hat{q} \geq \kappa > 0$,

$$|\hat{L} - L| \leq |e_a| + \rho_{\max} |e_t| / \kappa + \rho_{\max}^2 |t e_s| / \kappa^2.$$

For inversion at the true radiance, put $u = L - a = \rho t / q$, $D = t / q$, and $\hat{D} = \hat{t} + \hat{s}(L - \hat{a})$. Whenever $\hat{D} \neq 0$,

$$\hat{\rho} - \rho = -\frac{(1 - \rho \hat{s})e_a + \rho e_t + \rho u e_s}{\hat{D}}, \quad \hat{D} - D = e_t + u e_s - \hat{s} e_a. \quad (5)$$

If $B = |e_t| + |u e_s| + |\hat{s} e_a| \leq \eta D$ for some $0 \leq \eta < 1$, then $\hat{D} \geq (1 - \eta)D > 0$ and

$$|\hat{\rho} - \rho| \leq \frac{|1 - \rho \hat{s}| |e_a| + \rho |e_t| + \rho |u e_s|}{(1 - \eta)D}. \quad (6)$$

Proof. Subtract the two forward maps and use $1/\hat{q} - 1/q = \rho e_s / (q \hat{q})$ to obtain (4). Since $u = \rho D$, $\hat{D} = D + e_t + u e_s - \hat{s} e_a$ and $\hat{\rho} - \rho = (u - e_a - \rho \hat{D}) / \hat{D}$, which gives (5). The triangle inequality yields $|\hat{D} - D| \leq B$ and then (6). $\square$

When $\hat{q} \neq 0$, the same calculation gives the exact forward-inverse link

$$\hat{\rho} - \rho = -\frac{\hat{q}}{\hat{D}}(\hat{L} - L). \quad (7)$$

Consequently, on a common set where $|\hat{q}| \leq Q$ and $|\hat{D}| \geq \delta > 0$, the corresponding unweighted L^2 errors obey $\|\hat{\rho} - \rho\|_2 \leq (Q/\delta)\|\hat{L} - L\|_2$. Without the denominator condition there is no uniform implication from small spectral relative radiance error to small reflectance error. For example, take two bands with $a = s = 0$, $t = (1, \varepsilon)$ and $0 < \rho < 1$. A predictor with $\hat{t} = t$, $\hat{s} = 0$ and $\hat{a} = (0, c\varepsilon)$, where $0 < c < \rho$, has relative spectral radiance error $c\varepsilon / (\rho\sqrt{1 + \varepsilon^2}) \rightarrow 0$ but retrieves $(\rho, \rho - c)$, so its reflectance RMSE remains $c/\sqrt{2}$. Every denominator is positive for $\varepsilon > 0$ and both retrieved reflectances remain physical. This is a finite synthetic example, not an additional EMIT experiment.

The forward bound need not deteriorate when t tends to zero, whereas the inverse bound requires error control relative to $D = t/q$. Thus a small spectral radiance error does not certify stable reflectance recovery. At $t = 0$, $L = a$ is independent of ρ wherever the forward formula is defined, so reflectance is unidentifiable from this band. The identities are exact-arithmetic statements; floating-point cancellation in $L - a$ can further affect an implemented inverse.

A retrieval assessment should therefore specify a common application passband and a conditioning threshold before comparing predictors, report the retained fraction, and count invalid inversions instead of hiding them by zero filling or clipping. The utility `code/conditioned_reflectance.py` implements this diagnostic with a threshold relative to a positive flux scale determined on training data and a common truth-based evaluation mask. The latter is available in a simulation benchmark, not automatically in a real retrieval. Prediction failures on that mask are reported separately. This diagnostic is supplied for reruns with the original arrays; conditioned EMIT results cannot be reconstructed from the public aggregate metrics and are not claimed here. The all-band tail analysis in Section 4.1 uses only quantities already stored in the public run records.

3 Models

All models are trained per component and evaluated in physical units after undoing the standardization; unless stated otherwise they regress the 64 leading principal coefficients of the standardized component, with the training variance criterion and projection qualification of Section 2. Every hyperparameter named below is selected on the validation set of the seed at hand. Everything runs on CPU.

Reference family. Cubic polynomial ridge regression, acting on all monomials of degree at most three in the six standardized inputs, with the ridge weight selected from $\{10^{-10}, 10^{-8}, 10^{-6}, 10^{-4}, 10^{-2}\}$. It provides a low-cost, low-degree reference; its empirical error alone does not prove a smoothness class for the underlying function.

Exact kernel ridge regression. Three constructions, all Matérn kernels on the standardized inputs with the length scale expressed as a multiple of the median pairwise distance of the fit set. The listed nugget is γ in $(K + n\gamma I)\alpha = Z$, not the unscaled diagonal term. The first construction is kept for comparison: a Matérn-5/2 kernel fitted on a uniformly drawn subset of 4,000 training points, one eigendecomposition per length scale, multiplier in $\{1, 2\}$, nugget in $\{10^{-8}, 10^{-6}, 10^{-4}, 10^{-2}\}$. The second is the exact solve on every training row: smoothness $\nu \in \{3/2, 5/2\}$, multiplier in $\{0.5, 1, 2, 4\}$ and the same nugget grid, tuned on validation from a 6,000-row training subsample and the winner refit on all 18,884 rows by one Cholesky factorization. The third gives each input its own length scale: starting from the isotropic winner, a coordinate search over per-input multipliers in $\{1/4, 1/2, 1, 2, 4\}$, two sweeps, scored on validation, after which the smoothness, scale and nugget are re-tuned at the found metric and the solve is refit on every row. The input relevance measured in Section 2 spans more than an order of magnitude across the six coordinates, which is what the third construction is for.

Fully connected networks. Three ELU hidden layers of 512 units per component, trained on the 285 standardized bands with AdamW at learning rate 10^{-3} , cosine decay, batch size 1024, weight decay 10^{-6} , up to 150 epochs with early stopping on validation loss (patience 25). Five networks are trained per component and seed from different initializations; the table reports the first and the mean of the five.

Kernel corrections of the neural mean. The construction of [6], with the exact stages now fitted on every training row. (i) *Residual correction*: an exact Matérn regression, on the standardized input space, of the network’s training residuals in coefficient space, added back to the network’s projection into that space, with smoothness, scale and nugget selected by the validation error of the corrected prediction; applied to the single network and to the five-network mean. (ii) *Kernel on features*: the same exact regression with the input metric replaced by the network’s last hidden layer, standardized, so that the kernel’s exactness is spent in coordinates the network has already straightened. This is the post-hoc counterpart of deep kernel learning [25], in which the representation and the kernel are trained together, and of the kernel-flow regularization of a network’s inner layers [12]; here the representation is fixed first and the kernel is solved exactly on it. (iii) *Per-coordinate selection*: for each of the 64 output coefficients independently, pick on validation among the exact kernel, the network, the ensemble, the two corrected predictors and the feature kernel, and assemble the test prediction from the

model	component rel. L^2 [%]				radiance	reflectance
	Y_1	Y_2	Y_3	Y_4	rel. L^2 [%]	med. $ \hat{\rho} - \rho $ [pp]
Cubic ridge	1.16±0.02	2.18±0.04	1.11±0.03	2.98±0.06	0.760±0.011	0.471±0.010
Matérn KRR, 4000-point fit	0.17±0.01	0.61±0.04	0.29±0.02	2.44±0.87	0.332±0.029	0.178±0.030
Matérn KRR, exact on all rows	0.08±0.01	0.42±0.03	0.20±0.00	2.62±0.18	0.239±0.016	0.141±0.018
Matérn KRR, per-input scales	0.08±0.01	0.16±0.02	0.08±0.01	2.18±0.16	0.095±0.009	0.036±0.004
FC-DNN (3×512)	0.48±0.02	0.97±0.04	0.41±0.02	3.02±0.18	0.387±0.012	0.245±0.008
Five-network mean	0.37±0.01	0.67±0.03	0.28±0.01	2.51±0.09	0.251±0.006	0.163±0.003
DNN + residual KRR	0.10±0.00	0.30±0.02	0.16±0.01	2.43±0.08	0.141±0.006	0.089±0.008
Ensemble + residual KRR	0.08±0.00	0.27±0.03	0.13±0.00	2.42±0.17	0.123±0.007	0.080±0.008
Kernel on network features	0.18±0.01	0.17±0.01	0.10±0.01	1.99±0.08	0.098±0.007	0.054±0.004
Per-coordinate selection	0.19±0.01	0.44±0.10	0.13±0.02	2.25±0.18	0.193±0.022	0.123±0.023
Convex stack	0.07±0.00	0.16±0.01	0.09±0.00	1.94±0.06	0.087±0.006	0.048±0.003

Table 1: Test errors under the protocol of Section 2, mean $\pm$ standard deviation over the ten splits. Component columns are mean relative L^2 errors; the radiance column evaluates (1) at $\rho = 0.7$; the last column is the median absolute reflectance error of the inversion (2), in reflectance percentage points, under the zero-fill evaluation of Section 2. The median is used because the sensitivity of (2) to the path radiance and the two flux components diverges continuously as the transmitted flux vanishes, so mean-square and tail-percentile summaries are dominated, for every model alike, by near-opaque entries; the exact components put through the same evaluation give an RMSE of about half a percentage point there (Section 2), which is a property of those entries and not a bound on any emulator. The tail summaries are also displayed in Table 3; a small median does not certify retrieval performance.

winners. (iv) *Convex stack*: per component, a convex combination of the same six heads with weights fitted on the validation points in the physical relative metric, in the manner of stacked generalization [23] and stacked regression [24]. The choice of weighting rule for pools of this kind, and its relation to constrained minimum-variance weighting, is studied separately and is not part of this comparison.

What is evaluated at other budgets. Section 5 repeats the isotropic exact kernel, the network, its residual correction and the feature kernel at nested training sizes, at other output ranks and at a wider network; the table of Section 4 is the full-size, rank-64, width-512 cell of that grid.

4 Results

Table 1 collects the test errors. Every entry is a mean over the ten splits with its standard deviation, and the spreads are small enough that most of the orderings discussed below hold seed by seed; where they do not, the text says so.

The scale of the problem. The cubic ridge reaches 1.86% mean component error and 0.76% on the radiance: the map is smooth, and a polynomial in six variables recovers it to two percent. The earlier construction of the kernel regression, a Matérn-5/2 on a 4,000-point subset, divides both numbers by more than two at seven of the ten splits (0.74% and 0.31%); at the other three its albedo fit collapses (3.3–4.3% against 1.9%), which is what a fit capped at a random subset can do, and its ten-split mean (0.88%, 0.33%) carries that fragility.

Exact kernels. Fitting the same kernel on every training row, one Cholesky factorization of the $18,884 \times 18,884$ Gram matrix per component, lowers the radiance error to 0.24% and the path-radiance error by half (0.08% against 0.17%), and it has no collapsing split (2.3–2.9% on the albedo at every seed), but where the capped fit is sound it is worse on the albedo (2.6% against 1.9%), so the mean component error moves the wrong way at those splits (0.83% against 0.74%). The albedo is the component whose principal spectrum decays slowest and whose values sit within a few percent of zero over much of the grid, and it is also the one where the transfer of a scale and nugget tuned on a 6,000-row subsample to the full solve is least reliable. A length scale per input removes most of this: the per-input kernel reaches 0.63% on the components and 0.095% on the radiance, with the best reflectance column of the table (0.036 percentage points), and its albedo column (2.2%) sits between the two isotropic solves. The learned metric is consistent with the expected component anisotropy: for the two flux components and the spherical albedo the coordinate search gives the relative azimuth a length scale sixteen to thirty-two times longer than the others — a median of sixteen on the direct flux and the albedo and of thirty-two on the diffuse flux, at every one of the ten seeds for the two flux components and at nine of the ten for the albedo, whose search leaves the azimuth unscaled at the tenth — and the view zenith two to sixteen times longer, while for the path radiance, which does depend on the viewing geometry, the metric stays nearly isotropic. The weak fitted azimuth dependence is consistent with the physical interpretation of these components; a finite length-scale search is not a proof of exact invariance. On this table the per-input exact kernel and the kernel on network features (below) are, within the seed-to-seed spread, the strongest single models.

Networks. The fully connected network lands at 1.22% and 0.39%; the mean of five independently initialized networks at 0.96% and 0.25%, a third off the radiance error for five times the training cost, the usual decorrelation gain. Neither reaches the isotropic exact kernel on any of the first three components, and the network’s advantage on the albedo that the single-split study reported does not survive replication: on Y_4 the five-network mean (2.5%) is behind the 4,000-point kernel wherever that kernel’s fit is sound (1.9% at seven of the ten seeds) and behind the exact kernel with per-input scales at every seed.

Kernel corrections of the neural mean. An exact Matérn regression of the network’s residuals on the standardized inputs takes the single network from 0.39% to 0.14% on the radiance and the five-network mean from 0.25% to 0.12%; the kernel fitted in the network’s last-layer features (DKR) reaches 0.61% and 0.098%, indistinguishable from the per-input kernel on the components and on the radiance — 0.098 against 0.095, the lower at six seeds of ten — but half again worse on the reflectance median, 0.054 against 0.036 points, and worse there at every one of the ten. The correction therefore does exactly what it did on the single split — it hands the network the kernel’s accuracy on the smooth components while keeping the network’s own representation — but with the kernel now allowed its own input metric the corrected network is no longer the lowest row: the ordering by radiance error is stack, per-input kernel, DKR, corrected ensemble, corrected network, with the first four within 0.04 percentage points of one another.

Selection and stacking. Both combiners here draw on the same six heads — the isotropic exact kernel, the network, the five-network mean, the two corrected predictors and the feature kernel — and not on the per-input kernel, which is the strongest single row of the table. The

per-coordinate selection among those six, made on validation, does not transfer: its test radiance error (0.19%) is worse than three of the heads it selects among, and its spread over seeds (± 0.02) is among the largest in the table. The per-component convex stack, with weights fitted on the validation rows in the physical relative metric, is the best aggregate row (0.57% components, 0.087% radiance, 0.048 reflectance points) and wins the path radiance and the albedo columns outright; its gain over the strongest single row of the table is eight percent on the radiance and ten percent on the components, of the order of one standard deviation of either; measured against the best model the stack actually contains, the kernel on network features, it is eleven percent on the radiance and seven percent on the components. The paired view is sharper than the marginal one: split by split, the stack beats the per-input kernel on the components in nine of the ten splits (mean difference 0.06 points, spread 0.05) and beats DKR in all ten (0.047 ± 0.006), while the per-input kernel and DKR trade places (six of ten). The stack is the best row, and the residual correlations among the members it actually contains say why it can be: on the path radiance the residuals of the isotropic exact kernel and of the five-network mean are correlated at 0.27 ± 0.09 , below the ratio of their errors (0.34) at every one of the ten splits, which is the classical admission condition for a combination to improve on its better member [26], in the form used in [6]. This admission diagnostic concerns a quadratic error metric; it does not directly guarantee improvement in the mean-relative-norm objective used by this stack or in the nonlinear downstream inverse. Not every pair in the pool admits: the two residual-corrected predictors are correlated at 0.88 on the path radiance and at 0.99 on the albedo, and are largely two views of the same predictor. The observed gain is compatible with complementary errors among the heads. It does not make every pair useful, and the best albedo head in this pool is the feature kernel, not the residual-corrected networks. The coordinatewise rule performs worse under this validation budget; both sampling variability and its different selection objective may contribute.

Comparison with the single-split study. The earlier version of this study reported the corrected network as the lowest row, at 0.25% radiance error, and the exact kernel on the raw state at 0.41%; both numbers came from a single split whose test block had been consulted during exploration and from a kernel fitted on 4,000 points with one length scale. Under the fresh-seed protocol the corrected network reads 0.14% and the exact kernel with per-input scales 0.095%. The smoothness of the map, the complementarity of the heads across components and the mechanism of the residual correction hold under replication; the claim that the corrected network is the best model on this table does not.

The albedo. Every family is worst on Y_4 by an order of magnitude, and the per-band errors show why: the spherical albedo’s values in the short-wave infrared are close to zero, so relative error is large for every model there, and the component is hard for reasons of signal-to-noise rather than of model class. The stack’s 1.94% is the best number any construction reaches at this training size.

4.1 Paired comparisons and the radiance–reflectance trade-off

The saved per-seed records permit a comparison that marginal standard deviations alone do not: subtract each candidate’s error from its reference’s error on the same split. Table 2 reports those differences for the main and wide pipelines. Positive gains favor the candidate. The wide feature kernel improves on the input-scaled kernel at all ten splits, by 0.0165 ± 0.0066 radiance

candidate / reference	mean gain $\pm$ SD [pp]	gain range [pp]	wins/10
Input-scaled / isotropic KRR	0.1436 $\pm$ 0.0122	[0.1196, 0.1585]	10
Residual correction / network	0.2464 $\pm$ 0.0090	[0.2365, 0.2615]	10
Narrow stack / input-scaled KRR	0.0083 $\pm$ 0.0066	[-0.0054, 0.0185]	9
Wide feature kernel / input-scaled KRR	0.0165 $\pm$ 0.0066	[0.0054, 0.0295]	10
Wide / narrow feature kernel	0.0191 $\pm$ 0.0079	[0.0096, 0.0352]	10
Wide stack / wide feature kernel	0.0006 $\pm$ 0.0026	[-0.0053, 0.0037]	8

Table 2: Paired radiance contrasts on seeds 101–110. Gain is the reference error minus the candidate error, in percentage points (pp). SD is the sample standard deviation of the ten paired differences; the range and win count describe those same ten differences. They are not confidence intervals or tests of a population-level superiority claim.

model	radiance [%]	median [pp]	95th percentile [pp]	RMSE [pp]
Cubic ridge	0.760 $\pm$ 0.011	0.471 $\pm$ 0.010	70.68 $\pm$ 0.02	4.18e+03
Isotropic KRR	0.239 $\pm$ 0.016	0.141 $\pm$ 0.018	69.53 $\pm$ 0.42	4.18e+03
Input-scaled KRR	0.095 $\pm$ 0.009	0.036 $\pm$ 0.004	31.16 $\pm$ 6.21	1.34e+04
Narrow feature kernel	0.098 $\pm$ 0.007	0.054 $\pm$ 0.004	7.75 $\pm$ 1.13	1.78e+03
Narrow convex stack	0.087 $\pm$ 0.006	0.048 $\pm$ 0.003	29.44 $\pm$ 5.89	1.25e+03
Wide feature kernel	0.079 $\pm$ 0.009	0.045 $\pm$ 0.006	7.57 $\pm$ 0.95	1.40e+03
Wide convex stack	0.078 $\pm$ 0.011	0.043 $\pm$ 0.005	22.71 $\pm$ 4.42	5.81e+03

Table 3: All-band metrics from the complete ten-seed main and wide campaign records. Radiance is mean relative L^2 error. Reflectance median, 95th-percentile absolute error and RMSE are in percentage points at $\rho = 0.7$. The first three columns show mean $\pm$ sample SD across splits; RMSE shows its mean only, with its spread retained in the machine-readable summary. The legacy inverse is neither physically clipped nor conditioned on an application passband. All inverses in these records are finite, so their zero-fill branch does not explain the large errors. Large finite values near denominator zeros dominate RMSE. These are diagnostic all-band errors, not operational retrieval accuracy.

percentage points. By contrast, the wide stack’s improvement over its own feature kernel is only 0.0006 ± 0.0026 points, with eight wins and two losses. The latter is a small mean improvement within the repeated-split variability, not compelling evidence of a practically superior predictor. These are descriptive paired summaries: repeated splits of one finite table, overlapping test rows and adaptive study development do not supply ten independent external validation experiments.

Radiance and reflectance also order the heads differently. Table 3 reports the already stored all-band reflectance tails alongside the medians. Each entry averages a within-split statistic; the mean of ten 95th percentiles is not a pooled 95th percentile. For the width-512 pipeline, the convex stack lowers radiance error relative to the feature kernel, but raises the average 95th-percentile reflectance error from 7.75 to 29.44 percentage points. The tail is worse at all ten splits, by factors between 2.63 and 4.94. The width-2000 comparison repeats the reversal: the feature kernel has 7.57 points at the 95th percentile and the stack 22.71, with the stack worse at all ten splits by factors between 2.36 and 3.79. Thus the near-tie at 0.079% versus 0.078% radiance error conceals a substantial difference in the inverse metric.

training rows	seeds	Cubic ridge	KRR, isotropic	KRR, per-input	FC-DNN	DNN + residual KRR	kernel on features
500	10	0.846±0.031	0.639±0.022	0.415±0.024	1.961±0.057	0.815±0.035	0.456±0.020
1,000	10	0.802±0.024	0.519±0.034	0.260±0.066	1.862±0.015	0.590±0.027	0.316±0.040
2,000	10	0.780±0.016	0.424±0.051	0.217±0.052	1.232±0.034	0.376±0.041	0.240±0.049
4,000	10	0.769±0.010	0.328±0.028	0.154±0.037	0.830±0.019	0.256±0.021	0.164±0.033
8,000	10	0.763±0.010	0.279±0.020	0.119±0.027	0.584±0.020	0.190±0.014	0.121±0.017
18,884	10	0.760±0.011	0.239±0.016	0.095±0.009	0.387±0.012	0.141±0.006	0.098±0.007

Table 4: Learning curves: radiance relative L^2 error [%] against the number of training rows, mean $\pm$ standard deviation over the ten splits reported in the archived aggregate. At $n = 8,000$, one individual record per kernel campaign is absent from the public archive; the aggregate values are preserved rather than recomputed from nine records. Each rung uses the first n rows of the seed’s training block; the validation and test blocks are those of Table 1, so the top rung is that table. The per-input kernel repeats its coordinate search at every rung. Component columns are in Figure 3 and in the per-seed records.

The input-scaled kernel illustrates the same distinction: its median error is only 0.0365 points but its mean 95th percentile is 31.16 points. Conversely, the feature kernels have slightly larger medians and appreciably smaller tails. No head should therefore be called best without naming the metric, width, candidate pool and spectral domain. The identities of Section 2.1 explain why a radiance or componentwise convex optimization need not control the inverse: the predicted denominator must remain separated from zero. They do not attribute the observed tail difference to any unmeasured component-error correlation. The result motivates a separately specified retrieval objective and conditioning analysis; it does not retroactively validate a mask chosen to improve these test scores.

All entries in these two tables are regenerated by `code/make_revision_tables.py` from the twenty individual records, with no model retraining and no access to raw EMIT arrays. The reanalysis includes checks of the recorded data-file digests, configuration checks and equality of the paired deterministic baseline rows. These checks establish consistency of the archived summaries, not independent reproduction of the original simulator or network training.

5 Learning curves, output rank and network width

Table 1 is one cell of a larger design: one training size, one output rank, one network width. This section varies each in turn, at the same ten seeds, with the same validation-based selection and one test read per family, for the isotropic exact kernel, the network, its residual correction and the kernel on its features; the cubic ridge is carried as the smoothness reference. The per-input kernel and the ensembles are run at full size only.

Learning curves. Each rung uses the first n rows of the seed’s training block, $n = 500, 1,000, 2,000, 4,000, 8,000$ and the full 18,884, with the validation and test blocks unchanged, so the rungs of one seed are nested and the full block is the top rung (Table 4, Figure 3). Three things are visible. First, the cubic is flat: from 0.85% to 0.76% over a factor of 38 in data, a bias floor reached by a thousand rows. Second, on the three smooth components every other family decays at close to $n^{-1/2}$ (Table 5): the least-squares slopes of log error against log n over the six rungs lie between -0.38 and -0.61 for the kernel, the network, the corrected network and the feature kernel on Y_1 , Y_2 and Y_3 , while on the albedo every slope lies between -0.09 and $+0.01$: the




Figure 3: Learning curves on log-log axes for the radiance and for three of the four components, mean and standard deviation over ten splits. The cubic ridge is flat from the first rung; the kernel, the network, the corrected network and the feature kernel fall at close to $n^{-1/2}$ on the smooth components and not at all on the spherical albedo. The albedo’s spread at 1,000 and 2,000 rows comes from single splits whose albedo fit collapses, the failure Section 4 describes for a kernel capped at a random subset; the medians across splits are flat.

family	slope over the six rungs					slope of the last step				
	Y_1	Y_2	Y_3	Y_4	radiance	Y_1	Y_2	Y_3	Y_4	radiance
Cubic ridge	-0.03	-0.03	-0.03	-0.02	-0.03	-0.00	-0.01	-0.00	-0.08	-0.01
KRR, isotropic	-0.55	-0.38	-0.40	-0.01	-0.28	-0.38	-0.19	-0.17	0.10	-0.18
KRR, per-input	-0.45	-0.43	-0.43	0.01	-0.40	-0.32	-0.26	-0.29	0.12	-0.26
FC-DNN	-0.50	-0.49	-0.57	-0.09	-0.48	-0.59	-0.60	-0.55	0.07	-0.48
DNN + residual KRR	-0.52	-0.53	-0.61	-0.04	-0.50	-0.39	-0.33	-0.39	0.12	-0.35
Kernel on features	-0.45	-0.44	-0.46	-0.03	-0.43	-0.28	-0.22	-0.31	0.08	-0.25

Table 5: Learning-curve exponents: the least-squares slope of log error against log n over the six rungs of Table 4, and the slope of the last step alone, from 8,000 to 18,884 rows, per family and component.

albedo shows no sustained decrease over this budget range. This plateau does not distinguish approximation bias, optimization, relative-error weighting or limitations of the training data.

These are slopes of six mean errors over a finite sample-size range, not established asymptotic rates. Standard Sobolev or kernel regression rate formulas require sampling, source and regularization assumptions [16, 18]; neither those assumptions nor the smoothness norm of the fitted network residual is verified by this experiment. Third, the kernel’s curve bends at the top. Its slope over the first five rungs is -0.59 , -0.44 and -0.47 on the smooth components, and its last step, from 8,000 to 18,884 rows, is -0.38 , -0.19 and -0.17 , while the network’s last step is as steep as its average (-0.59 , -0.60 and -0.55 ; Table 5). Part of this is the tuning protocol: the kernel’s smoothness, scale and nugget are chosen on a 6,000-row subsample at every rung, so above 6,000 rows the hyperparameters are those of a smaller problem, and Section 5.1 measures what tuning with every row in the solve recovers: all of it on the radiance, most of it on the components. The consequence for the comparison is that the ratio of the network’s radiance error to the kernel’s falls from 3.1 at 500 rows to 1.6 at the full block, or to 1.8 with the kernel tuned on every row. The kernel’s margin over the network on this table is a margin at this training size; whether it would survive ten times the data depends on where the kernel’s curve settles, which the bend leaves open.

Fourth, and this is the one a single training size cannot show: against the isotropic kernel

PCA rank	variance outside the basis	KRR, isotropic	FC-DNN	DNN + residual KRR	kernel on features
16	7.5×10^{-4}	0.270 $\pm$ 0.015	0.390 $\pm$ 0.015	0.175 $\pm$ 0.007	0.148 $\pm$ 0.009
32	1.5×10^{-5}	0.243 $\pm$ 0.018	0.390 $\pm$ 0.015	0.139 $\pm$ 0.006	0.102 $\pm$ 0.009
64	8.2×10^{-8}	0.242 $\pm$ 0.018	0.390 $\pm$ 0.015	0.143 $\pm$ 0.006	0.100 $\pm$ 0.009
128	5.4×10^{-9}	0.242 $\pm$ 0.018	0.390 $\pm$ 0.015	0.150 $\pm$ 0.005	0.100 $\pm$ 0.009
285	0	0.242 $\pm$ 0.018	0.390 $\pm$ 0.015	0.172 $\pm$ 0.004	0.100 $\pm$ 0.009

Table 6: Output-rank ablation: radiance relative L^2 error [%] when the kernels and corrections regress the leading r principal coefficients, mean $\pm$ standard deviation over five splits. The second column is the fraction of the standardized variance the basis leaves out for the slowest-decaying component, the spherical albedo; at rank 285 it is zero. The network regresses the 285 bands directly and is the same predictor in every row.

the per-input metric changes the measured finite-range slope as well as the error level. Its coordinate search, repeated at every rung, gives a slope of -0.399 on the radiance against the isotropic kernel’s -0.279 , and the ratio between them falls from 0.65 at 500 rows to 0.40 at the full block, with one step against that direction (0.50 at 1,000 rows, 0.51 at 2,000) and a fall at every other. The observed discount increases over this sample-size range. The selected metric, smoothness and nugget all change with sample size, so these measurements do not isolate a change of asymptotic dimension or establish a faster universal rate. Against the kernel on network features the margin closes as the data grows. The per-input metric has the lower mean at all six rungs, but the split-by-split count runs 10, 9, 10, 7, 5 and 6 wins of ten from 500 rows to the full block, and the gap itself narrows from 0.415 against 0.456 at 500 rows to 0.095 against 0.098 at the top; the feature kernel’s own slope, -0.435 , is the steeper of the two. A metric searched on six validated numbers is the better buy on small data, and a representation learned from the data catches it up. On the albedo the per-input slope is $+0.01$: the one component no metric rescues.

The corrected network is below both of its parents on the radiance at every one of the ten splits from 2,000 rows on, and at no split at 500 or 1,000 rows, where its error sits between the network’s and the kernel’s. On Y_2 and Y_3 it is below both at every split at the two largest sizes; on Y_1 it never is, the isotropic kernel alone being the better head at every size (0.084% against 0.099% at the full block). Its exponent is close to the network’s and its error level is closer to the kernel’s in these runs. Section 6 gives the exact residual decomposition, while distinguishing it from a proof of these empirical slopes or of the component-specific exception.

Output rank. The kernels and corrections regress the leading r principal coefficients of the standardized component; the network regresses the 285 bands and is the same predictor at every rank. Table 6 refits at $r = 16, 32, 128$ and 285 at five of the seeds beside the rank-64 rows of those seeds. The kernel is indifferent to the rank above 32: its radiance error reads 0.243, 0.242, 0.242 and 0.242 at 32, 64, 128 and 285, and 0.270 at 16, where the truncation leaves 7.5×10^{-4} of the albedo’s variance outside the basis. The kernel on features behaves the same way. The residual correction does not: it is best at 32 (0.139), a little worse at 64 (0.143), and worse again at 128 and 285 (0.150 and 0.172). The difference is in what the two regressions fit. The kernel regresses the component itself, whose tail coefficients are small and smooth functions of the state, and including them changes these measured errors little; it still has a computational cost. The correction regresses the network’s residual, whose tail coefficients are the network’s own errors along directions that hold 10^{-6} of the variance, and a kernel regression of those directions

network	head	component rel. L^2 [%]				radiance	reflectance
		Y_1	Y_2	Y_3	Y_4	rel. L^2 [%]	med. $ \hat{\rho} - \rho $ [pp]
3×512	network	0.48 ± 0.02	0.97 ± 0.04	0.41 ± 0.02	3.02 ± 0.18	0.387 ± 0.012	0.245 ± 0.008
	network + residual KRR	0.10 ± 0.00	0.30 ± 0.02	0.16 ± 0.01	2.43 ± 0.08	0.141 ± 0.006	0.089 ± 0.008
	kernel on its features	0.18 ± 0.01	0.17 ± 0.01	0.10 ± 0.01	1.99 ± 0.08	0.098 ± 0.007	0.054 ± 0.004
	per-coordinate selection	0.19 ± 0.01	0.44 ± 0.10	0.13 ± 0.02	2.25 ± 0.18	0.193 ± 0.022	0.123 ± 0.023
	convex stack	0.07 ± 0.00	0.16 ± 0.01	0.09 ± 0.00	1.94 ± 0.06	0.087 ± 0.006	0.048 ± 0.003
3×2000	network	0.93 ± 0.20	2.11 ± 0.49	0.76 ± 0.19	3.27 ± 0.19	0.748 ± 0.116	0.508 ± 0.110
	network + residual KRR	0.12 ± 0.01	0.35 ± 0.03	0.19 ± 0.03	2.35 ± 0.13	0.160 ± 0.016	0.106 ± 0.010
	kernel on its features	0.13 ± 0.02	0.13 ± 0.02	0.08 ± 0.01	2.10 ± 0.34	0.079 ± 0.009	0.045 ± 0.006
	per-coordinate selection	0.14 ± 0.02	0.30 ± 0.26	0.09 ± 0.01	2.55 ± 0.26	0.187 ± 0.057	0.099 ± 0.033
	convex stack	0.08 ± 0.01	0.13 ± 0.02	0.08 ± 0.00	2.00 ± 0.24	0.078 ± 0.011	0.043 ± 0.005

Table 7: The pipeline of Section 3 run with a 3×2000 network in place of the 3×512 one, at the same ten splits: the network (one member, up to 500 epochs under the same optimizer, learning rate and early stopping), its residual correction, the exact kernel on its last-layer features, and the per-coordinate selection and convex stack over the wide pipeline’s heads (the isotropic kernel, the network, its correction and its feature kernel), beside the corresponding rows of Table 1; mean $\pm$ standard deviation over the ten splits.

adds a term the 2,098 validation rows select but the test block does not reward. A rank at the point where the retained variance stops moving, 32 to 64 on this table, is the right setting for the correction, and the rank-64 setting of Table 1 is on the safe side of it.

A wider network. A network of three hidden layers of 2,000 units, trained for up to 500 epochs under the same optimizer, learning rate and early stopping as the 512-unit network, was run at all ten seeds with the rest of the pipeline built on it (Table 7). It is a worse emulator at every seed, 0.75% against 0.39% on the radiance, and its early stopping fired between 42 and 210 epochs of the 500 allowed: at this learning rate and batch size this particular wider-network training configuration gives worse standalone predictions. Width and epoch budget change together; these results do not rule out a better-optimized wide network. Its representation is better. The exact kernel on its last-layer features reaches 0.079% on the radiance against 0.098% on the 512-unit features, below the per-input kernel’s 0.095% at all ten seeds and below the convex stack of Table 1 at nine of them, and on the two flux components, at 0.13% and 0.08%, it is at or below every row of that table. The convex stack over the wide pipeline’s heads, which puts 0.88 to 0.96 of its weight on that feature kernel, has the lowest mean radiance error in the ten-split comparison at 0.078% on the radiance and 0.043 reflectance points, below the stack of Table 1 at nine of the ten seeds; the per-coordinate selection over the same heads does not transfer here either. The residual correction of the wider network is a little worse than that of the narrower one, although that ordering is not guaranteed merely because its base network is worse (Section 6). A 3×1000 rung at the same ten seeds sits between the two: the network at 0.72%, the kernel on its features at 0.088% against 0.098% at width 512 and 0.079% at width 2,000. From 512 to 2,000 the network’s own error rises with width while the error of a kernel on its features falls: what a wider network learns on this table is a better coordinate system, not a better map. At 4,000 units, run at two seeds, the network is worse again at 0.95% and the kernel on its features turns back up to 0.097%, so the gain in the representation does not continue past 2,000.

If that reading is right, several networks should give a better coordinate system than one.

kernel	component rel. L^2 [%]				radiance	reflectance
	Y_1	Y_2	Y_3	Y_4	rel. L^2 [%]	med. $ \hat{\rho} - \rho $ [%]
isotropic, tuned on 6,000 rows	0.08±0.01	0.42±0.03	0.20±0.00	2.62±0.18	0.239±0.016	0.141±0.018
isotropic, tuned on every row	0.08±0.00	0.36±0.02	0.17±0.01	2.58±0.14	0.210±0.008	0.107±0.006
per-input search, tuned on 6,000 rows	0.08±0.01	0.16±0.02	0.08±0.01	2.18±0.16	0.095±0.009	0.036±0.004
per-input search, re-tuned on every row	0.08±0.00	0.16±0.02	0.08±0.01	2.09±0.12	0.096±0.009	0.036±0.004
per-input from the cubic's gradients, $p = 1$	0.34±0.01	0.19±0.02	0.14±0.01	2.17±0.14	0.159±0.007	0.078±0.004
per-input from the cubic's gradients, $p = 1/2$	0.11±0.00	0.19±0.02	0.09±0.01	2.27±0.15	0.125±0.008	0.055±0.004
per-input from the cubic's gradients, $p = 1/4$	0.08±0.00	0.25±0.02	0.14±0.01	2.38±0.20	0.157±0.015	0.077±0.013

Table 8: Where the kernel’s hyperparameters are chosen, and how its metric is obtained, at the full training block and the same ten splits. The first pair tunes smoothness, scale and nugget on a 6,000-row subsample and on every training row; the last three rows replace the coordinate search by a metric read off the cubic surrogate’s gradients, at three exponents. Selecting the exponent on validation, per seed, chooses $p = 1/2$ at all ten, so that row would duplicate the $p = 1/2$ row and is not printed; the per-seed choices are in the stored numbers. Mean $\pm$ standard deviation over the ten splits.

A three-member version of the same wide pipeline, run at three seeds, lets the exact kernel see the concatenated last-layer features of all three members at once, 6,000 coordinates rather than 2,000. It reads 0.074% on the radiance against 0.083% for the kernel on a single member’s features in the same lanes, lower at all three seeds, and 0.043 against 0.051 reflectance points, again at all three. That puts one exact solve level on the radiance with the convex stack of six heads fitted in those same lanes (0.074% against 0.074% on the mean, though the stack is the lower at two of the three seeds and on the reflectance median at all three). A combination of six predictors and a single kernel in a wider feature space buy about the same thing here, which is consistent with a useful role for representation diversity. The comparison does not isolate all optimization and kernel-selection effects. Three seeds provide exploratory evidence rather than a ten-split replication, and the construction is not run at the width of Table 1, so whether the same gain survives there is left open.

5.1 Tuning with every row in the solve, and a metric read off the data

Two further constructions test the two readings above: that the bend at the top of the kernel’s curve is its tuning protocol, and that the metric the coordinate search finds is the physics of the components. Both were run at all ten seeds (Table 8).

Tuning with every row in the solve. The kernel families of Section 3 choose smoothness, scale and nugget by the validation error of a solve on a 6,000-row subsample, and refit the winner on all 18,884 rows. Doing the tuning solve on every row instead lowers the isotropic kernel’s radiance error from 0.239% to 0.210% and its reflectance median from 0.141 to 0.107 points, better at nine of the ten splits, the ratio averaging 0.882 and never above 1.001. The mechanism is visible in the chosen scales: over the forty component-seed cells the full block picks a shorter length scale than the subsample in 26, the same in 12, and a longer one in 2. A tuning set a third of the size of the solve supports a longer length scale than the solve does, so the subsample protocol over-smooths, and it over-smooths more the larger the gap between the two.

That is the bend. Below 6,000 rows the tuning subsample *is* the whole training block and the rungs are already full-block tuned; at 8,000 rows it is three quarters of it; only the top rung is a third. Applying the measured ratio to the top rung of Table 4 takes the kernel’s last step

on the radiance from -0.183 to -0.334 , against a slope of -0.305 over the five lower rungs: the top-step slope becomes closer to that of the lower rungs. One retuned endpoint does not establish a power law over or beyond the observed range. On the two flux components the same correction takes the last step from -0.187 and -0.171 to -0.352 and -0.358 against averages of -0.437 and -0.468 , so there most of the bend is the protocol and a little of it is not. The consequence for the comparison of Section 4 is that the kernel’s margin over the network at the full block is wider than Table 4 shows: 1.8 times rather than 1.6, under this retuned endpoint. No conclusion about unobserved sample sizes follows.

The same re-tuning does nothing for the per-input kernel: 0.096% against 0.095%, better at four of the ten splits, ratio 1.005. Its coordinate search already selects a scale of once the median distance at the metric it finds, which is what the full-block tune selects too. Rescaling the coordinates removes the over-smoothing that the isotropic kernel suffers, so the per-input kernel of Table 1 was never paying the tuning cost, and its number stands as reported.

A metric read off the data. The coordinate search costs a sweep over the six inputs at five multipliers, twice, per component and per seed. The alternative sets the metric before any kernel is fitted: the multiplier of coordinate j is the root-mean-square gradient of the fitted cubic surrogate along j over the training rows, raised to a power p , normalized to geometric mean one and clipped to the search’s own range $[1/16, 4]$, so a coordinate the map barely responds to gets a long length scale and p sets how hard the metric leans on that reading. It costs one cubic fit and six finite differences.

Taken proportionally, at $p = 1$, it captures rather more than half the gap: on the radiance it reads 0.159%, below the isotropic kernel at every one of the ten splits by a factor 0.669 on average and above the searched metric at every one of them by a factor 1.680. But p is a free number, and the ladder is not monotone. The square root reads 0.125% and the fourth root 0.157%, so the best exponent is interior, and choosing it the way every other hyperparameter in this study is chosen — the lowest mean validation error over the four components, per seed, with the test block untouched — selects $p = 1/2$ at all ten seeds. At that exponent the metric closes 79% of the distance from the isotropic kernel to the searched one on the radiance and 82% on the reflectance median, for one cubic fit, six finite differences and a choice among three numbers.

The per-band reading says why the exponent matters and what the search still buys. The four components do not want the same one: Y_2 and Y_3 are best at $1/2$, the albedo at 1 (2.18% against 2.27%), and the path radiance at $1/4$. Y_1 is where the damage is, and it is worth seeing why. On that component the isotropic kernel already reads 0.084% and the searched metric 0.085%: there is no anisotropy to find. The proportional metric nonetheless reads 0.34% there, four times worse than doing nothing, because the surrogate’s gradient along a coordinate measures how fast the map moves with it and not how much predictive information it carries, and on the one component that depends on the viewing geometry those two disagree. Softening the exponent to $1/4$ takes Y_1 back to 0.085%, exactly the searched metric’s value, at the cost of Y_2 and Y_3 . So a metric read from first-order sensitivity, with its exponent validated like any other hyperparameter, is most of what the coordinate search buys at a small fraction of its cost. What the search still buys is the freedom to be anisotropic differently for each component, which one global exponent cannot be.

6 Residual correction: identities, comparisons and limits

The learning curves motivate a distinction between an identity for a fitted predictor and a claim about its out-of-sample rate. The former can be established without distributional assumptions; the latter requires assumptions not measured by these experiments. We first fix all fitted quantities, including any data-dependent kernel hyperparameters, and work with one output coordinate. The identities extend coordinatewise to vector outputs.

6.1 A frozen regression operator

Fix a design $X_n = (x_1, \dots, x_n)$, a positive-semidefinite kernel with Gram matrix K , and $\lambda > 0$. Define

$$(Sg)(x) = k(x, X_n)(K + \lambda I)^{-1}g(X_n), \quad R = I - S. \quad (8)$$

Here λ is the diagonal regularizer in the linear system. The campaign code uses $\lambda = n\gamma$, where γ is its reported nugget; the distinction matters when the training size changes. Although a selected kernel can depend on the targets, S is linear in the response only *after that selection is frozen*.

Proposition 2 (Frozen residual correction). *Let f be the target and h any fitted mean, and set $\hat{f}_c = h + S(f - h)$. Then:*

$$f - \hat{f}_c = R(f - h), \quad (9)$$

$$(Rg)(X_n) = \lambda(K + \lambda I)^{-1}g(X_n). \quad (10)$$

If $Ku_j = \mu_j u_j$, the design residual in direction u_j is multiplied by $\lambda/(\mu_j + \lambda) \in (0, 1]$. The factor is one when $\mu_j = 0$ and strictly below one when $\mu_j > 0$; this is shrinkage, not a hard spectral projection.

Suppose, in addition, that $\mathcal{S}$ is a normed vector space containing f and h and that the following bound holds uniformly on that space for the realized operator:

$$\|Rg\|_{L^2(P)} \leq C_n \|g\|_{\mathcal{S}} \quad (g \in \mathcal{S}).$$

Then

$$\|f - \hat{f}_c\|_{L^2(P)} \leq C_n \|f - h\|_{\mathcal{S}}. \quad (11)$$

Consequently, $\|f - h\|_{\mathcal{S}} < \|f\|_{\mathcal{S}}$ makes this displayed upper bound smaller than $C_n \|f\|_{\mathcal{S}}$ when $C_n > 0$. It is neither a necessary nor a sufficient condition for the corrected predictor's actual error to be smaller than that of Sf .

Proof. Subtract $h + S(f - h)$ from f to obtain (9). On the design, $I - K(K + \lambda I)^{-1} = \lambda(K + \lambda I)^{-1}$, which proves (10) and its spectral form. Apply the assumed uniform inequality to $g = f - h$ for (11). The final distinction is demonstrated below. $\square$

Why the bounds do not order the errors. On a two-point design take $K = \text{diag}(1, 99)$ and $\lambda = 1$, so $R = \text{diag}(1/2, 1/100)$ and $C_n = 1/2$ is its Euclidean operator norm. With $f = (1, 0)$ and $f - h = (0, 2)$, the residual norm is twice the target norm, yet $\|R(f - h)\| = 0.02 < 0.50 = \|Rf\|$. Conversely, with $f = (0, 1)$ and $f - h = (1/2, 0)$, the residual norm is smaller, but $\|R(f - h)\| = 0.25 > 0.01 = \|Rf\|$. These examples also show why a more accurate neural mean need not yield a more accurate corrected predictor: the orientation of the remaining error matters.

6.2 An exact criterion for improvement

The relevant comparison can be made exactly in any Hilbert error metric. This includes a finite weighted squared error and $L^2(P)$, but not directly the mean of per-sample relative norms or the nonlinear radiance and reflectance metrics in Table 1.

Proposition 3 (Error alignment). *For the same frozen operator S , put $a = Rf$ and $b = Rh$. Then*

$$\|f - \hat{f}_c\|^2 - \|f - Sf\|^2 = \|b\|^2 - 2\langle a, b \rangle. \quad (12)$$

Thus full correction improves on Sf exactly when $2\langle a, b \rangle > \|b\|^2$. More generally, the interpolation $p_t = Sf + tRh$, $0 \leq t \leq 1$, has minimum squared error at

$$t_* = \min\{1, \max\{0, \langle a, b \rangle / \|b\|^2\}\} \quad \text{if } b \neq 0; \quad (13)$$

if $b = 0$, every t is equivalent. If the direct fit uses S_0 and the correction uses S_1 , (12) instead holds with $a = (I - S_0)f$ and $b = (S_1 - S_0)f + (I - S_1)h$.

Proof. For a common operator the corrected error is $a - b$, so expansion of its squared norm gives (12). The error of p_t is $a - tb$; minimize the resulting quadratic on $[0, 1]$ to obtain (13). For distinct operators, $f - h - S_1(f - h) = (I - S_0)f - [(S_1 - S_0)f + (I - S_1)h]$, after which the same expansion applies. $\square$

The criterion is a diagnostic, not an oracle available at prediction time. Replacing its inner products by validation errors gives a validation-optimal interpolation, not a guarantee on new data. The reported experiments use separately tuned direct and residual kernels, so their comparison must use the final, two-operator form rather than silently identifying the operators. Neither the criterion nor the training-design contraction establishes that correction will beat both parents at every sample size.

6.3 What the principal-component implementation computes

The campaign trains a network on all standardized bands but projects its predictions into the retained principal-component basis before correction. Thus its corrected predictor is not, literally, the unprojected network plus a kernel residual. To state the implementation precisely, let f and h now denote the target and network output after the common fitted standardization and PCA centering. Let Π be the orthogonal projection onto the retained spectral coordinates. In these centered coordinates, the implemented prediction is

$$\hat{f}_{c,\Pi} = \Pi h + S\Pi(f - h).$$

The output mean and diagonal band scaling are restored afterwards. Since the same scalar kernel operator acts on each retained output coordinate,

$$f - \hat{f}_{c,\Pi} = (I - \Pi)f + (I - S)\Pi(f - h). \quad (14)$$

The terms on the right are orthogonal in the standardized spectral Euclidean norm at each query, hence their squared norms add there. They need not be orthogonal in physical units after unequal band rescaling. This decomposition separates truncation error from residual regression and explains why output rank is a genuine part of the model specification. It does not by itself prove the observed optimum near rank 32. Comparisons with the unprojected network include the effect of projection as well as residual fitting.

Rates and interpretation. Standard kernel learning-rate results impose source, spectrum, sampling and regularization conditions [16, 17, 18]. If a bound uniform over the relevant residual class holds with $C_n = O(n^{-\alpha})$ and one *separately establishes* $\|f - h_n\|_{\mathcal{S}} = O(n^{-\beta})$, then (11) gives an upper bound $O(n^{-(\alpha+\beta)})$. It gives neither equality of rates nor a lower bound. Empirical network loss does not establish convergence in $\mathcal{S}$, and a bound for a fixed target cannot automatically be applied to an adaptively fitted residual or a selected kernel. The required smoothness norms and uniform events are not measured here. Accordingly the slopes in Table 5 are finite-range empirical slopes, not estimates of a proved asymptotic exponent. Smoother or better-aligned residuals are possible interpretations of the successful corrections; they are not conclusions forced by the algebra.

The tests in `code/test_publication_revision.py` check the residual identities, regularizer convention, semidefinite boundary case, alignment formula, counterexamples and projected implementation on explicit and randomized finite designs. These numerical checks support the implementation of the displayed formulas; the proofs above supply their mathematical justification.

7 Where the constructions do not transfer: OCO-2

The companion paper [6] applies the same construction to the OCO-2 radiative-transfer emulation problem of Lamminpää et al. [5]: per spectral band (O_2 , weak CO_2 , strong CO_2), a reduced atmospheric state of twenty or twenty-four coordinates maps to forty reduced radiance coefficients, with errors on the standardized coefficients and on the reconstructed radiance, under the ten-split protocol of Section 2. That paper carries the tables; the findings that are the counterpart of the EMIT results are stated here so the two problems can be read against each other.

The problem sits at the opposite end of the structural spectrum. The input is higher-dimensional and the map rougher, and a kernel on the state trails the network by a factor of nine, four and five on the reduced coefficients: the isotropic Matérn reaches 40.6%, 58.6% and 43.4% on the three bands against 4.4%, 16.4% and 8.2% for the network. The residual correction adds no material gain in the reported OCO-2 comparisons. Little alignment with useful kernel directions is a possible explanation, but the public records do not measure the residual’s Gram-eigenvector coefficients. The design identity in Proposition 2 alone cannot establish that explanation on test data. The coordinate search that gives a kernel on the state its own length scales, a strong raw-input construction on the EMIT table, is not the strongest there either. Four learned metrics were run at the same ten splits on all three bands: a kernel-flow metric [27], a Mahalanobis metric under the same flow, the gradient outer product of the recursive feature machine [15], and per-input scales fitted by empirical Bayes. The last is the best of the four at every band, taking the isotropic kernel’s 40.6%, 58.6% and 44.3% on the coefficients to 18.7%, 25.3% and 16.9%, and on O_2 its radiance from 0.233% to 0.150%; the two flow metrics agree within half a point of each other at 26.9%, 29.8% and 23.9%, the recursive feature machine reads 28.2%, 47.9% and 31.3%, and an additive kernel 38.2%, 46.9% and 47.7%. The isotropic baseline agrees between these lanes and the protocol lanes on O_2 and weak CO_2 to two decimals; on strong CO_2 it reads 44.3% here and 43.4% there, the one band where they part, and also the one whose isotropic column carries a spread of ± 3.4 because two of its ten splits land three points below the other eight. So a learned metric does close a good part of the distance on the state, cutting the isotropic error to between two fifths and a half of itself

depending on the band. What it does not reach is the kernel inside the representation, and by how much is a property of the band rather than a constant: the exact Matérn regression on the network’s frozen features reads 4.1%, 16.3% and 8.1% at the same splits, so the best state metric is 4.6 times that on O₂ but only 1.6 and 2.1 times it on the two CO₂ bands. One caution on the best row: on O₂ the empirical-Bayes metric is bimodal across splits, eight of the ten landing near 17.2% and two at 24.7%, so its spread of ± 3.2 is thirty times the isotropic kernel’s ± 0.1 and its mean is not the value a single split is likely to show. The feature-kernel construction remains useful on both tasks: the exact Matérn regression on the network’s frozen features beats a ridge readout refitted on the same features at every split, on every band and under every training loss, and a per-coefficient choice among heads trained under different losses has lower mean error than the release’s stored emulator under both reported metrics. In the archived loss campaign it wins the reduced-coefficient comparison at all ten splits on each band and the radiance comparison at ten, nine and ten splits on O₂, weak CO₂ and strong CO₂, respectively. These companion comparisons are not an independent EMIT validation; their tables are in [6].

One observation about the release itself is recorded here because it does not depend on the protocol. The reference’s data release carries 976 retrieval records, each with an atmospheric state, the full-physics monochromatic radiance at that state, its rank-40 reconstruction, and the prediction of the release’s Gaussian-process emulator at the same forty coefficients. Every one of the 976 reduced states matches a stored test input within 6.67×10^{-16} ; the records map to 974 distinct test rows, and the nearest test-row index is the same in all three bands. The records are therefore not disjoint from the stored test set and cannot serve as an independent evaluation pool, though they do provide a different target, the full-physics radiance, at those states. The state reduction was replayed independently from the release’s two dimensional-reduction files and the `x_in` and `geo_in` fields of all 976 records; the replay code and its digests are archived with the per-run records.

8 Ten further emulation configurations

The two problems above stand at opposite ends of one axis: a smooth, six-dimensional, low-rank table where input-scaled and learned-feature exact kernels are strong predictors, and a rougher, twenty-dimensional one where the kernel is useful only inside a learned representation. If that axis orders the comparison, it should order other emulation problems too. Table 9 applies the same family set to ten further configurations, with no tuning per corpus beyond what each family selects on validation.

The corpora are the advection equation with discontinuous initial data from the operator suite of [19], which is also the source of the structural-mechanics problem of [6]; Burgers’ equation at three viscosities from PDEBench [20]; three one-step maps from The Well [21] — a two-dimensional turbulent radiative layer, an active-matter system, and a Helmholtz staircase; the ClimSim atmospheric physics emulation task [22]; and the paired correction-coefficient corpus of the physics-guided radiative-transfer work of [4], which is a third radiative-transfer emulation problem and the largest here at 487,790 rows. That last corpus is run in both of its configurations: from the atmospheric state alone, and with the low-fidelity 6S coefficients supplied as additional inputs, which is a different and easier problem.

The configurations are these. Advection is the 200-point input/output pair of the operator suite, used at full resolution with no reduction. Burgers is PDEBench’s `1D_Burgers_Sols_Nu/` release, the map from the initial condition to the solution at $t = 1$ on a grid subsampled by four

to 256 cells, the input used raw and the output reduced to 128 principal coefficients. The three Well datasets are one-step maps of the full multi-channel state with 256 principal coefficients in and out: the turbulent radiative layer on its 128×384 grid (density, pressure, velocity), active matter on 256×256 (concentration, velocity, and the two tensor fields), and the Helmholtz staircase on its two pressure channels. ClimSim is the low-resolution release, 124 inputs to 128 tendencies, with 100,000 training rows drawn from the 10.1M available and validation and test taken from the release’s own validation and scoring files. The correction-coefficient corpus is the paired 6S/libRadtran Sentinel-2 set, nine state variables and the band wavelength mapping to the three libRadtran coefficients, split by atmospheric state so that all bands of a state travel together; its multi-fidelity configuration adds the 6S coefficients of the same state as inputs.

The families are those of Section 3 with two differences forced by scale. The reference family is a ridge regression linear in whatever inputs the models receive, not the cubic of Table 1: a cubic expansion of a 256-dimensional input is not tractable, and on the two corpora whose inputs are six- and nine-dimensional the linear form is the one that was run. And the kernel solves are exact only when the training block has at most 20,000 rows; above that they use a Nyström approximation on 6,000 uniformly drawn landmarks. That cap binds on ClimSim, on the Helmholtz staircase at 20,384 rows, and on the two correction-coefficient configurations. Those entries are what a capped fit gives; they are not the exact solve, and we do not claim them as a bound on it in either direction, since restricting the fit set both discards information and regularizes. The convex stack and the per-coordinate selection are formed from the kernels, the networks, the residual correction and the feature kernel; the ridge is a reference row and is not one of their members.

What each row averages differs and is stated in the table. Advection, Burgers, ClimSim and the two correction-coefficient configurations have three to five random splits. The Well’s releases carry their own train, validation and test files, which are used unchanged, so the three turbulent-radiative-layer runs share one split and differ only in network initialisation; their spread, of order 0.01 points, measures initialisation and not sampling. The active-matter and Helmholtz-staircase rows are single runs.

Three things in the table repeat what the two main problems show. First, the ordering between families follows the structure rather than the corpus. Where the map is smooth and the sample small — Burgers at $\nu = 0.1$ — the exact kernel is at 2.81% against the network’s 20.34%, a factor of seven, of the same kind as the five- to six-fold gap between the per-input kernel and the network on the EMIT table’s smooth components; as the viscosity falls and the solutions roughen the gap closes to 23.40% against 36.65% and then to 31.66% against 39.98%. Where the sample is large and the map rough — ClimSim, and the correction-coefficient corpus at nearly half a million rows — the network leads instead, 25.10% against 27.20% and 3.75% against 62.93%, which is the OCO-2 regime. Where the operator is linear, on advection, the linear ridge is the lowest row and every other family lands within 0.4 points of it, as it should.

Second, a learned metric is what rescues a kernel on the state when the plain one fails, and it is worth most where the plain kernel is worst. On the correction-coefficient corpus from the state alone the isotropic kernel is at 62.93 ± 0.51 and a kernel-flow metric [27] takes it to 3.97 ± 0.42 , a factor of sixteen in the mean, which brings a kernel on the state to within six percent of the network on a corpus where the plain kernel was seventeen times worse. With the low-fidelity coefficients supplied as inputs the same pair reads 15.76 and 4.60, a factor of three: the easier the problem is made for the plain kernel, the less the learned metric adds. This is the same mechanism as the per-input length scales on the EMIT table and the empirical-Bayes metric of

corpus	n	d	q	runs	ridge (linear)	KRR	KRR, learned metric	network	network ensemble	network + residual KRR	kernel on features	convex stack
Advection	19,000	200	200	5	11.29±0.04	11.36±0.04	11.36±0.04	11.69±0.07	11.50±0.05	11.50±0.05	11.64±0.07	11.36±0.04
Burgers $\nu = 0.1$	7,200	256	128	3	13.89±0.53	2.81±0.28	2.90±0.33	20.34±0.61	19.28±0.09	2.78±0.29	4.92±0.30	2.78±0.29
Burgers $\nu = 0.01$	7,200	256	128	3	33.68±0.28	23.40±0.54	23.58±0.50	36.65±1.95	35.31±0.70	23.12±0.56	29.20±0.59	23.01±0.54
Burgers $\nu = 0.001$	7,200	256	128	3	38.38±0.12	31.66±0.16	33.23±1.73	39.98±0.65	38.89±0.11	31.82±0.11	35.59±0.48	31.34±0.14
Turbulent radiative layer	7,200	256	256	3	31.73	30.31	30.26±0.01	32.17±0.02	31.58±0.02	30.40±0.01	31.08±0.01	30.12
Active matter	14,000	256	256	1	37.24	35.81	35.18	37.31	36.93	35.77	36.22	35.05
Helmholtz staircase	20,384	256	256	1	0.00	0.26	0.07	0.40	0.24	0.23	0.45	0.07
ClimSim	100,000	124	128	5	32.24±0.06	27.20±0.05	26.31±0.04	25.10±0.18	24.43±0.17	24.33±0.13	24.71±0.25	24.25±0.17
RT corrections, state inputs	487,773–487,823	8	3	5	71.54±0.59	62.93±0.51	3.97±0.42	3.75±0.25	3.24±0.11	3.13±0.05	3.52±0.07	3.09±0.05
RT corrections, multi-fidelity inputs	487,773–487,823	11	3	5	33.83±0.26	15.76±1.98	4.60±0.31	2.79±0.14	2.40±0.11	2.34±0.09	2.80±0.09	2.32±0.08

Table 9: Mean relative L^2 test error [%] on ten further emulation configurations, mean $\pm$ standard deviation over the runs available for each. n is the number of examples — a range where the state-level split of the correction-coefficient corpus makes it differ by run — d the input dimension the models receive and q the output dimension they regress. The runs are random splits except for the turbulent radiative layer, whose three runs share The Well’s own split and differ in network initialisation, and the two single-run rows. The family set and the validation-based selection are common to every row; the network budget is not, and is 200 epochs at width 512 except on ClimSim (60 at 512) and the two correction-coefficient configurations (120 at 384). The ridge column is linear in the inputs the models receive, unlike the cubic reference of Table 1, and is not a member of the convex stack; its Helmholtz-staircase entry is 0.004, below the two decimals shown. Kernel solves are exact when the training block has at most 20,000 rows and otherwise use 6,000 Nyström landmarks, which is the case on ClimSim, the Helmholtz staircase and the two correction-coefficient rows.

Section 7, at a scale where its effect is unambiguous.

Third, the combination is at or near the best row of its own pool everywhere, and where it is not the lowest row of the table the reason is membership rather than weighting. The convex stack is the lowest row on eight of the ten configurations and the second lowest on the other two; the residual-corrected network is in the lowest two on five. Both exceptions are the linear problems, and on both the row that beats the stack is the ridge, which the stack does not contain. On advection the ridge is lower by 0.07 points while the stack, at 11.36%, matches the best member it has — the isotropic kernel, on which it puts weight 0.87 — to two decimals. On the Helmholtz staircase the ridge reaches 0.004%, sixteen times below the next row of the table and about a hundred times below the weakest member of the pool, and the stack reaches 0.066%, again a shade below its own best member, the learned-metric kernel at 0.067%, on which it puts weight 0.94. In both cases the stack does what it is asked to do inside its pool; what it cannot do is recover a family the pool excludes. The statement this supports is correspondingly narrow: a combination is a good default when no member of its own pool dominates it, and a combination can only be read against the whole comparison when the pool contains the families the comparison includes.

Three limits on scope. These rows use one training budget per corpus and one architecture family, and several of these corpora have published results from methods tuned for them specifically; the comparison here is between families under a common protocol, not against the state of the art on each problem. With three to five runs, and one on two rows, the spreads support the large orderings above and not fine distinctions between adjacent families. And every kernel column here inherits the tuning protocol Section 5.1 measures: the length scale is chosen on a 6,000-row subsample and the winner refit on every row, which over-smooths by an amount that grows with the ratio of the two. On the EMIT table that costs the isotropic kernel a factor 0.882; on ClimSim, where the subsample is a sixteenth of the solve, re-choosing the scale against the solve that actually runs halves it and lowers the kernel’s error by 3.6%; on a 20,000-row corpus it changes nothing. These examples show that tuning can affect the comparison. They

do not imply that every reported kernel error is pessimistic or that a reported gap is an upper bound: validation retuning can also increase test error. The columns of Table 9 describe the specified training and selection procedures, not the optimum attainable by either family.

9 Discussion

The strongest reproducible conclusion is metric- and budget-specific. On the six-dimensional EMIT table, an exact kernel with input-specific scales is substantially more accurate in radiance than the tested isotropic kernel, cubic polynomial and width-512 network. Residual fitting substantially improves the network under the implemented PCA projection. Widening the network to 2,000 units worsens its standalone radiance prediction under this training configuration but improves its kernel feature representation at all ten splits. This last contrast is more robust than the very small further radiance gain of the wide convex stack. The width and epoch-budget change is not a controlled test of an intrinsic disadvantage of wide neural networks.

The central qualification is that forward and inverse metrics do not select the same head. The input-scaled kernel has the best median reflectance error in the main comparison, whereas feature kernels have smaller 95th-percentile inverse errors. Both narrow and wide stacks worsen those tails relative to their feature kernels at every split. A single mean radiance ranking would miss this distinction. The finite-perturbation bound identifies the additional quantity that must be controlled for retrieval: the inverse denominator relative to the component errors. All-band medians alone are inadequate, and the unconditioned RMSE is strongly affected by near-pole predictions. A physically specified passband, explicit coverage and failure reporting, variable surface reflectances and a separate retrieval evaluation are needed before drawing operational conclusions.

The residual analysis establishes exact identities, an error-alignment criterion, and a separate PCA truncation term. It does not establish an asymptotic learning rate, automatic improvement over either parent, or decay of a neural residual in an unmeasured smoothness norm. The learning curves and rank ablations are empirical observations over the measured budgets. The full-row tuning experiment supports sensitivity to subsample tuning, but not a universal assertion that retuning lowers test error. The OCO-2 comparison and the additional corpora similarly delimit where the tested couplings are useful; they do not establish a universal structural taxonomy or state-of-the-art performance against task-specific models.

Computational cost is an additional qualification. An exact dense solve requires a Gram matrix with n^2 entries and a factorization costing order n^3 , followed by response solves; one $18,884 \times 18,884$ double-precision matrix alone occupies about 2.85 GB, before factorization workspaces and other arrays. Feature kernels also require evaluation of the network and the kernel against the training design at prediction time. A stack must evaluate all of its active members. The archived campaign times are not matched per-model throughput benchmarks, so the present comparison establishes accuracy under specified budgets, not an operational speed advantage or superiority to the existing sRTMnet retrieval pipeline.

Finally, repeated partitions of an already explored table quantify split sensitivity, not external validation. The main and wide tables and the new paired/tail analysis have complete per-seed public summaries. Some secondary ablations have aggregate records but incomplete individual archives; none has been reconstructed by inventing missing runs. Original data access, simulator provenance, numerical precision and training-environment records remain important for independent reproduction. A future confirmation should freeze the candidate pool and evaluation

protocol before using a genuinely untouched dataset. The present paper supports a narrower practical conclusion: representation learning and exact kernels can complement each other, but a coupling should be assessed in every downstream metric that matters, rather than promoted from a single favorable forward-error average.

Acknowledgements

The research presented in this paper was supported by the European Research Council (ERC) under the European Union’s Horizon Europe research and innovation programme (grant agreement No. 101041711), by the Simons Foundation as part of the Collaboration on the Mathematical and Scientific Foundations of Deep Learning, by Heights Labs, by Convex Nexus Capital, by the Israel Science Foundation (grant number 2258/19), and by the Israel Science Foundation (ISF Grant 4101/25).

Declaration of competing interest

The author holds equity in Heights Labs and in Convex Nexus Capital and has a management role at Convex Nexus Capital; both entities supported this work and are named above. These entities and the public funders had no role in the study design, analysis, manuscript preparation, or decision to publish.

Declaration of generative AI and AI-assisted technologies in the manuscript preparation process

During the preparation of this work the author made substantial use of generative AI tools, and describes their role here. The software — the data pipeline, the model comparison, the diagnostics, and the figures — was written and executed by Claude Code (Anthropic) and Codex (OpenAI), which also prepared the first draft of this manuscript, working from the author’s data and from the method and experimental recipe of [6]. The measurements of the data’s structure, the identification of the ill-posed reflectance entries, the scaling runs of Section 5 and the identity check of Section 6 were produced in the same AI-assisted sessions and verified against the run summaries. ChatGPT (OpenAI) subsequently assisted with the mathematical corrections, reanalysis of archived metrics, diagnostic tests, and manuscript revision. The author is responsible for final review of the code, results and text and for the content of the manuscript.

Data and code availability

The public repository is <https://github.com/yspennstate/emit-oco2-kernel-emulators>. It contains the manuscript, model drivers, stored summaries and regeneration code. The main EMIT campaign and the width-2000 pipeline each have all ten per-seed JSON records, including data-array SHA-256 digests, selected hyperparameters, split sizes and evaluation metrics. The paired and tail tables can be independently regenerated from those twenty records without retraining. The supplementary OCO-2/corpora tables are generated by `code/make_tables.py`. `docs/REPRODUCE.md` maps tables by label, separates aggregate from individual evidence, and records the commands and limitations.

Three secondary EMIT summary files retain the reported ten-run aggregates: `scaling_numbers.json`, `retune_numbers.json`, and `width_numbers.json`. Their individual archives are incomplete: at 8,000 rows one isotropic-campaign record and one input-scaled record are missing; only three width-1000 and five full-tuning records are present. Their aggregate findings are preserved as such, not relabeled as new reproductions. The three-member wide-feature observation is exploratory and does not supply ten-split evidence. The reproduction manifest identifies available and absent records.

The underlying EMIT arrays (`data_EMIT_24k`), attributed to the Jet Propulsion Laboratory in the supplied study materials, are not redistributed in this repository. The original simulator version, generation configuration and redistribution permissions are not specified in the public run records and must accompany an independently reproducible data release. The repository describes raw arrays, trained weights and prediction dumps as available from the author on request; access to these is not replaced by the public JSON summaries. Earlier campaign code exported predictions as float32 after scoring in float64; the revised export preserves float64, and the new conditioning diagnostic requires explicit opt-in for a precision-altered analysis of older dumps. No raw-array retraining or conditioned EMIT evaluation is represented by the new summary reanalysis.

The OCO-2 data and reference predictions originate in the release accompanying [5]; the additional corpora are the cited public releases. The method companion is [6]. Its comparison tables are not a substitute for establishing raw-data reproducibility of this EMIT study.

References

- [1] A. Berk, P. Conforti, R. Kennett, T. Perkins, F. Hawes, and J. van den Bosch. MODTRAN6: a major upgrade of the MODTRAN radiative transfer code. In *Proc. SPIE 9088, Algorithms and Technologies for Multispectral, Hyperspectral, and Ultraspectral Imagery XX*, 2014.
- [2] D. R. Thompson, V. Natraj, R. O. Green, M. C. Helmlinger, B.-C. Gao, and M. L. Eastwood. Optimal estimation for imaging spectrometer atmospheric correction. *Remote Sensing of Environment*, 216:355–373, 2018.
- [3] P. G. Brodrick, D. R. Thompson, J. E. Fahlen, M. L. Eastwood, C. M. Sarture, S. R. Lundeen, W. Olson-Duvall, N. Carmon, and R. O. Green. Generalized radiative transfer emulation for imaging spectroscopy reflectance retrievals. *Remote Sensing of Environment*, 261:112476, 2021.
- [4] M. A. A. Mazid and N. Rishe. Multi-fidelity emulation of atmospheric correction coefficients with physics-guided Kolmogorov–Arnold networks. *Remote Sensing*, 18(11):1826, 2026.
- [5] O. Lamminpää, J. Susiluoto, J. Hobbs, J. McDuffie, A. Braverman, and H. Owhadi. Forward model emulator for atmospheric radiative transfer using Gaussian processes and cross validation. *Atmospheric Measurement Techniques*, 18:673–694, 2025.
- [6] Y. Shmalo. Neural means and kernel corrections for operator learning. Preprint, ResearchGate, 2026. Code and per-run summaries at <https://github.com/yspennstate/neural-means-kernel-corrections>.
- [7] L. Berlyand, E. Sandier, Y. Shmalo, and L. Zhang. Enhancing accuracy in deep learning using random matrix theory. *Journal of Machine Learning*, 3(4):347–412, 2024.

- [8] Y. Shmalo, J. Jenkins, and O. Krupchytskyi. Deep learning weight pruning with RMT-SVD: increasing accuracy and reducing overfitting. *arXiv:2303.08986*, 2023.
- [9] P. Batlle, M. Darcy, B. Hosseini, and H. Owhadi. Kernel methods are competitive for operator learning. *Journal of Computational Physics*, 496:112549, 2024.
- [10] L. Lu, P. Jin, G. Pang, Z. Zhang, and G. E. Karniadakis. Learning nonlinear operators via DeepONet based on the universal approximation theorem of operators. *Nature Machine Intelligence*, 3:218–229, 2021.
- [11] Z. Li, N. Kovachki, K. Azizzadenesheli, B. Liu, K. Bhattacharya, A. Stuart, and A. Anandkumar. Fourier neural operator for parametric partial differential equations. In *International Conference on Learning Representations*, 2021.
- [12] G. R. Yoo and H. Owhadi. Deep regularization and direct training of the inner layers of neural networks with kernel flows. *Physica D: Nonlinear Phenomena*, 426:132952, 2021.
- [13] R. O. Green et al. The Earth surface mineral dust source investigation: an Earth science imaging spectroscopy mission. In *IEEE Aerospace Conference*, 2020.
- [14] C. E. Rasmussen and C. K. I. Williams. *Gaussian Processes for Machine Learning*. MIT Press, 2006.
- [15] A. Radhakrishnan, D. Beaglehole, P. Pandit, and M. Belkin. Mechanism for feature learning in neural networks and backpropagation-free machine learning models. *Science*, 383(6690):1461–1467, 2024.
- [16] A. Caponnetto and E. De Vito. Optimal rates for the regularized least-squares algorithm. *Foundations of Computational Mathematics*, 7(3):331–368, 2007.
- [17] D. J. Sutherland. Fixing an error in Caponnetto and de Vito (2007). *arXiv:1702.02982*, version 2, 2021.
- [18] S. Fischer and I. Steinwart. Sobolev norm learning rates for regularized least-squares algorithms. *Journal of Machine Learning Research*, 21(205):1–38, 2020.
- [19] M. V. de Hoop, D. Z. Huang, E. Qian, and A. M. Stuart. The cost-accuracy trade-off in operator learning with neural networks. *Journal of Machine Learning*, 1(3):299–341, 2022.
- [20] M. Takamoto, T. Praditia, R. Leiteritz, D. MacKinlay, F. Alesiani, D. Pflüger, and M. Niepert. PDEBench: an extensive benchmark for scientific machine learning. In *Advances in Neural Information Processing Systems 35*, 2022.
- [21] R. Ohana, M. McCabe, L. Meyer, R. Morel, F. Agocs, M. Beneitez, M. Berger, B. Burkhart, S. Dalziel, D. Fielding, et al. The Well: a large-scale collection of diverse physics simulations for machine learning. In *Advances in Neural Information Processing Systems 37*, pages 44989–45037, 2024.
- [22] S. Yu, W. M. Hannah, L. Peng, J. Lin, M. A. Bhouri, R. Gupta, B. Lütjens, J. C. Will, G. Behrens, et al. ClimSim: a large multi-scale dataset for hybrid physics-ML climate emulation. In *Advances in Neural Information Processing Systems 36*, 2023.

- [23] D. H. Wolpert. Stacked generalization. *Neural Networks*, 5(2):241–259, 1992.
- [24] L. Breiman. Stacked regressions. *Machine Learning*, 24:49–64, 1996.
- [25] A. G. Wilson, Z. Hu, R. Salakhutdinov, and E. P. Xing. Deep kernel learning. In *Proceedings of the 19th International Conference on Artificial Intelligence and Statistics*, pages 370–378, 2016.
- [26] J. M. Bates and C. W. J. Granger. The combination of forecasts. *Operational Research Quarterly*, 20(4):451–468, 1969.
- [27] H. Owhadi and G. R. Yoo. Kernel flows: from learning kernels from data into the abyss. *Journal of Computational Physics*, 389:22–47, 2019.